

* This set of notes follows several resources, including Prof. Lap Chi's and Prof. Halim's notes.
Unacademy CS3341 NUS CS3230

We start by considering two problems: the **traveling salesman problem** and the **Chinese postman problem**.

Say we have an undirected graph on n vertices and m edges, where each edge has some non-negative cost.

Traveling salesman: find minimum cost tour visiting every vertex at least once. \rightsquigarrow **naive:** $O(n!)$ \xrightarrow{DP} $O(2^n)$
NP-complete

Chinese postman: find minimum cost tour visiting every edge at least once. \rightsquigarrow surprisingly can be solved
in $O(n^m)$ (polynomial) time.

- Goal:**
- know well-known algorithms well \rightsquigarrow solve hard problems in polynomial time
 - design new algorithms
 - prove correctness and time complexity
 - use reductions to solve problems and prove hardness

Note on recursion: Many modern languages have recursion depth limits due to memory and stack unwinding constraints.

We cover **Master Theorem**, **(D&C, randomized, graph, greedy) algorithms**, **DP**, **amortized analysis**, and **NP-hardness**.
and recursion trees

But how do we define time complexity? What operations count as primitive?

Computation Models

We generally assume the **word-RAM model**: we can access an arbitrary position of an array in constant time, and that word-operations (addition, read/write, etc.) are constant-time.

Remark Suppose we have a graph with n vertices. We would need $\log n$ bits to identify each vertex

\rightsquigarrow we assume this fits in a word and assume primitive operations (e.g. ID comparisons) are $\Theta(1)$.

Sometimes this model is unrealistic \rightarrow we would analyse bit-complexity.

3SUM Given numbers a_1, \dots, a_n and c , we want to determine if there are i, j, k such that $a_i + a_j + a_k = c$.

Naive: enumerate all triples and check if sum is c . $\Theta(n^3)$

Alg. 2: Write $a_i + a_j + a_k = c \Leftrightarrow c - a_i - a_j = a_k$. Enumerate all pairs a_i, a_j and check if $c - a_i - a_j = a_k$ for some k .

1) Sort array $A = a_1 \leq \dots \leq a_n$. $\Theta(n \log n)$

$\Theta(n^2)$

sort + run binary search

2) Enumerate a_i, a_j and binary search for $a_k = c - a_i - a_j$.

$$\Rightarrow O(n \log n + n^2 \log n) = O(n^2 \log n).$$

Alg. 3: Reduce to a simpler problem. Notice $a_i + a_j + a_k = c \Leftrightarrow a_i + a_j = c - a_k$.

For fixed k , this is a **2SUM problem** \Rightarrow we reduce to n 2SUM problems.

one for each k

reduce to 2SUM

1) Sort $A = a_1 \leq \dots \leq a_n$. $\Theta(n \log n)$

2) Enumerate a_k and solve the 2SUM problem $a_i + a_j = c - a_k = b$. $\star O(n)$

$\Rightarrow O(n^2)$

We claim that 2SUM on a sorted array can be solved in $O(n)$.

Double-pointer approach. Keep a left-index L and right-index R .

Time complexity: stops when $L > R$, and L and R get closer

Proof.

every iteration \Rightarrow at most n iterations $\Rightarrow O(n)$.

While $L \leq R$:

if $a_L + a_R = b$, we are done
if $a_L + a_R > b$, $R--$
if $a_L + a_R < b$, $L++$

Correctness: If $\exists i, j$ such that $a_i + a_j = b$, we won't find any i, j .

Proof.

Say $\exists i, j$ such that $a_i + a_j = b$, for some $i \leq j$.

Alg. runs until $L=R \Rightarrow$ there is an iteration such that $L=i$ or $R=j$. at some point, one of these must be true.

WLOG suppose $L=i$ occurs first, and so $R>j$.

As long as $R>j$, we have $a_i + a_R > b$ (since A sorted) \Rightarrow we decrement R

\Rightarrow eventually we will have $R=j$.

General k-SUM Algorithms Exercise (see CS341 material)

Whether $O(n^2)$ was optimal for 3SUM has been a long-standing open problem.

Recently, there was an $O\left(\frac{n^2}{(\log n)^{2/3}}\right)$ algorithm discovered, but it remains to be proven if there is an $O(n^{2-\epsilon})$ time algorithm.

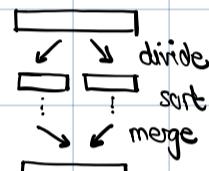
Solving Recurrences

Merge-Sort Classic divide-and-conquer sorting algorithm.

merge-sort:

```
sort(A[1..n]):
    if n=1: return
    sort(A[1..n/2])
    sort(A[n/2+1..n])
    merge(A[1..n/2], A[n/2+1..n])
```

reduce problem to smaller instances
of the same problem

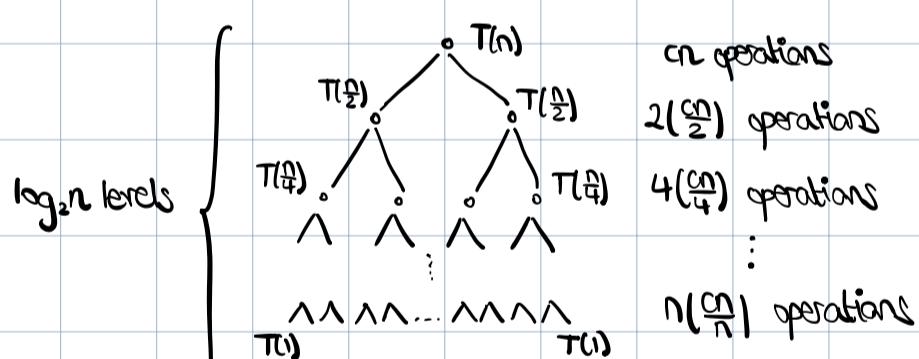


We have a recurrence

$$T(n) = \begin{cases} T\left(\frac{n}{2}\right) + T\left(\frac{n}{2}\right) + O(n) & \text{if } n > 1 \\ 1 & \text{if } n = 1 \end{cases}$$

$\Rightarrow T(n) = 2T\left(\frac{n}{2}\right) + O(n)$ sloppy recurrence.

One way to solve the recurrence is to draw the **recursion tree**.

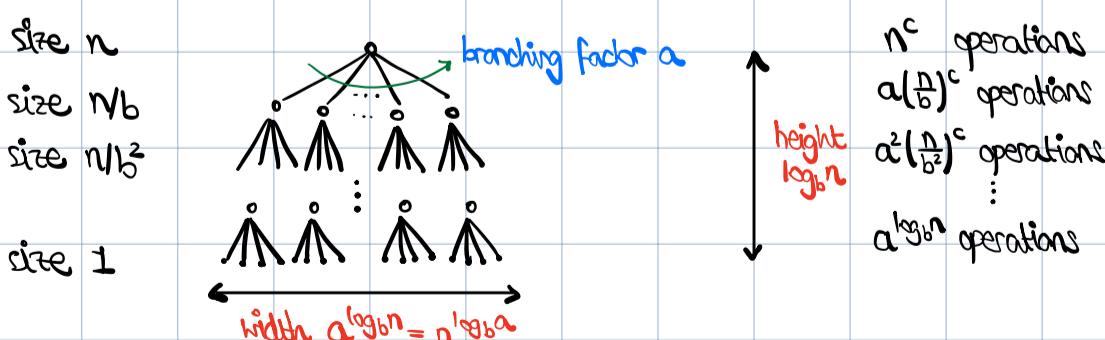


$\Rightarrow n$ operations for $\log_2 n$ levels $\approx O(n \log n)$ as expected

We can also solve this inductively.

Let's consider a more general setting.

Consider the relation $T(n) = aT\left(\frac{n}{b}\right) + n^c$ for some $a > 0$, $b > 1$, $c \geq 0$ constants.



We want to solve for

$$\sum_{i=0}^{\log_b n} a^i \left(\frac{n}{b}\right)^c$$

sum over all steps

\approx operations
in step i

$$\sum_{i=0}^{\log_b n} \left(\frac{a}{b}\right)^i n^c$$

a geometric series.

Case 1: $\left(\frac{a}{b}\right) = 1$: we get $\sum_{i=0}^{\log_b n} n^c = n^c (\log_b n + 1) \in \Theta(n^c \log_b n)$

Case 2: $\left(\frac{a}{b}\right) < 1$: decreasing geometric series dominated by the first term $\Rightarrow \Theta(n^c)$

Case 3: $\left(\frac{a}{b}\right) > 1$: increasing sequence dominated by last term.
 $\Rightarrow \Theta(a^{\log_b n}) = \Theta(n^{\log_b a})$

Theorem (Master): If $T(n) = aT(\frac{n}{b}) + n^c$ for constants $a \geq 0$, $b > 1$, $c \geq 0$, then:

$$T(n) \in \begin{cases} O(n^c) & \text{if } c > \log_b a \\ O(n^c \log n) & \text{if } c = \log_b a \\ O(n^{\log_b a}) & \text{if } c < \log_b a \end{cases}$$

Remark: Method more important than result.

Examples: 1 single subproblem.

$$T(n) = T\left(\frac{n}{2}\right) + 1, \text{ we have } T(n) \in O(\log n) \text{ (binary search)}$$

$$T(n) = T\left(\frac{n}{2}\right) + n, \text{ we have } T(n) \in O(n) \text{ geometric sequence (quick-select)}$$

$$T(n) = T(\sqrt{n}) + 1, \text{ we have } T(n) \in O(\log \log n) \text{ (counting levels)}$$

2 Non-even subproblems.

$$T(n) = T\left(\frac{2n}{3}\right) + T\left(\frac{n}{3}\right) + n \in O(n \log n)$$

$$T(n) = T\left(\frac{n}{2}\right) + T\left(\frac{n}{4}\right) + n \in O(n)$$

3 Exponential.

$$T(n) = 2T(n-1) + 1 \in O(2^n)$$

4 Fibonacci. Can we improve the runtime of $T(n) = T(n-1) + T(n-2) + 1$?

Consider recurrence $t_n = t_{n-1} + t_{n-2} \Rightarrow t_n - t_{n-1} - t_{n-2} = 0$. We use MATH 239 to create a closed-form expression to show that $t_n = \left(\frac{1+\sqrt{5}}{2}\right)^n$ so $T(n) \in O\left(\left(\frac{1+\sqrt{5}}{2}\right)^n\right) \approx O(1.618^n)$.

This is faster than $O(2^n)$ standard exponential time.

These sorts of recurrences appear often when analyzing algorithms (especially faster-exponential time algorithms).

Consider the **maximum independent set** problem: given a graph $G = (V, E)$, we want to find a maximum subset $S \subseteq V$ such that $u, v \in S \Rightarrow uv \notin E$.

Naive: enumerate all subsets $S \subseteq V$ to form powerset $P(V)$, and check each $X \in P(V)$
 $\Rightarrow 2^{2^{|V|}}$ time.

Variant: Pick $v \in V$ with maximum degree. Two possibilities: $v \in S$ or $v \notin S$.

$v \in S$: Reduce to $G-v$ (delete v) \rightsquigarrow reduce graph size by 1.

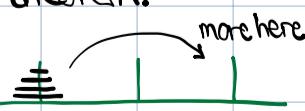
$v \notin S$: Then neighbours of v cannot be in S \rightsquigarrow delete v and all its neighbours \rightsquigarrow reduce graph size by ≥ 2 .

$$\Rightarrow T(n) \leq T(n-1) + T(n-2) + O(n) \rightsquigarrow T(n) \in O(n^2).$$

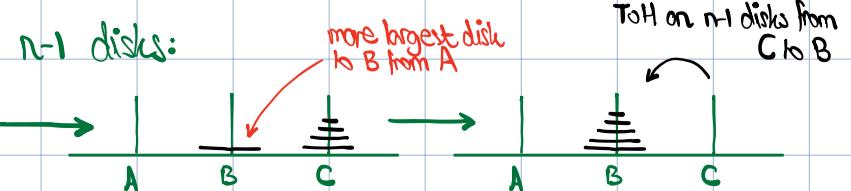
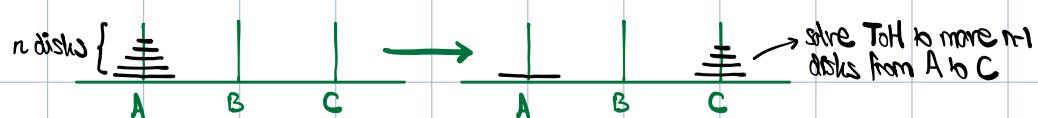
In general, we consider **4 methods for solving recurrences**:

telescoping series, substitution, recursion trees, and master theorem.

Substitution: Towers of Hanoi. • only one disk may be moved at a time
• no disk may be placed on a smaller disk



Notice a ToH on n disks can be reduced to a ToH on $n-1$ disks:



runtime
 $\Rightarrow T(n) = T(n-1) + T(n-1) + 1 = 2T(n-1) + 1$ with algorithm ToH(n, A, B, C):

intuitively, we can guess that $T(n) \in \Theta(2^n)$:

$$T(n) = 2T(n-1) + 1$$

$$= 2(2T(n-2) + 1) + 1 = 4T(n-2) + (2+1)$$

$$= \dots = 2^n T(n-n) + \sum_{i=0}^{n-1} 2^i = 2^n + \frac{2^n - 1}{2-1} \in \Theta(2^n)$$

From ToH(n, A, B, C):
 ToH(n-1, A, C, B)
 move A[i] to B
 ToH(n-1, C, B, A)

Consider $T(n) = aT(\frac{n}{b}) + f(n)$ where $f(n)$ is any generic function. We have three possible cases:

1 $f(n) \in O(n^{d-\varepsilon})$ for some $\varepsilon > 0 \Rightarrow T(n) \in \Theta(n^d)$

2 $f(n) \in \Omega(n^{d+\varepsilon})$ for some $\varepsilon > 0$ (and $a\frac{f(n)}{b} \leq c f(n)$ for some $c < 1$ \Leftarrow regularity condition) $\Rightarrow T(n) \in \Theta(f(n))$

3 $f(n) \in \Theta(n^d \log^k n)$ for some $k \geq 0$

} where $d = \log_b a$

Proof (1): $T(n) = aT(\frac{n}{b}) + f(n)$. Let $n = b^m$ so $T(b^m) = aT(b^{m-1}) + f(n) = a^m T(b^{m-m}) + \sum_{i=0}^{m-1} a^i f(\frac{n}{b^i})$.

We have $\underbrace{a^m T(b^{m-m})}_{a^{\log_b m} T(1)} + \underbrace{\sum_{i=0}^{m-1} a^i f(\frac{n}{b^i})}_{a^i O((\frac{n}{b^i})^{d-\varepsilon})} \leq a^{\log_b m} T(1) + \sum_{i=0}^{m-1} a^i c \left(\frac{n}{b^i}\right)^{d-\varepsilon}$

$$\begin{aligned} &= n^{\log_b a} T(1) + cn^{d-\varepsilon} \sum_{i=0}^{m-1} \left(\frac{n}{b^i}\right)^{d-\varepsilon} \\ &\quad \underbrace{cn^{d-\varepsilon} \sum_{i=0}^{m-1} \left(\frac{ab^\varepsilon}{b^i}\right)^i}_{= cn^{d-\varepsilon} \sum_{i=0}^{m-1} \left(\frac{ab^\varepsilon}{b^{i-\varepsilon}}\right)^i} = cn^{d-\varepsilon} \sum_{i=0}^{m-1} \left(\frac{ab^\varepsilon}{b^{\varepsilon}}\right)^i \\ &\quad \quad \quad \text{constant} \\ &= cn^{d-\varepsilon} \sum_{i=0}^{m-1} (b^\varepsilon)^i \\ &= cn^{d-\varepsilon} \left(\frac{b^m - 1}{b^\varepsilon - 1}\right) \\ &= cn^{d-\varepsilon} n^\varepsilon = Cn^d \end{aligned}$$

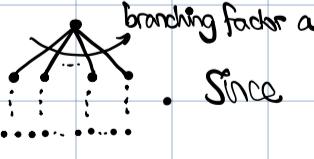
how does this prove a Θ bound?
 consider the other case of removing rightmost sum term.

Clearly this method is tedious, so we will again use the recurrence tree proof.

Theorem (Master): Let $T(n) = aT(\frac{n}{b}) + f(n)$ be a recurrence where $a > 0$, $b > 1$, and $f(n)$ is any function.

- 1 leaf-dominated $f(n) \in O(n^{d-\varepsilon})$ for some $\varepsilon > 0 \Rightarrow T(n) \in \Theta(n^d)$ for all n sufficiently large
- 2 root-dominated $f(n) \in \Omega(n^{d+\varepsilon})$ for some $\varepsilon > 0$ and $a\frac{f(n)}{b} \leq c f(n)$ for some $c < 1 \Rightarrow T(n) \in \Theta(f(n))$
- 3 balanced $f(n) \in \Theta(n^d \log^k n)$ for some $k \geq 0 \Rightarrow T(n) \in \Theta(n^d \log^{k+1} n)$ regularity condition

where $d = \log_b a$.

Proof: Consider the recurrence tree for $T(n)$:  Since $n \rightarrow \frac{n}{b}$ each level, we have height $h = \log_b n - 1$.

Notice that each node has "a" children

$$\Rightarrow \# \text{leaves} = a^{\log_b n} = n^{\log_b a} = n^d. \text{ Since each node has } \geq 1 \text{ child, we have } \# \text{nodes} \in \Theta(n).$$

Each node has runtime $f(\frac{n}{b^i})$ at level $i = 0, \dots, h$

$$\Rightarrow \text{total runtime across node operations is } \sum_{i=0}^h a^i f(\frac{n}{b^i}).$$

#nodes at level h
 per-node runtime

and total runtime across leaf operations is $n^d T(1)$.

$$1 \quad f(n) \in O(n^{d-\varepsilon}). \text{ Dominant term becomes leaf operations} \Rightarrow T(n) = n^d T(1) + \sum_{i=0}^h a^i f\left(\frac{n}{b^i}\right)$$

$$\Rightarrow T(n) \in \Theta(n^d). \quad \triangle$$

Cnd from above

2 $f(n) \in \Omega(n^{d+\varepsilon})$ and $a^i f\left(\frac{n}{b^i}\right) < c f(n)$. Dominant term becomes node operations.

$$n^d T(1) + \sum_{i=0}^h a^i f\left(\frac{n}{b^i}\right) \leq n^d T(1) + f(n) \sum_{i=0}^h c^i \leq n^d T(1) + \left(\frac{1}{1-c}\right) f(n) \in \Theta(f(n))$$

$\leq \sum_{i=0}^h c^i = \frac{1}{1-c}$ constant since $c < 0$

regularity $\begin{cases} f(n) + af\left(\frac{n}{b}\right) + a^2 f\left(\frac{n}{b^2}\right) + \dots \\ \leq f(n) + cf(n) + c^2 f(n) + \dots \end{cases} \Rightarrow T(n) \in O(f(n)).$

$$\text{and } n^d T(1) + \sum_{i=0}^h a^i f\left(\frac{n}{b^i}\right) = n^d T(1) + f(n) + \sum_{i=1}^h a^i f\left(\frac{n}{b^i}\right) \geq n^d T(1) + f(n) \in \Theta(f(n))$$

dominant

$$\Rightarrow T(n) \in \Omega(f(n)).$$

Thus $T(n) \in \Theta(f(n))$. \triangle

3 $f(n) \in \Theta(n^d \log^k n)$. Both terms are comparable.

$$\begin{aligned} \text{We have } T(n) &= n^d T(1) + \sum_{i=0}^h a^i f\left(\frac{n}{b^i}\right) = n^d T(1) + \sum_{i=0}^h a^i \left(\frac{n}{b^i}\right)^d \log^k\left(\frac{n}{b^i}\right) \\ &= n^d T(1) + n^d \sum_{i=0}^h \left(\frac{a^i}{b^{di}}\right) [\log^k n - i \log^k b] \\ &\leq n^d T(1) + n^d \log^k n \sum_{i=0}^h \left(\frac{a}{b^d}\right)^i = n^d \log^k n (h+1) \end{aligned}$$

$$n^d \log^k n (h+1) = n^d \log^k n \log_b n$$

$$= n^d \log^k n \underbrace{\frac{\log n}{\log b}}_{\text{constant}} \in n^d \log^{k+1} n \Rightarrow T(n) \in O(n^d \log^{k+1} n).$$

★ how to formulate a meaningful LB?

Exercise. \triangle

Therefore, we are done through a case-by-case analysis. \square

Remark: The Master Theorem does NOT cover all recurrences of the form $T(n) = aT\left(\frac{n}{b}\right) + f(n)$.

There are functions that will not satisfy any 3 of the cases.

Ex. Let $T(n) = T\left(\frac{n}{2}\right) + 2^{\sqrt{\log n}}$ so $a=1$, $b=2$, $d=\log_b a=0$.

Claim: $T(n) \notin O(n^{d-\varepsilon})$, $T(n) \notin \Omega(n^{d+\varepsilon})$, $T(n) \notin \Theta(n^d \log^k n)$ for any $k \geq 0$. ★ Exercise

Remark: We can write the general form of a recurrence $T(n) = aT\left(\frac{n}{b}\right) + f(n)$ as $n^d T(1) + \sum_{i=0}^h a^i f\left(\frac{n}{b^i}\right)$.

leaf ops. internal ops.

Exercise (AIQ2) Resolve recurrence $T(n) = 16T\left(\frac{n}{4}\right) + n^4$.

$$\begin{aligned} \text{We have (from recursive algo. general form) that } T(n) &= n^{\log_b a} T(1) + \sum_{i=0}^{\log_b n-1} a^i f\left(\frac{n}{b^i}\right) \\ &= n^2 + \sum_{i=0}^{\log_b n-1} 16^i \left(\frac{n}{4^i}\right)^4 \end{aligned}$$

$$n^4 \sum_{i=0}^{\log_b n-1} \left(\frac{16}{4}\right)^i = n^4 \sum_{i=0}^{\log_b n-1} \left(\frac{1}{2}\right)^i \in \Theta(n^4 \log n)$$

bounded by constants

Say we have a recurrence $T(n) = aT(\frac{n}{b}) + \underbrace{cn^j \log^k n}_{f(n)}$ for some

$$\begin{array}{l} a>0 \\ b>1 \\ c>0 \\ j>0 \\ k>0 \end{array}$$

When do we have some $\delta < 1$ such that $aT(\frac{n}{b}) \leq \delta f(n)$ for sufficiently large n (regularity)?

Notice $af(\frac{n}{b}) = a(\frac{n}{b})^j c \log^k (\frac{n}{b}) = ab^{-j} n^j c (\log n - \log b)^k < (ab^{-j}) n^j c \log^k n = (ab^{-j}) f(n)$ so we want $(ab^{-j}) < 1$.
In particular, $(ab^{-j}) < (ab^{-d}) = 1$ so $j > d$ needed.

Corollary: Say $T(n) = aT(\frac{n}{b}) + f(n)$ satisfying $f(n) \in \Omega(n^{d+\varepsilon})$ for some $a>0$, $b>1$, $\varepsilon>0$, where $f(n) = cn^j \log^k n$ for some $c>0$, $j \geq 0$.

Then $f(n)$ always satisfies the regularity condition.

Proof: Since $cn^j \log^k n \in \Omega(n^{d+\varepsilon})$ for some $\varepsilon>0$, the idea is that $n^j \log^k n \geq n^d n^\varepsilon \Rightarrow j > d = \log_b a$.

Formally, we have $cn^j \log^k n \geq Cn^{d+\varepsilon} = Cn^d n^\varepsilon > Cn^d \log^k n$ for sufficiently large n .

so $\Theta(f(n)) = \Theta(n^j \log^k n) \subseteq \Omega(n^{d+\varepsilon})$ and $\Theta(n^{d+\varepsilon}) \subseteq \underbrace{\Omega(n^{d+\varepsilon})}_{\text{given } n^j \geq n^{d+\varepsilon}} \subseteq \omega(n^d \log^k n)$ since $n^{d+\varepsilon} > n^d \log^k n$
 $\Rightarrow \Theta(n^j \log^k n) \subseteq \omega(n^d \log^k n)$ so $j > d$.

Therefore, we have $af(\frac{n}{b}) = a(\frac{n}{b})^j \log^k (\frac{n}{b}) = ab^{-j} n^j (\log n - \log b)^k < ab^{-j} n^j \log^k n = (ab^{-j}) f(n) = \delta f(n)$

where $\delta < 1$ as required. \square

Lastcode

1 merge-k-sorted-lists on k sorted singly-linked-lists.

hard

Naive: create a pointer for each list L_i (say pointers p_i). $O(k)$

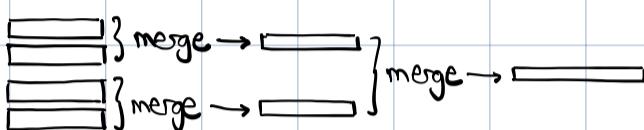
for each element, find $\min(L_i(p_i))$ and append element to a new linked list. $O(k) n$ times

increment that pointer p_i . $O(1)$

$\Rightarrow \Theta(nk)$ time, $\Theta(n)$ space

Merge: modify the merge step from merge-sort. We divide-and-conquer on k .

merge step from merge-sort: essentially naive algorithm on 2 lists.



since each SLL already sorted, we don't need to do the sorting step, just the merging step.

why can't we? because each step needs a merge step

given k sorted SLLs:

base case $k \leq 2$: naive merge

first half \rightarrow merge first $\lceil \frac{k}{2} \rceil$ lists } each list has n/k elements
second half \rightarrow merge last $\lceil \frac{k}{2} \rceil$ lists } n/k elements
merge first and second half

running analysis; would then prove by induction

}

$$\text{we have a runtime } T(k) = \begin{cases} 2T\left(\frac{k}{2}\right) + cn & \text{if } k > 2 \\ O(1) & \text{if } k \leq 2 \end{cases}$$

on average, n total in k lists $\Rightarrow \frac{n}{k}$ per list

\Rightarrow let $k = 2^m$ for some m , so we have

$$T(2^m) = 2T(2^{m-1}) + cn$$

$$= 2(2T(2^{m-2}) + cn) + cn$$

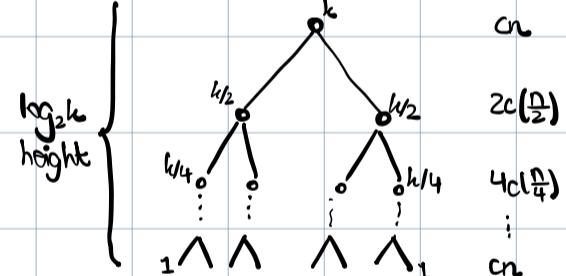
$$= \dots = 2^m T(2^0) + cn \sum_{i=0}^{m-1} 2^i. \quad \left. \begin{array}{l} k=2^m \Rightarrow m=\log k \\ \hline \end{array} \right.$$

$$\Rightarrow T(2^m) = 1 \cdot 2^{\log_2 k} + cn(2^m - 1) = n + cn \log k \in \Theta(n \log k)$$

from base case

We can merge with $O(1)$ auxiliary space

$\Rightarrow \Theta(n \log k)$ time, $O(1)$ space



Heaps Idea: in naive approach, we want to find smallest element from k pointers p_1, p_k :

We can do this using a min-heap instead of a linear search.

initialize all p_i to start of L_i :

add all $L_i(p_i)$ to a min-heap

for n iterations:

say p_j is the min in heap

pop min. element from heap and add to new list $O(1)$

increment p_j $O(1)$

add $L_i(p_j)$ to heap $O(\log k)$ since heap always has k elements

$\Rightarrow \Theta(n \log k)$ time, $O(n)$ space

can be reduced to $O(n)$

2 sort-colours: given an array of n elements in {0, 1, 2}, design an algorithm to sort the array.

med

Ideas: bucket-sort. $\Theta(n)$ time, $\Theta(n)$ space.

or radix sort

counting-sort: count occurrences of each category and override old array $\Rightarrow \Theta(n)$ time, $\Theta(1)$ space.

but what if each element $x \in A$ has additional information? counting-sort would lose info.

Modified: start with a first-pass counting each category ($\Theta(n)$), say n_0, n_1, n_2 counts.

then, initialize indices $i_0 = 0, i_1 = n_0, i_2 = n_0 + n_1$.

initialize temporary variable $t_1 = A[i_0], t_2 = \text{NULL}$

for n iterations: notice $t_i \in \{0, 1, 2\}$

$t_2 = A[i_{t_1}]$ the element that $A[t_1]$ should replace

$\Rightarrow \Theta(n)$ time, $\Theta(1)$ space without

$A[i_{t_1}] = t_1$

using any information

$t_1 = t_2$ next iteration will find space for the popped element
increment i_{t_1}

We must show the correctness of this algorithm, since it is not immediately obvious.

*Exercise

drop-sort: we want to use swapping to design a 1-pass, $\Theta(n)$ time, $\Theta(1)$ space algorithm.

Idea: initialize $i_0 = 0, i_2 = n - 1$

for $j = 0, j < n, j++:$

if $A[j]$ is 0: swap $A[j]$ and i_0

i_0++

if $A[j]$ is 2: swap $A[j]$ and i_2

i_2--

else pass 1s will naturally end up in the right place

3 remove-duplicates-sorted-ii: A: sorted array of n elements. remove dups such that each unique element appears at most twice without changing sorted order.

med

read/write approach: one pointer goes through entire array to read; when we hit a new element, we write.

initialize read index and write index $r_i = 1, w_i = 1$

for $r_i < n, r_i++:$

if $A[r_i] \neq A[w_i - 1]$: current element not same as previously written element

$A[w_i] = A[r_i]$ write

w_i++

otherwise r_i incremented (we hit a dupe \Rightarrow move on)

return $A[0:w_i - 1]$

4 k-radius-subarray-averages: array $A[0:n-1]$, some integer $k \geq 1$. Create array of average of subarray of A (subarray size k , centered at index i). Essentially sliding window average problem.

med

Divide and Conquer Algorithms

Counting Inversions

An **inversion** in an array A is a pair of indices (i, j) such that $i < j$ but $A[i] > A[j]$.

Consider the problem of **counting inversions**. \rightsquigarrow how "inverted" is a sequence?

\Rightarrow can we use D&C here? If we count #inversions in first half, then second half, can we use this info to count inversions of both halves put together?

\rightsquigarrow must count inversions

with one element in each half

\rightarrow cross-inversion pairs are easier to find because we know their relative positions

\rightsquigarrow sort each half (not lossy because we already counted in-half inversions!)

then we can just count elements $\geq a_{j+1}$ in second half for each $i \rightarrow$ can do this equiv. on the other half instead.

\Rightarrow reduced the problem to efficient counting problem. Let halves be H_1, H_2 .

Case 1: H_1 sorted \Rightarrow binary search for number $> a_j$ where $a_j \in H_2$ (\Rightarrow inversion)

$\Theta(\log n)$ time for each $a_j \Rightarrow \Theta(n \log n)$ time to find all cross-inversions.

Case 2: The info we want can be determined using a merge from merge-sort.

When we insert $x \in H_2$ to the merged list

\rightarrow naive merge

$\Rightarrow \Theta(n)$ time. **count($A[1..n]$):**

```
if n=1: return 0
divide { count(A[1..n/2]) ← T(n/2)
           count(A[n/2+1..n]) ← T(n/2)
conquer { sort(A[1..n/2]) ← Θ(n log n) ←----- Is this step necessary?
           sort(A[n/2+1..n]) ← Θ(n log n)
           merge-count(A[1..n/2], A[n/2+1..n]) ← Θ(n)
```

Therefore $T(n) = 2T\left(\frac{n}{2}\right) + f(n)$ where $f(n) \in \Theta(n \log n)$

$\Rightarrow T(n) \in \Theta(n \log^2 n)$ by **Masters Theorem**.

Merge Count **count-and-sort($A[1..n]$):**

```
if n=1: return 0
S1 = count-and-sort(A[1..n/2])
S2 = count-and-sort(A[n/2+1..n])  $\leftarrow$  remove redundant sorting step because we already sort
S3 = merge-count(A[1..n/2], A[n/2+1..n])
return S1 + S2 + S3
```

$\Rightarrow T(n) = 2T\left(\frac{n}{2}\right) + O(n) \in \Theta(n \log n)$. Theoretically best (based on comparison-based LB).

*Exercise: write merge-count pseudocode.

Maximum Subarray

Given array A of elements a_1, \dots, a_n , we want to find (i, j) that maximizes

$$\sum_{k=i}^j a_k.$$

Naive For every i , iterate j and compute $\sum_{k=i}^j a_k$.

Use sliding window for each addition to be $O(1) \Rightarrow$ we have time $O(n^2)$.

Would D&C help? If we had max subarrays for halves H_1, H_2 , would this help solve for A ?

\rightarrow optimal solution either contained in H_1, H_2 , or crosses between H_1 and H_2 .

→ after solving subproblem, we just need to find max. subarray that crosses $\frac{n}{2}$ (so $i < \frac{n}{2} < j$).

Assume that (i, j) is optimal with $i < \frac{n}{2} < j \Rightarrow [i, \frac{n}{2}]$ must be a max. subarray at $\frac{n}{2}$.

SEAC not, so $[i, \frac{n}{2}]$ max. subarray.

then $\sum_{k=i}^j a_k > \sum_{k=i}^j a_k$ contradicts (i, j) -maximality ↴

Claim: For $i < \frac{n}{2} < j$, we have $[i, j]$ is a maximum subarray crossing midpoint

↔ $[i, \frac{n}{2}]$ max. subarray ending at $\frac{n}{2}$ and $[\frac{n}{2}+1, j]$ max. subarray starting at $\frac{n}{2}+1$.

Proof: Above argument. □

Finding max. subarray with a fixed starting/ending point can be done in $O(n)$ time.

Therefore:

start with $A[1..n]$

find (i_1, j_1) for H_1 and compute $s_1 = \sum_{k=i_1}^{j_1} a_k$ $T(\frac{n}{2})$
find (i_2, j_2) for H_2 and compute $s_2 = \sum_{k=i_2}^{j_2} a_k$ $T(\frac{n}{2})$
compute whether $[i_1, \frac{n}{2}]$ max. subarray } $\theta(n)$
 $[\frac{n}{2}+1, j_2]$ max. subarray }

$$\Rightarrow T(n) = 2T\left(\frac{n}{2}\right) + \Theta(n)$$

$\in \Theta(n \log n)$.

yes $\Rightarrow [i_1, j_2]$ max. for A .
else $\max(s_1, s_2)$ max. for A . } $\Theta(1)$

* **Exercise:** Find an $\Theta(n)$ time algorithm for maximum subarray problem.
HARD

* **Exercise:** extend algo. to work in circular settings? ↗ just run algorithm
on concatenated array AA

Median Finding ↗ randomized
quick-select (see CS240) works; $\Theta(n)$ expected time. But this is a randomized algo.

problem with quick-select: the pivot choice needs to be good, but might be terrible.

Deterministic Algo. There is an interesting deterministic algorithm that would always return a number with

rank $\frac{3n}{10} \leq r \leq \frac{7n}{10}$ (somewhat close to median) in $O(n)$ time.

→ we can find a reasonable pivot in $O(n)$.

$$\Rightarrow T(n) \leq T\left(\frac{7n}{10}\right) + P(n) + c_2 \Rightarrow T(n) \in O(n).$$

↗ pivot selection time

problem size reduces at least by $\frac{3n}{10}$

Good Pivot Finding

The idea is to **find median of medians**.

1) Divide $A[1..n]$ into $\frac{n}{5}$ groups (each group has 5 numbers) in $O(n)$.

2) Find medians $b_1, \dots, b_{\frac{n}{5}}$ for each group, each $O(1)$ time $\Rightarrow O(n)$ time.

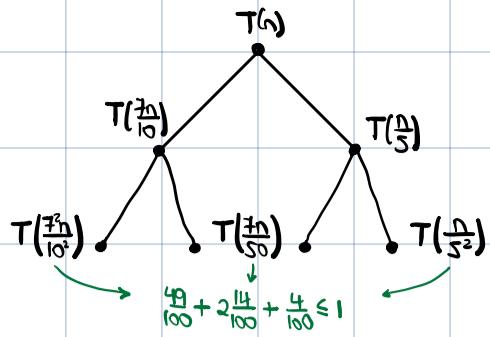
3) Find the median of $b_1, \dots, b_{\frac{n}{5}}$ by recursive call $\Rightarrow T\left(\frac{n}{5}\right)$ time. → we use mutual recursion
between pivot and median finding

Lemma: Let r be the rank of the median of medians. Then $\frac{3n}{10} \leq r \leq \frac{7n}{10}$.

Proof: See L03.pdf from CS341 notes.

Therefore, we have $P(n) = T\left(\frac{n}{5}\right) + c_2 n \Rightarrow T(n) \leq T\left(\frac{7n}{10}\right) + P(n) + c_2 n = T\left(\frac{7n}{10}\right) + T\left(\frac{n}{5}\right) + (c_1 + c_2)n$

Recursion Tree Analysis:



$c_3 n$ operations

$$c_3 \left(\frac{7}{10} + \frac{9}{10}\right) = c_3 \left(\frac{96}{100}\right) \text{ operations}$$

$$\vdots \leq c_3 n \text{ operations}$$

* Exercise complete.

Closest Pair

Input: n points $(x_1, y_1), \dots, (x_n, y_n) \in \mathbb{R}^2$.

Output: indices $1 \leq i < j \leq n$ such that Euclidean distance $\sqrt{(x_i - x_j)^2 + (y_i - y_j)^2}$ minimized.

Naive: Enumerate all pairs, $O(n^2)$.

D&C: Split the problem in half. Would finding the closest pair in each half help?

→ we can easily take the closer of the two pairs
→ suffices to find cross-half pairs

Suppose we halve by x -coordinate median.

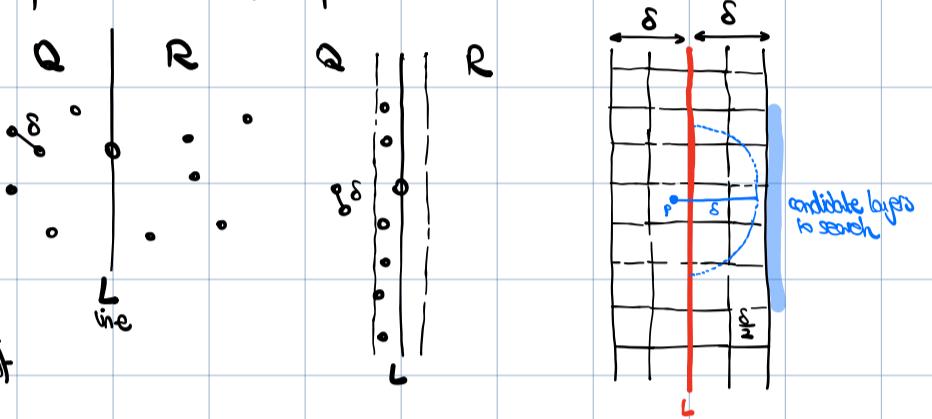
Notice we only need to check crossing pairs for points that are $< \delta$ distance from L

but this can still be all the points!

Idea: divide the δ -neighbourhood of L into a grid of

squares with side length $\frac{\delta}{2}$.

* We assume all points have distinct x -coord.



Lemma: Each box contains at most one point.

Proof: SFAC not \Rightarrow we have a pair within Q or R (non-crossing) with distance at most $\frac{\delta}{2} < \delta$. \square

Lemma: Each point only needs to compute distances with points within two horizontal layers (up and down).

Proof: Vertical squares are non-crossing. Diagram above shows that ± 2 horizontal layers stay δ -close, else not. \square

\Rightarrow each candidate point must compare with at most 11 other points \Rightarrow search space $O(n)$ squares.

Algorithm 1

we quickselect

1. Find line L by finding median. $O(n)$
2. Recursively solve problem in Q and R (halves) $2T\left(\frac{n}{2}\right)$
3. Using a linear scan, remove all points outside δ -neighbourhood. $O(n)$
4. Sort points by y -value $O(n \log n)$ bottleneck \leftarrow we don't need to do this in the recursion
5. For each point, compare distance to next 11 points in y -ordering. $O(n)$
6. Return minimum distance found.

$$\Rightarrow T_1(n) = 2T\left(\frac{n}{2}\right) + \Theta(n \log n) \stackrel{\text{Notes}}{\Rightarrow} T_1(n) \in \Theta(n \log^2 n).$$

If we sort beforehand we get $T(n) = 2T\left(\frac{n}{2}\right) + O(n) \in O(n \log n)$.

\rightarrow also allows us to remove the median-finding step

Remark: There is a randomized algorithm for closest pair with expected $O(n)$ time.

Arithmetic Operations with Dic

This is where Dic is most powerful. Eventually, this also allows us to parallelize operations for high performance.

See CUDA and tensor operations for instance.

A Integer Multiplication

Given n -bit numbers $a = a_0 \dots a_n$ and $b = b_0 \dots b_n$, we want to compute ab .

Grade-school multiplication requires $\Theta(n^2)$ operations.

Dic: Say we can do n bit multiplication efficiently. How can this help $2n$ bit multiplication?

Let x, y be $2n$ bit. Write $x = x_1 x_2$ and $y = y_1 y_2$ where x_1, x_2, y_1, y_2 are n -bit.

$$\Rightarrow x = x_1 2^n + x_2 \text{ and } y = y_1 2^n + y_2$$

$$\Rightarrow xy = (x_1 2^n + x_2)(y_1 2^n + y_2) = x_1 y_1 2^{2n} + (x_1 y_2 + 2x_2 y_1) 2^n + x_2 y_2$$

So $T(n) = 4T\left(\frac{n}{2}\right) + O(n)$ $\Rightarrow T(n) \in O(n^2)$ no improvement. Notice we did nothing clever to combine subproblems
extra (use bitshift)
extra (bitshift)
extra (use bitshift)
extra (bitshift)

So what can we do?

Karatsuba's Algorithm Reduce to three subproblems: $x_1 y_1$, $x_2 y_2$, and $(x_1 + x_2)(y_1 + y_2)$.

Compute middle term using $(x_1 + x_2)(y_1 + y_2) - x_1 y_1 - x_2 y_2 = x_1 y_2 + x_2 y_1$.

$$\Rightarrow T(n) = 3T\left(\frac{n}{2}\right) + O(n) \stackrel{\text{Master}}{\Rightarrow} T(n) \in O(n^{\log_2 3}) = O(n^{1.59}).$$

B Polynomial Multiplication

Same as Karatsuba. * Exercise

C Matrix Multiplication ($n \times n$)

Given two $n \times n$ matrices, suppose we can multiply them efficiently.

Say we have two $2n \times 2n$ matrices A, B .

$$\Rightarrow A = \begin{pmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{pmatrix} \text{ and } B = \begin{pmatrix} B_{11} & B_{12} \\ B_{21} & B_{22} \end{pmatrix} \text{ block matrices where blocks are } n \times n.$$

$$\Rightarrow AB = \begin{pmatrix} A_{11}B_{11} + A_{12}B_{21} & A_{11}B_{12} + A_{12}B_{22} \\ A_{21}B_{11} + A_{22}B_{21} & A_{21}B_{12} + A_{22}B_{22} \end{pmatrix} \Rightarrow \text{eight subproblems.}$$

$$\Rightarrow T(n) = 8T\left(\frac{n}{2}\right) + O(n^2) \in O(n^3) \text{ no improvement.}$$

Strassen's Algorithm

Smart reduction for matrix multiplication. Details omitted, but reduces matmul to $8 \rightarrow 7$ subproblems.

$$\Rightarrow T(n) = 7T\left(\frac{n}{2}\right) + O(n^2) \Rightarrow T(n) \in O(n^{\log_2 7}) = O(n^{2.81}).$$

Many combinatorial problems can be reduced to matmul $\xrightarrow{\text{Strassen}}$ these can be solved in $O(n^3)$ time.

* Ex. Determine whether $G = (V, E)$ has a triangle.

Fast Fourier Transform solves integer and polynomial multiplication in $O(n \log n)$. See CS487.

D Exponentiation given $a, n \in \mathbb{N}$, compute a^n .

Näive: iterative algorithm $O(n)$ (slow).

D&C: we can consider the rough idea of $a^n = a^{\frac{n}{2}} a^{\frac{n}{2}}$. we have 2 main cases here:
 n even and n odd.

if n even: $a^n = a^{\frac{n}{2}} a^{\frac{n}{2}}$

$\text{O}(1)$ operation*

if n odd: $a^n = a^{\lfloor \frac{n}{2} \rfloor} a^{\lceil \frac{n}{2} \rceil}$ or $a^n = a^{\lceil \frac{n}{2} \rceil} a^{\lfloor \frac{n}{2} \rfloor}$

$\text{O}(1)$ operation* → we assume these are $\text{O}(1)$ word operations

$$\Rightarrow T(n) = 2T\left(\frac{n}{2}\right) + \text{O}(1) \in \text{O}(\log n) \text{ exponential improvement over naive.}$$

E Quake Fast Inverse Square Root

* Exercise.

General Structures for Proof of Correctness

In CS245, we analyzed computational correctness proofs for simple algorithms. We generalize here.

Iterative Algorithms

Of the form initial conditions → for/while loop → termination.

- We do 3 things:
- 1 Show that the loop invariant is true at the start of the iterative process
 - 2 Show that the loop invariant is maintained (inductively)
 - 3 Prove that the iterative process terminates

→ choose a loop invariant such that the loop invariant itself helps prove correctness

Recursive Algorithms

Similar inductive approach.

- 1 Show correctness for base case of recursion
- 2 Show correctness for inductive step (usually strong induction)

More Arithmetic Operations

Polynomial Multiplication (Karatsuba's)

Say we have $A(x) = \sum_{k=0}^n a_k x^k$ and $B(x) = \sum_{k=0}^n b_k x^k$ so $\deg(A) = \deg(B) = n$.

We want to compute $C(x) = A(x)B(x)$ with $\deg(C) = 2n$ efficiently. Denote $C(x) = \sum_{k=0}^{2n} c_k x^k$.

Naive: Compute $c_k = \sum_{i+j=k} a_i b_j = \sum_{i=0}^k a_i b_{k-i}$ for every $k \in \{0, \dots, 2n\}$.

→ each c_k computation takes $k+1$ operations, so we have a runtime

$$\sum_{k=0}^{2n} (k+1) \in \Theta(n^2).$$

Idea (D&C): Say we can compute $\deg(\frac{n}{2}) \times \deg(\frac{n}{2})$ polynomial multiplication fast. How can this help?

Let $A(x) = A_1(x) + A_2(x)$ where $A(x) = \sum_{k=\lfloor \frac{n}{2} \rfloor + 1}^n a_k x^k + \sum_{k=0}^{\lfloor \frac{n}{2} \rfloor} a_k x^k$ and similarly $B(x) = B_1(x) + B_2(x)$.

$$\Rightarrow C = (A_1 + A_2)(B_1 + B_2)$$

$$= A_1 B_1 + A_1 B_2 + A_2 B_1 + A_2 B_2. \quad \text{However, we can do better subproblem division.}$$

The key idea is we want all subproblems to be $(\frac{n}{2}) \times (\frac{n}{2})$.

Notice if we instead write $A(x) = \sum_{k=0}^{\lfloor \frac{n}{2} \rfloor} a_k x^k + \sum_{k=0}^{\lceil \frac{n}{2} \rceil} a_k x^k = x^{\lfloor \frac{n}{2} \rfloor+1} \underbrace{\sum_{k=0}^{\lceil \frac{n}{2} \rceil} a_{k+\lfloor \frac{n}{2} \rfloor+1} x^k}_{A_1(x)} + \underbrace{\sum_{k=0}^{\lfloor \frac{n}{2} \rfloor} a_k x^k}_{A_2(x)}$
 then we have $C(x) = (x^{\frac{n}{2}} A_1(x) + A_2(x)) (x^{\frac{n}{2}} B_1(x) + B_2(x))$
 $= x^n \underbrace{A_1 B_1}_{\text{Master (leaf)}} + x^{\frac{n}{2}} \underbrace{A_1 B_2 + A_2 B_1}_{\text{Master (leaf)}} + x^{\frac{n}{2}} \underbrace{A_2 B_2}_{\text{Master (leaf)}} \text{ where each subproblem is } (\frac{n}{2}) \times (\frac{n}{2}).$

Since $x^{\frac{n}{2}} \sum_{i=0}^{\frac{n}{2}} d_i x^i = \sum_{i=0}^{\frac{n}{2}} d_i x^{i+n}$ can be done in $\mathcal{O}(n)$ for arbitrary coeffs. d_i , we have

$$T(n) = 4T(\frac{n}{2}) + \Theta(n) \text{ with } \log_2 4 = 2 \Rightarrow \Theta(n) \subseteq \mathcal{O}(n^{\log_2 4 - \varepsilon})_{\varepsilon = \frac{1}{2} > 0} = \mathcal{O}(n^{\frac{1}{2}})$$

$$\Rightarrow T(n) \in \Theta(n^{\log_2 4}) = \Theta(n^2). \text{ No improvement.}$$

The issue is that our subproblem space is too large. Karatsuba's algorithm reduces this space.

Karatsuba's

Notice that it suffices to compute $(A_1 + A_2)(B_1 + B_2) = A_1 B_1 + A_1 B_2 + A_2 B_1 + A_2 B_2$

$$\xrightarrow{\text{somehow}} x^0 A_1 B_1 + x^{\frac{n}{2}} A_1 B_2 + x^{\frac{n}{2}} A_2 B_1 + A_2 B_2.$$

However for the final transformation, we need access to each subproblem term individually.
almost!

$$= x^0 A_1 B_1 + x^{\frac{n}{2}} (A_1 B_2 + A_2 B_1) + A_2 B_2. \text{ We can reduce as follows:}$$

- compute $A_1 B_1$ term 1 and $A_2 B_2$ term 3 ($2T(\frac{n}{2})$)
- compute $(A_1 + A_2)(B_1 + B_2) - A_1 B_1 - A_2 B_2 = (A_1 B_2 + A_2 B_1)$ term 2 $T(\frac{n}{2}) + \Theta(n)$

$$\Rightarrow \text{we have } d = \log_2 3 > 1 \Rightarrow T(n) \in \Theta(n^{\log_2 3}) \approx \Theta(n^{1.57}).$$

Quake Fast Inverse Sort we want to compute $\frac{1}{\sqrt{x}}$.

Idea: 1 alias x to an integer to compute approx. $\log_2 x$.

2 use $\log_2 x$ to compute $-\frac{1}{2} \log_2 x = \log_2 \frac{1}{\sqrt{x}}$.

3 alias back to float to approximate $2^{\log_2 \frac{1}{\sqrt{x}}} = \frac{1}{\sqrt{x}}$. ← by this step, we have a (bad) approx of $\frac{1}{\sqrt{x}}$

4 refine using Newton's method.

```
float Q_sqrt(float number)
{
    long i2;
    float x2, y2;
    const float threehalves = 1.5f;

    x2 = number * 0.5f;
    y2 = number;
    i2 = +1; //long +1 kys
    i2 = 0x3f320000f - 1; i2 += 1; // evil floating point bit-level hacking
    y2 = +1.0f; //float +1.0f
    y2 = y2 * threehalves - (x2 + y2 * y2); // 1st iteration
    y2 = y2 / threehalves - (x2 + y2 * y2); // 2nd iteration, this can be removed
    return y2;
}
```

The details are in the type aliasing steps that require strong understanding of float representations.

Addition Machine

Consider a restriction of the word-RAM model called the addition machine that supports only $(+, -)$ and comparison.

Fact: $\text{div}(x, y) = x // y$ can be implemented in $\Theta(\log(\frac{x}{y}))$ time.

is-odd(n) can be implemented using $\text{div}(n, 2) \Rightarrow \Theta(\log n)$ time.

Goal: implement $\text{square}(n) = n^2$.

Nature: iteratively add n to a variable n times $\Rightarrow \Theta(n)$.

Diff: how can we break n^2 into recursive subproblems? Say $(\frac{n}{2})^2$ can be computed fast.

$$n^2 = 4(\frac{n}{4})^2 = 4(\frac{n}{2})^2 \text{ where we can add } (\frac{n}{2})^2 \text{ to itself 4 times in } \mathcal{O}(1).$$

If $n=2^k$ for some k , we clearly have $T(n) = T(\frac{n}{2}) + O(1) \Rightarrow T(n) \in \Theta(\log n)$.

However, if $n \neq 2^k$, then we would have to use `is_odd(n)` before recursive call, so

$T(n) = T(\frac{n}{2}) + \Theta(\log n)$ balanced $\Rightarrow T(n) \in \Theta(\log^2 n)$. This is the high-level idea.

Leetcode

5 binary search on offset sorted array (search-rotated-sorted-array) med

Goal: Given sorted array rotated arbitrarily, search for an element k

→ e.g. [6, 7, 8, 1, 2, 3, 4, 5] If we know shift value (e.g. smallest element index), we can solve easily

⇒ we want to find smallest element in $O(\log n)$ first.

D&C: Say we split into L and R, and can find minimum for both. Then $m = \min(m_L, m_R)$.

$$\Rightarrow T(n) = 2T\left(\frac{n}{2}\right) + O(1) \in O(\log n).$$

Then we simply offset our binary search by the index of this minimum element.

D&C 2: At least one half of A must be fully sorted. → possible for both halves sorted, e.g. [4, 5, 6, 1, 2, 3]

1 Check which half sorted ($O(1)$) time by checking first and last elements). WLOG A_L sorted.

2 If $k \in A_L$ sorted, regular binary search. Else recursive call on A_R .

$$\Rightarrow T(n) \leq \max\left(T\left(\frac{n}{2}\right), c\log n\right) + O(1).$$

If $\max\left(T\left(\frac{n}{2}\right), c\log n\right) = c\log n$, then $T(n) \in \Theta(\log n) \Rightarrow T\left(\frac{n}{2}\right) \in \Theta(\log n)$.

else $\max\left(T\left(\frac{n}{2}\right), c\log n\right) = T\left(\frac{n}{2}\right) \Rightarrow T(n) \in O(\log n)$.

∴ $T(n) \in O(\log n)$.

6 pow-2c: calculate x^n given $x \in \mathbb{Z}$ and $n \in \mathbb{N}_{>0}$ (not both zero). med

Naïve: $O(n)$ by iteration.

D&C: If we know $x^{\frac{n}{2}}$ and $x^{\frac{n}{2}}$, then $x^n = x^{\frac{n}{2}} x^{\frac{n}{2}}$. Therefore we do such a binary divide.

This gives runtime $T(n) = 2T\left(\frac{n}{2}\right) + O(1) \in \Theta(\log n)$.

pow(x, n):

```
if n=0: return 1
if n=1: return x
return pow(x, ⌈n/2⌉) * pow(x, ⌊n/2⌋)
```

7 perfect-square: given n , identify if n is a perfect square without using `sqr()`. med

Naïve: Iterate i from 1 to n . Calculate i^2 , check against n .

$i^2 < n \Rightarrow$ continue

$i^2 = n \Rightarrow$ perfect square $\Rightarrow T(n) \in O(\sqrt{n})$ because we stop at $i = \lceil \sqrt{n} \rceil$.

$i^2 > n \Rightarrow$ terminate

Binary Search: Search for $\lceil \sqrt{n} \rceil$ in $[1..n]$. Say search range is $[1..y]$.

Calculate $(\lceil \sqrt{n} \rceil)^2$. If $> n$, then $\lceil \sqrt{n} \rceil$ in left half-range. If $< n$, in right half-range.

else $= n \Rightarrow$ done.

Recursive case: we search desired half-range.

$$\Rightarrow T(n) = T\left(\frac{n}{2}\right) + O(1) \Rightarrow T(n) \in O(\log n).$$

8 spells-and-potions: input: two positive integer arrays S and P, length n and m; integer "success". med

a potion pair is successful if $S[i]P[j] \geq \text{success}$.

output: integer array A of length n where $A[i] = \# \text{potions that will form a successful pair with } S[i]$.

Nature: Direct-address table $A[i..n]$. Enumerate all pairs and update $A \Rightarrow O(mn)$.

Sort First: Sort S and P . Then we know that $S[i]P[j] \geq \text{success} \Rightarrow S[i..n]P[j..m] \geq \text{success}$.

\Rightarrow for each $S[i]$, we just need $P[j]$ such that $P[j] \geq \frac{\text{success}}{S[i]}$.

Binary search on P to find j such that $P[j] \geq \frac{\text{success}}{S[i]}$.

$$\text{so } T(n) = \underbrace{\Theta(n\log n)}_{\text{sort } S} + \underbrace{\Theta(m\log m)}_{\text{sort } P} + \underbrace{\Theta(n\log m)}_{\text{binary search}} \text{ times}$$

Notice we don't use the fact that $S[i]$ sorted. \Rightarrow we can avoid sorting S .

$$\xrightarrow{\text{reduce}} T(n) = \Theta(m\log m) + \Theta(n\log m).$$

Sorting both and using some smart algorithm does not help, since $n\log n$ or $m\log m \geq m\log n$ and $n\log m$ based on $n > m$ or $n \leq m$.

What if S and P were sorted? Can we speed things up more?

Consider the observation that $S[i]P[j] \geq \text{success} \Rightarrow S[i..n]P[j..m] \geq \text{success}$. Let i be fixed.

Let j_i be the smallest index such that $S[i]P[j_i] \geq \text{success}$

$\Rightarrow j_i$ is a lower-bound index for $i..l$, so we only must look for $j_i \geq j_i$ from P .

in practice, potentially limits search space

9 snapshot-array: data structure `SnapshotArray(int length)` that:

- starts with array of zeros of specified length

- `set(index, val)` sets $A[\text{index}] \leftarrow \text{val}$

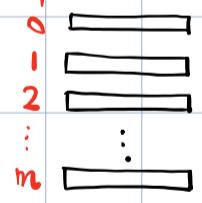
starts at 0, increases by 1 every snap

`snap()` takes a "snapshot" of the array and returns a snapshot ID

`get(index, snapId)` gets index at $A[\text{index}]$ at that snapshot

Binary Search every snapshot stored as an array \Rightarrow we have a matrix.

`snapid`



\Rightarrow binary search on `snapid` to retrieve array, then return element at index `idx`.

`get()` time $O(\log m)$

`set()` time $O(1)$

`snap()` time $O(1)$ amortized

} space $\Theta(mn)$

Direct Addressing Store data in matrix as above.

Then `get(idx, snapid)` returns $M[n * \text{snapid} + \text{idx}]$.

`get()` time $O(1)$
`set()` time $O(1)$
`snap()` time $O(1)$ amortized

} space $\Theta(mn)$

Sparse Space Optimization Idea: each index i keeps a sorted list of pairs `(snapid, value)` so

$A = [i \downarrow \cdot \cdot \cdot \cdot \cdot]$ where each $A[i]$ only updates when that value updates.

$[(s_1, v_1), \dots, (s_n, v_n)]$

for `get(i, s)` we get the sorted pairs list for $A[i]$ ($O(1)$)

then binary search $A[i]$ for the largest $\text{snapid} \leq s$, and return corresponding value

- 10** **2D-matrix search:** $M \in \mathbb{Z}^{m \times n}$ with each row of M sorted, and last element of row $r_i <$ first element of r_{i+1} .
 med Search for value k in M . \leadsto we can binary search
 1 Search largest $v \leq k$ along $r_i[0]$ (first element of each row)
 2 Once some r_i identified, binary search through r_i 's elements for k
 $\Rightarrow \Theta(\log n + \log m) = \Theta(\log mn)$ runtime

- 11** **Find-peaks:** find local maxima of an array A (elements greater than their neighbours).
 easy Observation: worst-case we have $A = [0, 1, 0, 1, \dots, 0]$ so there are $\Omega(n)$ peaks
 \Rightarrow our algorithm will be at best $\Omega(n)$.

Linear Search: l_i, r_i, m_i indices left, right, and middle.

```

 $l_i = 0, r_i = 2, m_i = 1$ 
peaks = []
while  $m_i < n-1$ :
    if  $A[l_i], A[r_i] < A[m_i]$ :
        peaks.append(m_i)
        notice  $r_i$  cannot be a peak  $\Rightarrow$  skip one
         $l_i, r_i, m_i += 2$ 
    else:
        not a peak
         $l_i, r_i, m_i += 1$ 

```

- 12** **Find-peak-element:** given A , find any peak of A .
 med

Näive: same as **11 Find-peaks**, runtime $\Theta(n)$.

Idea: Rec. Given arrays A_L and A_R , how can we find a peak?

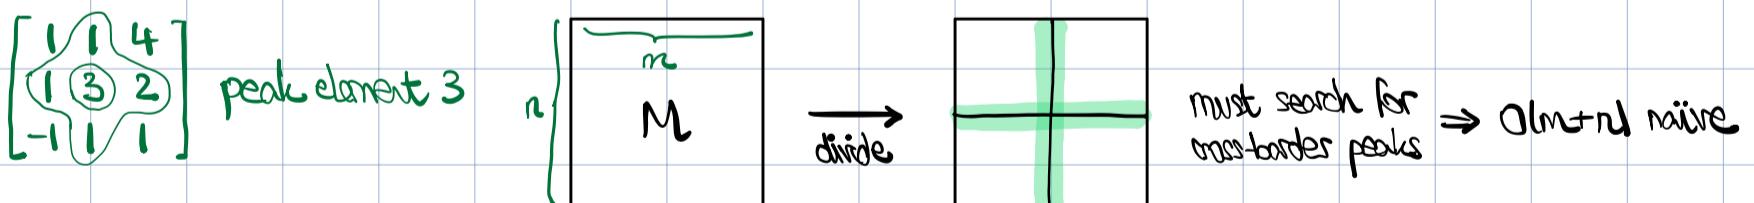
either A_L has a peak \Rightarrow return

or A_R has a peak \Rightarrow return

or there is a peak in $A_L[-2], A_L[-1], A_R[0], A_R[1]$ \Rightarrow runtime $T(n) = 2T\left(\frac{n}{2}\right) + O(1) \in \Theta(\log n)$

else no peaks

- 13** **Find-peak-element-2D:** find a local maximum of a matrix where neighbours are top, left, right, bottom.
 med



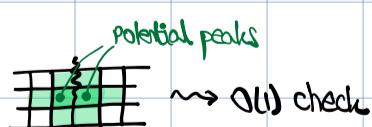
\leadsto can we find cross-border peaks in $\Theta(\log(m) + \log(n))$?

Re-use **Find-peak-element** and modify for 2D case. Goal: given $S = \overbrace{\begin{array}{|c|c|c|} \hline \end{array}}^n \{ }^3$ shape, find peaks (in middle row)

Rec: split into S_L and S_R .

S_L or S_R has a peak \Rightarrow done

else peak must be crossing between:
 or no peak



\Rightarrow finding cross-block peaks is $\underbrace{2T_S(n)}_{\text{search block crossings (vert.)}} + \underbrace{2T_S(m)}_{\text{search block crossings (horiz.)}}$ where $T_S(n) = 2T_S\left(\frac{n}{2}\right) + O(1) \in \Theta(\log n)$

\Rightarrow cross-border peaks

searchable in $\Theta(\log m + \log n)$ so $T(m, n) = 4T\left(\frac{n}{2}, \frac{m}{2}\right) + \Theta(\log m + \log n) \in \Theta(\max(n \log m, m \log n))$

more *

14 med group-anagrams: given an array of strings, group all anagrams together into sublists.

Say we have n strings and $m = \max_{s \in A} \text{len}(s)$ (longest string length).

Naive: Generate a character-count hashmap for each string $\Rightarrow n$ hashmaps $O(nm)$

Then check pairwise equality and group accordingly. $O(n^2)$

Prop.: Let h_1 and h_2 be hashmaps. Then we can compute whether $h_1 \neq h_2$ in $O(1)$ expected amortized time.

Idea: check equality of the hashes of h_1 and h_2 instead of element-wise. "D"

Remark: holds for arbitrary size but since our hashmaps are $O(1)$ size, we can guarantee $O(1)$ equality-checking

$\Rightarrow \Theta(n^2 + mn)$ time; bottleneck is equality-checking.

Hashing Arrays: Initialize arrays A_1, \dots, A_n of size 26 to act as our character counters.

Then construct a hashmap where keys are A_i and values are string arrays.

group-anagrams(A):

initialize $A_1, \dots, A_n = [0, \dots, 0]$ character count arrays $O(n)$

initialize $H = \text{dict}()$ 26

for $s \in A$:

build A_s by character-counting s $O(m)$ due to linear scan

if $A_s \in H.\text{keys}()$: $O(1)$ check for existence

$H[A_s].\text{append}(s)$

else:

$H[A_s] = [s]$ $O(1)$

don't use sets here because we want to preserve duplicates

* NOTE: arrays are an unhashable type so we encode A_i as strings (e.g. $[0, 1, 0] \mapsto "0..0"$) and use these as keys

$\Rightarrow O(nm)$ runtime and $O(n)$ space

15 med top-k-frequent: integer array $A[0..n-1]$, and some $k \leq n-1$. Return k most frequent elements in A .

Naive: linear scan and counting \rightsquigarrow element count hashmap \rightsquigarrow sort by frequency.

linear scan $O(n)$

$\Rightarrow O(n \log n)$ time

sort (k, v) pairs by v $O(n \log n)$

Idea: we must solve the sorting bottleneck.

1 randomized-quick-select

Notice we don't need to fully sort, just finding element of rank- k suffices

\Rightarrow modify randomized-quick-select to also return the left (\geq) subarray in $\Theta(n)$ expected time

rq-s(A, k):

pivot = rand(len(A)-1)

pivot-val = $A[\text{pivot}]$

$A_L = [], A_R = []$

for $a_i \in A$:

add to A_L or A_R based on $\leq \text{pivot-val}$

return recursive call

expected $\Theta(n)$ time but worst-case $\Omega(n^2) \Rightarrow$ problem.

k-frequent(A, k):

$H = \text{dict}()$ $O(1)$

count { for num in A:

$H[num] \leftarrow H[num] + 1$ else initialize $H[num] = 1$ } $O(n)$

return rq-s(H, k) $O(n)$ expected

2 bucket-sort frequency of an element can be at most n

\Rightarrow create n buckets B_1, \dots, B_n

perform linear occurrence count using hashmap $O(1)$

iterate through hashmap adding element to corresponding occurrence bucket $n \times O(1)$

read top k non-empty buckets from n to 1 $O(n)$

so runtime $\Theta(n)$ and space $\Theta(n)$.

not algo. design

16 encode-decode: A is a list of n strings with $m = \sum_{s \in A} \text{len}(s)$. Encode s into one string and then decode.

med

Näive: choose some non-(UTF-8) character as a delimiter and concat/split naïvely.

We want to not depend on this.

Idea: use some information about string lengths and delimit accordingly.

what if we had numbers representing string block lengths?

\rightsquigarrow how to avoid confusing with numbers in the string?

" $\ell_1 \# \underbrace{s_1}_{\ell_1} \ell_2 \# s_2 \dots \ell_n \# s_n$ " where $\#$ delimiter and $\ell_i = \text{len}(s_i)$

number of non-leading
characters before ℓ_{i+1}

17 sqrt-x: given $x \in \mathbb{N}_{\geq 0}$, compute \sqrt{x} rounding down.

easy

Binary Search: Consider $A = [0..x]$. Split A halfway to form A_L and A_R . Then $\sqrt{x} \in A_L$ or $\sqrt{x} \in A_R$.

If $A_R[0]^2 > x$ then $\sqrt{x} \in A_L \Rightarrow$ recurse. Say $n = x$.

else $\sqrt{x} \in A_R \Rightarrow$ recurse.

$$\Rightarrow T(n) = T\left(\frac{n}{2}\right) + O(1) \in \Theta(\log n)$$

18 peak-in-mountain: given an integer "mountain" array with one peak, find index of peak.

med

Näive linear scan $\Theta(n)$.

RFC: peak must either be in A_L or A_R . We can reduce this to a max. search.

search-max(A):

$L_{\max} = \text{search-max}(A_L)$

$R_{\max} = \text{search-max}(A_R)$

return $\max(L_{\max}, R_{\max})$

values increase, peak, then decrease

19 h-index: given citations array C, where $C[i] = \# \text{citations for } i^{\text{th}}$ paper, return h-index of C.

med

Näive: sort and then linear search $\Rightarrow \Theta(n \log n)$ time. Many similar approaches have $\Theta(n \log n)$ time.

Quickselect: can we modify quickselect for this problem?

choose pivot p randomly

$A_L < p$ and $A_R \geq p$

if $A_R[\text{size}] < p$: \rightsquigarrow want largest p such that $A_R[\text{size}] \geq p$

$P \leftarrow P/2$ so $h \leq p$

retry with new p

does not seem like this would be any faster

Binary Search $h \in [0, n]$ so we can run a binary search with some subroutine $\text{is-h}(k)$ that determines if

A has k elements $\geq k \Rightarrow \Theta(n \log n)$ binary search.

Notice $h \in [0, \min(\max(A), n)]$ so if $\max(A) \ll n$ then this can be even faster ($\Theta(n \log \max(A))$).

Faulting Sort: buckets $B[i]$ where we have $M = \min(\max(A), n)$ buckets

\Rightarrow add papers to buckets and identify h -index using a linear scan. $\Theta(n)$ time, $\Theta(n)$ space.

20 med product-array-except-i array of integers A , we want to return an array P such that $P[i] = \prod_{j \neq i} A[j]$.

Naive: Precompute $\prod_{j=0}^n A[j]$ full product. Then calculate $B[i] = p / A[i]$ where $p = \prod_{j=0}^n A[j]$ pre-computed.

Constraint: cannot use division.

Idea: compute forward and backward product arrays $F[i] = \prod_{j=0}^{i-1} A[j]$ and $B[i] = \prod_{j=i+1}^n A[j]$

$\Rightarrow P[i] = F[i] B[i]$ $\Theta(n)$ time, $\Theta(n)$ space

Notice $|P| \in \Theta(n) \Rightarrow$ our algorithms will always be $\Omega(n)$ so $\Theta(n)$ is optimal.

21 med longest-consecutive-sequence array A of integers, determine the length of the longest consecutive numeric sequence

that can be formed using A .

Naive: $\Theta(n^2)$ checking, or Naive: sort and check, $\Theta(n \log n)$. We want $\Theta(n)$. \leftarrow but without radix sort, that would be too easy

Hashing: Convert A so that we have a hashmap H where $H[a_i] = 1$.

Find $m = \min(A)$ and start at $H[m]$. Pop $H[m]$ and look for $H[m+1]$.

If $m+1 \in H$, pop and continue with $m+2$

else: continue to next element? * how can this be done efficiently though

Identify all potential sequence starts: for $a_i \in A$, check if $a_{i-1} \in H$. If no, then a_i potential start.

let S be the set of sequence starts

for $s \in S$:

do above pop-and-find process to identify sequence length

\Rightarrow get longest such sequence

$\Theta(n)$ time, $\Theta(n)$ space

22 med edit-distance word1 and word2, and 3 valid operations: insert character, delete character, edit character
Find least number of operations to convert word1 \rightarrow word2.

Idea: If $|word1| < |word2|$ then we must add \Rightarrow must find most optimal place to add

$|word1| > |word2|$ we must remove \Rightarrow find most optimal deletions

ops = 0

Look for chars of word2 in-order within word1

ops += number of missing word2 chars

ops += length difference

cat \underline{bot} \rightarrow $1+0 = 1$ operation

cats \underline{bot} \rightarrow $1+1 = 2$ operations

intention
execution

010001
11001
↓
010001
010001 } 2 operations

Dynamic Programming Problem

PA1.Q1.1 bloom-day-bouquets med integers array A, and some $m, k \in \mathbb{Z}$.

We want to make m bouquets. To make a bouquet requires k adjacent flowers from a garden.

The garden has n flowers, and the i^{th} flower blooms at $A[i]$ and can then be used in exactly one bouquet.

Find the minimum number of days to make m bouquets of size k , and return -1 if impossible.

(a) The largest output depends on the values in A. Consider $A = [1, M]$ and $M=2$ and $k=1$.

Then the algorithm will return $\max(M, 1)$.

(b) Say $\text{can}(d)$ is an $\Theta(n)$ subroutine that returns

We can binary search for $d \in [0, \overbrace{\max(a)}^{M}]$

\Rightarrow runtime $\Theta(T(\text{can}) \log M) = \Theta(n \log M)$.

{ True if can make m bouquets using k flowers on day d
False otherwise }

1 Find $M = \max(a) \rightarrow O(n)$

2 Set $l, r = 0, M$

$$s = \frac{r-l}{2}$$

$\log M \times$

```

while l <= r:
    if l == r and can(s) True:
        return s
    if can(s) True:
        r = s
    else:
        l = s
        s = (r-l)/2
  
```

(c) $\text{can}(d)$: need $\Theta(n)$ algorithm to decide whether we can make m k -bouquets on day d .

Linear search (?) keep a flower counter C, more linearly, reset if a flower not bloomed.

```

C = 0 current flower count
b = 0 bouquet count
for i = 0 to n-1:
    if A[i] <= d: C += 1
    else: C = 0
    if C >= k:
        b += 1
        C = 0
    if b >= m:
        return True
if b < m:
    return False
  
```

PA1.Q1.2 minimum-repair-time A: integers representing ranks of m mechanics.

mechanic with rank r can repair n cars in $r n^2$ minutes.

$n = \text{number of cars}$

largest possible output:

We have 1 mechanic of rank $R \Rightarrow$ takes $R n^2$ time to repair n cars. Make R arbitrarily large.

Say $\text{can}(t)$ Θ(n) subroutine: whether n cars can be repaired by time t .

We can binary search $t \in [0, \max(r) n^2]$

problem: this might be large
this is okay because algorithm will absorb any exponents

$\text{Can}(t)$:

$\text{cap} = 0$ repair capacity by time t

for r in A:

return \sqrt{r}
 $x = r // t$ $x = n^2$
 $cap += \text{floor}(\sqrt{x})$ $O(1)$ using quick sqrt
if $cap \geq n$:
 return True } can make $\geq n$ cars by time t
return False

Binary search use subroutine to binary search $t \in [0, \max_{r \in A}(r)n^2]$.

$$\Rightarrow \text{runtime } \Theta(n \log(Rn^2)) = \Theta(n \log R + 2n \log n) = \Theta(n \log(nR))$$

Randomized Algorithms We study more randomized algorithms beyond CS240.

Freivald's Algorithm Given three $n \times n$ matrices $A, B, C \in \mathbb{R}^{n \times n}$, we want to check whether $AB = C$.

Nature: Compute AB using Strassen's algorithm and check if $AB = C$ elementwise $\Rightarrow \Theta(n^{2.8})$ time.

Idea: Use randomization to reduce the problem to a simple equality checking.

Choose $v = \begin{pmatrix} v_1 \\ v_2 \\ \vdots \\ v_n \end{pmatrix}$ randomly such that $v_i \in \{0, 1\}$ for each i . e.g. $v_i = \begin{cases} 0 & \text{with prob. } \frac{1}{2} \\ 1 & \text{with prob. } \frac{1}{2} \end{cases}$

Check if $ABv = Cv$ (notice $ABv \in \mathbb{R}^n$ and $Cv \in \mathbb{R}^n$) $\rightsquigarrow \Theta(n^2)$ time with matrix-vector mult.

$ABv = Cv \Rightarrow$ predict $AB = C$ notice $ABv = (AB)v = A(Bv)$ so AB need not be computed.
 $ABv \neq Cv \Rightarrow$ $AB \neq C$ deterministically instead we can check if $ABv - Cv = 0$ using $A(Bv) - Cv = 0$ in $\Theta(n^2)$.

$$\begin{array}{c} \cancel{\Theta(n^2)} \\ \cancel{\Theta(n^2)} \\ \Theta(n^2) \end{array}$$

The non-trivial case is when $AB \neq C$ but $ABv = Cv$ which we must analyze.

Correctness Analysis Consider $D = (AB - C)$ so $AB = C \Rightarrow \ker(D) \supseteq \mathbb{R}^n$. Denote $Dv = \begin{pmatrix} u_1 \\ u_2 \\ \vdots \\ u_n \end{pmatrix}$.

If $\begin{pmatrix} u_1 \\ u_2 \\ \vdots \\ u_n \end{pmatrix} = Dv = 0 \in \mathbb{R}^n$ but $AB \neq C$, then Freivald's produces an incorrect answer.

Freivald's produces an incorrect answer $\Rightarrow ABv = C$ but $AB \neq C$.

so Freivald's incorrect $\Leftrightarrow ABv = Cv$ but $AB \neq C \Leftrightarrow v \notin \ker(D)$ but $D \neq 0 \in \mathbb{R}^{n \times n}$.

$\Rightarrow \exists (i, j)$ such that $D_{i,j} \neq 0$. Notice $u_i = D_{1,i}v_1 + D_{2,i}v_2 + \dots + D_{i,i}v_i + \dots + D_{n,i}v_n = z + D_{i,j}v_j$.

and $Dv = \begin{pmatrix} u_1 \\ u_2 \\ \vdots \\ u_n \end{pmatrix} = 0 \in \mathbb{R}^n \Rightarrow u_i = z + D_{i,j}v_j = 0$ with $D_{i,j} \neq 0$.

Once we reveal our random choices $v_1, \dots, v_{i-1}, v_{i+1}, \dots, v_n$, term z is fixed.
 (Principle of Deferred Decision)

Notice at most one choice between $v_j = \{0, 1\}$ can yield $u_i = 0$ since $D_{i,j} \neq 0$.

This choice has probability $\frac{1}{2} \Rightarrow \Pr(u_i \neq 0) \geq \frac{1}{2}$

$\Rightarrow \Pr(\text{Freivald's correctness}) \geq \frac{1}{2}$.

\rightsquigarrow amplify success probability by repeating k times independently. Let X_1, \dots, X_k be results.

Then $\Pr(X_1 \text{ failure}, X_2 \text{ failure}, \dots, X_k \text{ failure}) = \prod_{i=1}^k \Pr(X_i \text{ failure}) \leq \prod_{i=1}^k \frac{1}{2} = \frac{1}{2^k}$

failure: algorithm produces wrong output
 (i.e. false positive) $\Rightarrow \Pr(\text{correctness}) \geq 1 - \frac{1}{2^k}$.

Therefore Freivald's has $\Theta(n^2)$ time with $\geq 1 - \frac{1}{2^k}$ success probability.

Coupon Collector Problem There are n types of coupons. Each box of cereal contains one random coupon.

How many boxes of cereal to collect all n coupons?

{reframe}

m balls randomly placed in n bins, what is $\Pr(\text{every bin has } \geq 1 \text{ ball})$?

Consider one bin B_i . Notice $B_i \sim \text{Binomial}(m, \frac{1}{n})$ so $\Pr(B_i = 0) = \underbrace{(1 - \frac{1}{n})^m}_{\text{failure prob.}} = \underbrace{\left(1 - \frac{1}{n}\right)^m}_{1+x \leq e^x} \leq e^{-m/n}$

$$\begin{aligned} &\Rightarrow 1-x \leq e^{-x} \text{ for } x \leq 1 \\ &\Rightarrow (1-x)^m \leq e^{-mx} \end{aligned}$$

$$\Rightarrow \Pr(B_1 = 0 \text{ or } B_2 = 0 \text{ or } \dots \text{ or } B_n = 0) = \Pr\left(\bigcup_{i=1}^n (B_i = 0)\right)$$

$$= \sum_{i=1}^n \Pr(B_i = 0) \leq n \left(1 - \frac{1}{n}\right)^m \leq n e^{-\frac{m}{n}}$$

so $m = 2n \ln(n) \in \Theta(n \ln n)$ boxes guarantees $\Pr(\text{collect } n \text{ types}) \geq 1 - \frac{1}{n}$.

We introduce a useful result for calculating expected runtime and success probability.

Lemma (Markov's Inequality): Let X be a non-negative random variable and $a > 0$. Then $P(X \geq aE[X]) \leq \frac{1}{a}$.

Proof: By definition $E[X] = \int_{\mathbb{R}} xf(x) dx \stackrel{x \text{ non-neg.}}{=} \int_{[0, \infty)} xf(x) dx$ so we have

$$E[X] = \underbrace{\int_{[0, a]} xf(x) dx}_{\geq 0} + \underbrace{\int_{[a, \infty)} xf(x) dx}_{\geq 0} \geq \int_{[a, \infty)} xf(x) dx \geq \min_{x \in [a, \infty)} (x) \underbrace{\int_{[a, \infty)} f(x) dx}_{\text{def.}} = a \int_{[a, \infty)} f(x) dx = aP(X \geq a).$$

Thus $E[X] \geq aP(X \geq a) \Rightarrow \frac{E[X]}{a} \geq P(X \geq a)$.

Using $b = \frac{a}{E[X]} > 0$ gives $\frac{E[X]}{a} = \underbrace{\frac{1}{b}}_{\text{done}} \geq P(X \geq bE[X]) = P(X \geq a)$. \square

This is very useful in converting expected runtime to probabilistic runtime.

$P(T \geq aE[T]) \leq \frac{1}{a}$ where T is runtime.

$\Rightarrow P(T \geq 100E[T]) \leq \frac{1}{100}$. In particular, there is an upper bound on the probability of "bad" T .

If expected runtime $E[T] \in \Theta(n \log n)$ then $P(T \geq 100 \Theta(n \log n)) \leq 0.01$

$\Rightarrow P(T \leq 100 \Theta(n \log n)) \geq 0.99$ so $P(T \in \Theta(n \log n)) \geq 0.99$.

In general, this gives rise to two types of randomized algorithms:

Las Vegas algorithms: $P(\text{correctness}) = 1$ but time complexity guarantee is only in expectation.
randomized-quicksort
i.e. we have $E[T] \in O(f(n))$

Monte Carlo algorithms: $P(\text{correctness}) \leq 1$ but time complexity guarantee always holds.
Freivalds algorithm
i.e. we have $T \in O(f(n))$

Lemma: All Las Vegas algorithms are Monte Carlo algorithms.

Proof: Consider using Markov's inequality. Let \mathcal{A} be a Las Vegas algorithm.

$\stackrel{\text{def.}}{\Rightarrow} E[T(\mathcal{A})] \in O(f(n))$ and $P(\mathcal{A} \text{ correct}) = 1$

so we have $\frac{1}{a} \leq P(T(\mathcal{A}) \geq a \underbrace{E[T(\mathcal{A})]}_{\substack{\text{Markov} \\ \text{inequality}}} \leq P(T(\mathcal{A}) \geq acf(n)) \leq c f(n)$

$\Rightarrow P(T(\mathcal{A}) \in O(f(n))) \geq 1 - \frac{1}{a}$.

Modify \mathcal{A} by setting a time limit $a = 100E[T(\mathcal{A})]$ so we have a new \mathcal{A}' with

$P(\mathcal{A}' \text{ correct}) \geq 1 - \frac{1}{a}$ and $P(T(\mathcal{A}') \in O(f(n))) = 1$ so \mathcal{A}' is Monte Carlo. \square

from STAT 240 Notes

Proposition (Markov's Inequality): For any random variable X with $E[|X|^k] < \infty$, we have $P(|X| \geq c) \leq \frac{E[|X|^k]}{c^k}$ for all $c > 0$, for all $k \in \mathbb{Z}_{>0}$.

Proof: Say X_i are continuous. Similar proof for X_i discrete.

We have $E[|X|^k] = \frac{1}{c^k} \int_{-\infty}^{\infty} |x|^k f_X(x) dx$ and $R = \{|x| \geq c\} \cup \{|x| < c\}$:

$$= \frac{1}{c^k} \left(\underbrace{\int_{|x| \geq c} |x|^k f_X(x) dx}_{\geq 0} + \int_{|x| < c} |x|^k f_X(x) dx \right)$$

definition

$$\geq \frac{1}{c^k} \int_{\{|x| \geq c\}} |x|^k f_x(x) dx \geq \frac{1}{c^k} \underbrace{\min_{\substack{x \in \{|x| \geq c\} \\ x=c}}}_{x=c} \int_{\{|x| \geq c\}} f_x(x) dx = \int_{\{|x| \geq c\}} f_x(x) dx = P(|X| \geq c). \quad \square$$

Designing Random Algorithms

Reservoir Sampling

We are given some collection (a_1, \dots, a_n) where n is unknown apriori, and want to return a value uniformly at random.

We can use a strategy called reservoir sampling:

e.g. when sampling from a linked list

```
ans = None
for a_i in (a_1, ..., a_n):
    prob = 1/i
    if r_i:
        ans = a_i
return ans
```

This method gives $P(\text{ans} = a_i) = \frac{1}{n}$ at termination.

future states

$$\text{Notice } P(\text{ans} = a_i) = P(r_1=1, r_{i+1}=0, \dots, r_n=0)$$

$$\stackrel{\text{indep}}{=} P(r_1=1) \prod_{k=i+1}^n P(r_k=0)$$

$$= \frac{1}{i} \prod_{k=i+1}^n \left(1 - \frac{1}{k}\right) = \frac{1}{i} \left(\frac{i}{i+1}\right) \left(\frac{i+1}{i+2}\right) \cdots \left(\frac{n-1}{n}\right) = \frac{1}{n}.$$

Fisher-Yates Shuffle

We have an array $A = [a_1, \dots, a_n]$ and want to shuffle A such that any permutation is equally likely.

This is useful to generate random permutations of $[1, \dots, n]$ for arbitrary n .

```
for (i=n-1; i>0, i-):
    j ← randint(1..n-i)
    swap A[i] and A[j]
```

⇒ we shuffle in $O(n)$ with $O(1)$ auxiliary space (optimal).

It remains to show that the algorithm is correct.

Proposition (Fisher-Yates): The Fisher-Yates shuffle on $A = [1..N]$ is a random variable X such that

$$P(X = B_N) = \frac{1}{N!} = \frac{1}{N!} \text{ for all } B_N \in \mathcal{I}_N.$$

Proof: WLOG start with $A = [1, \dots, N]$. Consider arbitrary $B = [b_1, \dots, b_N] \in \mathcal{I}_N$.

$$\text{Then } P\left(\bigcap_{i=1}^N \{b_i = i\}\right) \stackrel{\text{indep}}{=} \prod_{i=1}^N P(b_i = i \mid \bigcup_{k=0}^{i-1} \{b_k = k\}) \text{ by conditional probability.}$$

For the first random swap, clearly a_1 is only touched once (when $i=0$ in the algorithm)

$$\Rightarrow P(b_1=1) = \frac{1}{N}. \text{ Clearly } P(b_2=2 \mid b_1=1) = \frac{1}{N-1}. \text{ Inductively, } P(b_i=i \mid b_{i-1}=i-1, \dots, b_1=1) = \frac{1}{N-i+1}.$$

$$\Rightarrow P(B=A) = P\left(\bigcap_{i=1}^N \{b_i = i\}\right) \stackrel{\text{indep}}{=} \prod_{i=1}^N P(b_i = i \mid \bigcup_{k=0}^{i-1} \{b_k = k\}) = \prod_{i=1}^N \frac{1}{N-(i-1)} = \prod_{i=0}^{N-1} \frac{1}{N-i} = \frac{1}{N!} \text{ as required. } \square$$

Thus the Fisher-Yates shuffle is correct.

Random Point Generation

Say we want to generate a random point in a circle of radius r centered at 0. There are two main approaches:

1. Generate (r, θ) pair. Notice we want the density of generated points to be uniform in the circle, so

we sample $x \sim \text{Unif}(0, \sqrt{r})$ and $\theta \sim \text{Unif}(0, 2\pi)$ so (x^2, θ) defines a random point.

2. What if we want some random point in a slightly more complex region?  or ?

We can use randomization.

Idea: Define a simple cover for our target region R (e.g. cover with a rectangle $[x_1, x_2] \times [y_1, y_2]$).

Sample $p \sim [x_1, x_2] \times [y_1, y_2]$ uniformly. If $p \in R$, return p . Else, retry.

target region R , area A_R

bounding box $B = [x_1, x_2] \times [y_1, y_2]$, area A_B

while True:

sample $x \sim B$ uniformly

If $x \in R$:
return x } probability $\frac{A_R}{A_B}$

Therefore our runtime depends on the success $P = \frac{A_R}{A_B} \leq 1$.

$$\text{Notice } E[T] = \sum_{t=1}^{\infty} t P(T=t)$$

$$= \sum_{t=1}^{\infty} t P(x_t \in R, x_{t-1} \notin R, \dots, x_1 \notin R)$$

$$\stackrel{\text{indep}}{=} \sum_{t=1}^{\infty} t P(x_t \in R) \prod_{k=1}^{t-1} P(x_k \notin R)$$

$$= \sum_{t=1}^{\infty} t \frac{A_R}{A_B} \left(1 - \frac{A_R}{A_B}\right)^{t-1} = L \text{ converges.}$$

Example

For instance, say R unit circle and B bounding square, so

$$E[T] = \sum_{t=1}^{\infty} t \frac{\pi}{4} \left(1 - \frac{\pi}{4}\right)^{t-1} = \frac{\pi}{4} \sum_{t=1}^{\infty} t \left(\frac{4-\pi}{4}\right)^{t-1} = \underbrace{\frac{\pi}{4} \sum_{t=1}^{\infty} (t-1) \left(\frac{4-\pi}{4}\right)^{t-1}}_{\frac{\pi}{4} \sum_{t=1}^{\infty} t \left(\frac{4-\pi}{4}\right)^t} + \underbrace{\frac{\pi}{4} \sum_{t=1}^{\infty} \left(\frac{4-\pi}{4}\right)^{t-1}}_{\frac{\pi}{4} \left(1 - \left(\frac{4-\pi}{4}\right)^2\right)}$$

$\Rightarrow E[T] \in O(1)$ expected.

Moderately Small Element

Recall earlier, we saw a deterministic median-finding algorithm that runs in $O(n)$.

We also know that smallest-element search takes $\Omega(n)$.

Can we find a moderately small element in $O(n)$?

Let A be an array of n integers. We say $a \in A$ is moderately small if $\text{rank}(a) \leq \frac{1}{10}$.

Notice $P(a \in A \text{ moderately small}) = P(\text{rank}(a) \leq \frac{1}{10}) = \frac{1}{10}$ for uniformly random $a \in A$.

Choose some set of indices $i_1, \dots, i_s \in \{0, \dots, n-1\}$ randomly with replacement, and returning $\min_k (a_{i_k})$

$$\begin{aligned} \text{gives } P(\min_k (a_{i_k}) \text{ moderately small}) &= P\left(\bigcup_{k=1}^s \{a_{i_k} \text{ moderately small}\}\right) \leftarrow \text{at least one is moderately small} \\ &= 1 - P\left(\bigcap_{k=1}^s \{a_{i_k} \text{ not moderately small}\}\right) \\ &\stackrel{\text{indep}}{=} 1 - \prod_{k=1}^s P(a_{i_k} \text{ not moderately small}) = 1 - \left(1 - \frac{1}{10}\right)^s. \end{aligned}$$

Say we want our algorithm to have a success rate $P(\text{success}) \geq 1 - \frac{1}{n}$.

\Rightarrow we can use randomization and require $1 - \left(1 - \frac{1}{10}\right)^s \geq 1 - \frac{1}{n}$ so $\left(1 - \frac{1}{10}\right)^s \leq \frac{1}{n}$

$$\Leftrightarrow \left(\frac{9}{10}\right)^s \leq \frac{1}{n} \Leftrightarrow s \log\left(\frac{9}{10}\right) \leq -\log n \Leftrightarrow s \log\left(\frac{10}{9}\right) \geq \log n \Leftrightarrow s \geq \log\left(\frac{10}{9}\right) \log n$$

so choose $s = C \log n \in \Theta(\log n)$. Selecting $A[0:s]$ and returning $\min(A[0:s])$ has $P(\text{success}) \geq 1 - \frac{1}{n}$.

$\Rightarrow \Theta(\log n)$ runtime. We can randomize subarray selection and repeat to amplify success rate.

Graph Algorithms I

We study some simple graph searching algorithms. We start with **breadth-first search** and **depth-first search**.
 (BFS) (DFS)

Let $G = (V, E)$ be an undirected graph.

From CS135/136, we know there are many representations of a graph symbolically.

We will use the convention $n = |V|$, $m = |E|$.

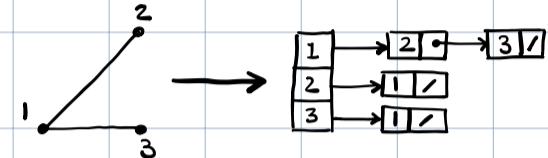
Adjacency matrix: represent vertices as rows/columns of an $n \times n$ matrix.

For G unweighted (edges do not have an associated weight), we can define $A \in \mathbb{F}^{n \times n}$ by $a_{ij} = \begin{cases} 1 & \text{if } ij, ji \in E \\ 0 & \text{if } ij \notin E \end{cases}$

If G is weighted with $w_{ij} \in \mathbb{R}_{\geq 0}$ for all $i, j \in V$, then define $a_{ij} = \begin{cases} w_{ij} & \text{if } ij \in E \\ \infty & \text{if } ij \notin E \end{cases}$

$\Rightarrow \Theta(n^2)$ space.

undirected \Rightarrow both same



Adjacency list: each vertex maintains a linked list of its neighbours.

$$\Rightarrow |V| + \sum_{v \in V} \deg(v) = 2|E| + |V| \in \Theta(m+n) \text{ space.}$$

we will frequently default to this representation.

Some basic questions, given a graph G , are:

- 1) Is G connected?
- 2) What are the connected components of G ?
- 3) Are some $u, v \in V$ connected?
Find a shortest uv -path.

* Graph BFS search through neighbours before depth

Input: graph $G = (V, E)$, source $s \in V$. **Output:** all vertices reachable by s (i.e. its component).

```

visited[v] = false  $\forall v \in V$ 
queue Q empty
enqueue(Q, s), visited[s] = true
while Q ≠ ∅:
    u = dequeue(Q)
    for each neighbour v of u:
        if visited[v] = false:
            enqueue(Q, v)
            visited[v] = true
    
```

how we know which vertex to search next in our BFS.

breadth-first



Time: each $v \in V$ enqueued at most once.
 when $v \in V$ dequeued, the loop executes $\deg(v)$ times.
 $\Rightarrow O(n + \sum_{v \in V} \deg(v)) = O(n+m)$ time

Lemma: For $v \in V$, $\exists sv$ -path in $G \Leftrightarrow \text{visited}[v] = \text{true}$ at the end of the algorithm.

Proof: (\Leftarrow) Induction on #steps of algorithm. We show $\text{visited}[v] = \text{true} \Rightarrow \exists sv$ -path.

Base case step 0, when only $\text{visited}[s] = \text{true}$.

Suppose $\text{visited}[v]$ is set to true at step k (within the for-loop for $u \in V$).

$\Rightarrow u \in V$ was dequeued $\Rightarrow u$ was enqueued at some previous step $\Rightarrow \text{visited}[u] = \text{true}$

IH $\Rightarrow \exists su$ -path in G . Since v is added in the for-loop for u , we have $uv \in E$.

Extend su -path by $uv \Rightarrow \exists sv$ -path. Done by induction. \triangle

(\Rightarrow) Let $U = \{v \in V \mid \text{visited}[v] = \text{true}\}$. We want to show that s and $V \setminus U$ are disconnected.

Notice $\text{cut}_u(V \setminus U) = \emptyset$ since otherwise, some endpoint in $V \setminus U$ would have been enqueued

$\Rightarrow U$ and $V \setminus U$ are disconnected components \Rightarrow $s \in U$ and $V \setminus U$ disconnected. \square

Correctness: **Lemma above.** Thus we see that this basic version of BFS can be used to answer:

- whether G is connected (check $\text{visited}[v] = \text{true}$ $\forall v \in V$);
- what the component containing s is;
- whether $\exists v$ -path (check $\text{visited}[v] = \text{true}$).

* **Exercise:** Find all components of G in $O(m+n)$ time.

Path Backtracking (BFS):

add an array $\text{parent}[v]$ that tracks which vertex enqueued v

$\Rightarrow \text{parent}[v]$ appears directly before v in the path found. We just need to keep finding vertex parents from the array to backtrack.

BFS Tree: Let $H \ni s$ be the component of s .

Lemma: The graph formed by edges $(v, \text{parent}[v])$ is a spanning tree of H .

Proof: Exercise (easy). Show acyclic.

BFS finds shortest paths This is a major feature of the algorithm.

interesting ↗ relatively intuitive result We can modify the BFS algorithm to compute path distances.

* **Check Bipartiteness (BFS):**

Given source $s \in V$ in a connected graph G , notice if $s \in A$ then $N(s) \subseteq B$ and $N(N(s)) \subseteq A$, etc.

\Rightarrow split V by path length parity from $s \in V$. Put even distances in A and odd in B .

$$A = \{v \in V \mid \text{dist}(v) \text{ even}\}, B = \{v \in V \mid \text{dist}(v) \text{ odd}\}$$

perform BFS starting at $s \in A$,

during BFS check if edge has endpoints in distinct positions

\Rightarrow time $O(m+n)$ (same as BFS)

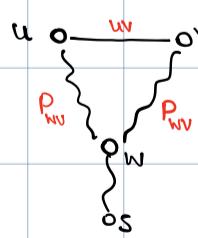
Correctness: Clearly correct when G bipartite. How about when the alg. returns false?

Proof: G bipartite $\Leftrightarrow G$ has no odd cycles.

Idea: alg. says non-bipartite \Rightarrow there is an odd cycle in G .

WLOG say $\exists u, v \in A$ with $uv \in E$ (since alg. finds one).

exists because $s \in V$ is definitely one in T .



Consider BFS tree T . let $w \in V(T)$ be the lowest common ancestor of u, v in T .

\Rightarrow paths P_{wu} and P_{vw} are even length

\Rightarrow cycle $P_{wu} uv P_{vw}$ odd length. \square

Remark: above is also an algorithmic proof of bipartiteness characterization.

also gives a linear-time alg. to find an odd cycle of an undirected graph.

* Graph DFS

In some situations, a DFS is more natural than a BFS. Consider solving a maze.

Maze \rightarrow graph where squares of the maze are vertices, and $uv \in E \Leftrightarrow$ squares u and v are one step apart.

Finding a solution is equivalent to finding some st-path.

Moving back/forth makes a BFS inefficient. We want to explore a path until a dead end
 \rightsquigarrow backtrack to last "good" vertex
 \rightsquigarrow try new path

analogy: trace chalk as you solve maze
 \rightarrow dead-end \rightarrow go back and only explore new paths \Rightarrow essentially DFS

input: undirected graph $G = (V, E)$, source $s \in V$. output: all $v \in V$ reachable from s

```
visited[v] = false  $\forall v \in V$ 
visited[s] = true, explore(s)
explore(u):
    for each neighbour v of u:
        if visited[v] = false:
            visited[v] = true, explore(v)
```

← recursive

Time: $\text{explore}(v)$ is called at most once for each $v \in V$,
with $\deg(v)$ subcalls
 \Rightarrow time $O(n + \sum_{v \in V} \deg(v)) = O(n + m)$.

We can also write a DFS non-recursively using a stack
 \Rightarrow BFS and DFS are identical, except one uses a queue and the other uses a stack
(two fundamental search methods correspond to two fundamental data structures)

All basic results (e.g. connectivity) hold for DFS.

But we will see that DFS can solve some interesting problems that BFS cannot.

* ^{important} DFS Tree:

Constructed in the same way as for BFS (using parent[v] relations).

Notice DFS trees may be different based on neighbour exploration order.

We induce tree ordering starting at root $s \in T$ and define regular terminology as in MATH 239.

A non-tree edge uv is called a **back edge** if either u is an ancestor or descendant of v in G .

Property (backedges): In an undirected graph, all non-tree edges are back edges.

Proof: SFAC $\exists u \in E$ but uv is not an ancestor-descendant pair.

WLOG say u is visited before v .

$uv \in E \Rightarrow v$ explored before u is finished $\Rightarrow u$ is an ancestor of v . \square

Starting Time and Finishing Time

We record the time when a vertex is first visited and the time when its exploration is finished.

\rightarrow information important in the design and analysis of algorithms.

We modify our algos. to include this:

```
visited[v] = false  $\forall v \in V$ 
time = 1, visited[s] = true, explore(s)
```

explore(u):

$\text{start}[u] = \text{time}$
 $\text{time} \leftarrow \text{time} + 1$

{ \dots }

Property (parentesis): Let $u, v \in V$. Then either:

- One of $[\text{start}[u], \text{finish}[u]], [\text{start}[v], \text{finish}[v]]$ is contained in the other
- $[\text{start}[u], \text{finish}[u]] \cap [\text{start}[v], \text{finish}[v]] = \emptyset$

and the former occurs precisely when u, v are ancestor-descendant pairs.

```

for  $v \in V(G)$ :
    if  $\text{visited}[v] = \text{false}$ :
         $\text{visited}[v] = \text{true}$ 
         $\text{explore}(v)$ 
     $\text{finish}(v) = \text{time}$ 
     $\text{time} \leftarrow \text{time} + 1$ 

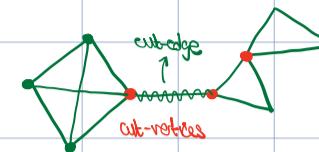
```

Remark: we can generalize DFS on non-connected graphs.
In this case, DFS would generate a forest.

Let G graph. In MATH 239, we defined when $e \in E$ was a **bridge**: when $\text{comp}(G-e) > \text{comp}(G)$.

Similarly, we say $e \in E$ is a **cut-edge** if $\text{comp}(G-e) > \text{comp}(G)$, and

$v \in V$ is a **cut-vertex** if $\text{comp}(G-v) > \text{comp}(G)$.

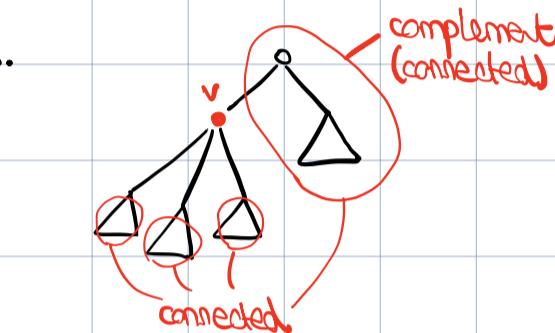


Idea: use a DFS tree to identify all cut-vertices and cut-edges.

Consider DFS tree $T \subseteq G$, and non-root $v \in V$. We want to know if v is a cut-vertex of G .

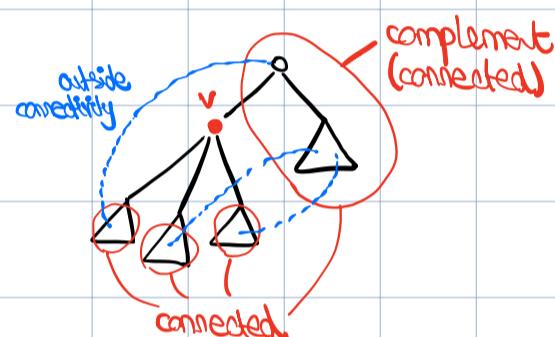
Recall **Property (back-edges)** \Rightarrow all non-tree edges are back-edges.

\Rightarrow the only way for a subtree below v to be connected outside that subtree is to have edges going to some ancestor of v .



Claim (cut-vertex): A subtree T_i below v is a connected component

in $G-v \Leftrightarrow \exists e = uv \in T_i$ with one endpoint in T_i and another endpoint in a (strict) ancestor of v .



Proof: (\Leftarrow) all non-tree edges are back-edges \Rightarrow there are no edges going to another subtree below v , nor edges going to another subtree of the root. \triangle

(\Rightarrow) If $\exists e \in E$ with one endpoint in T_i and another in a (strict) ancestor of v , then T_i and the complement are still connected $\Rightarrow T_i$ not a component $\not\in \Delta$ \square

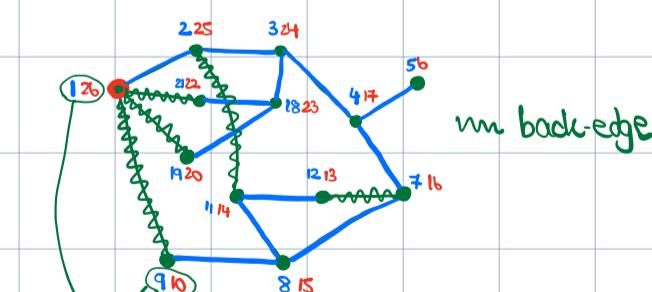
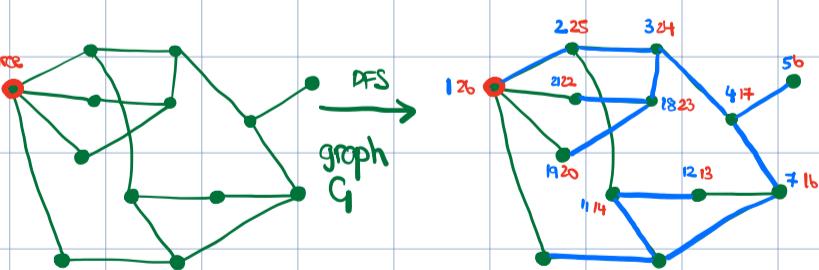
The same argument applies to all subtrees below v , so we get this characterization:

\hookrightarrow components are maximally connected.

Lemma: For a nonroot vertex v in a DFS tree,

v is a cut vertex \Leftrightarrow there is a subtree below v with no edges going to a strict ancestor of v .

Ex.



notice discovery times $910 \leq 126$
 \Rightarrow ancestor-descendant pair \Rightarrow back-edge

Let T be a DFS forest.

In general, we consider 4 edge classes on directed G :

tree edges are $e \in E(T)$

back edges are edges that connect $v \in V$ to an ancestor } where direction is implied; in undirected G ,

forward edges are edges that connect $v \in V$ to a descendant } we call both classes **back-edges**

cross edges are all other edges.

Proof: (\Rightarrow) Say v is a cut vertex. SFAC every subtree below v has an edge going to an ancestor of v .

Then every subtree of v is connected to the complement structure $\Rightarrow \text{comp}(G-v) = \text{comp}(G)$

$\Rightarrow v$ not cut vertex.

Thus there must be at least one subtree below v with no edge going to an ancestor of v . Δ

(\Leftarrow) Say some subtree T_i below v has no edges going to an ancestor of v

$\xrightarrow[\text{see above}]{\text{Claim}}$ $\Rightarrow T_i$ will be a connected component of $G-v \Rightarrow v$ is a cut vertex. $\Delta \square$

* Ex: extend this result to the root vertex:

Lemma: For root v of a DFS tree, v is a cut vertex $\Leftrightarrow v$ has at least 2 children.

Proof: (Intuition) Say v has 2 children in DFS tree. Say rooted subtrees of these children are T_1, T_2 .

Since we search depth-first, WLOG say T_1 generated first.

Notice T_2 created from $v \Rightarrow$ no exploration of T_1 brought us to T_2 (else T_2 not a subtree of v)

$\Rightarrow T_2$ unreachable from T_1 in $G-v \Rightarrow v$ is a cut vertex. \square

With these lemmas, we know how to determine if a vertex is a cut vertex by looking at a DFS tree.

Algorithm: Find all cut vertices $O(n+m)$

Idea: process the vertices of a DFS following a bottom-up ordering, keep track of how "far up" the back edges of a subtree can go (i.e. how close to the root these back edges go).

\Rightarrow use Lemma: for non-root v , we have v not a cut vertex \Leftrightarrow all subtrees below v has an edge that goes above v .

How can we keep track of "how far up" we can go?

use starting/finishing time since descendant time \leq ancestor time.

In particular, des.start \geq anc.start

Define $\text{early}[v]$ for each v in our DFS tree:

$\text{early}[v] = \min_{\text{current start}} \left\{ \text{start}[v], \min_{\substack{\text{all } w \\ \text{such that } vw \text{ is a back edge}}} \left[\text{start}[w] \mid w \text{ is a descendant of } v, \text{ or } w=v \right] \right\}$

earliest start time among all vertices connected to somewhere below v through back edge
if this w has $\text{start}[w] < \text{start}[v]$ then w ancestor of $v \Rightarrow$ back edge that goes above v in rooted subtree

Informally, $\text{early}[v]$ records how far up we can go from the subtree rooted at v .

Suffices to prove/show: ① We can compute $\text{early}[v]$ for all $v \in V$ in $O(n+m)$ time

② We can identify all cut vertices using $\text{early}[v]$ in $O(n+m)$ time.

Computing $\text{early}[v]$

① We compute $\text{early}[v]$ values from leaves to root in the DFS tree.

Prof: Base case: $v \in V$ is a leaf of the DFS tree. We want to compute $\text{early}[v]$.

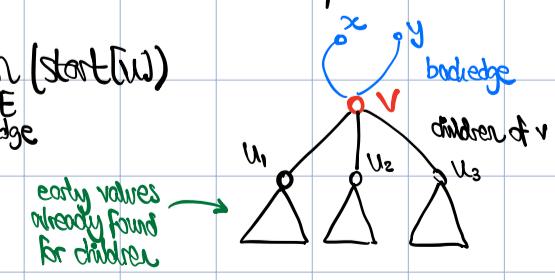
\rightsquigarrow consider all the edges incident on v (so say $uv \in E$), and take

$\text{early}[v] = \min_{uv \in E} (\text{start}[u])$ in $O(\deg(v))$ time (base case)

Inductively suppose, for $v \in V$, all $\text{early}[u]$ values for children of v has been computed.

$\Rightarrow \text{early}[v] = \min_{u \in \text{children of } v} \left(\min_{u \in \text{children of } v} (\text{early}[u]) \text{ and } \min_{u \in \text{children of } v} (\text{start}[u]) \right)$

$= \min_{\substack{u \in \text{children of } v \\ \subseteq \text{edges of } v}} \left(\min_{u \in \text{children of } v} (\text{early}[u]), \min_{u \in \text{children of } v} (\text{start}[u]) \right)$ $O(\deg(v))$ time



early values of children correct \Rightarrow early[v] correct * can be proven
inductively
bottom-up ordering \Rightarrow every $v \in V$ only computed once

so total time complexity is $\underbrace{O(n + \sum_{\text{ver}} \deg(v))}_{2|E|} = O(n+m)$. \square

Finding cuts

(2) To check if non-root $v \in V$ is a cut vertex, we need to check whether $\text{early}[u_i] < \text{start}[v]$ for all children u_i of v .

Proof: $\text{early}[u_i] < \text{start}[v]$ for all u_i child $\Rightarrow v$ not a cut vertex
i.e. we could have gone a different way from above v to each subtree

since all subtrees below v have a back edge going above v .

otherwise if $\exists u_i \text{ child with } \text{early}[u_i] \geq \text{start}[v]$ then u_i -rooted subtree will be a connected component of $G-v \Rightarrow v$ cut vertex.

Handle root vertex separately. Convince yourself this is $O(n+m)$. \square

* Exercise: Extend algorithm to find cut edges of a graph G .

Remarks: We looked at (undirected) graphs G and what BFS/DFS can do.

A highlight is very clever algorithm to detect strongly connected components in linear time.

Now we look at directed graph generalizations and algorithms

Directed Graph Algorithms

A directed graph is a graph G where $e \in E$ has an order property (in particular, $uv \neq vu$).

If $uv \in E$, we call u the tail and v the head of the edge. $u \rightarrow v$

We define $\text{indeg}(v)$ and $\text{outdeg}(v)$ accordingly.

\rightsquigarrow useful in modeling asymmetric relations. We are interested in connectedness properties of directed G .

G directed graph, $s, t \in V$. We say t is reachable from s if \exists directed st-path in G . $s \rightarrow \dots \rightarrow t$

We say G is strongly connected if for all pairs $u, v \in V$, u reachable from v and v reachable from u .



not strongly connected

A subset $S \subseteq V$ is strongly connected similarly.

$S \subseteq V$ is a strongly connected component if S is a maximally strongly connected subset

(i.e. $S \subseteq V$ strongly connected and $S+V$ not strongly connected for any $v \in V \setminus S$).

G directed is a directed acyclic graph if G has no directed cycles.

Remark: G undirected and G acyclic $\Rightarrow |E|$ is bounded.

The same does not hold for directed graphs because of edge direction property.



acyclic directed graph

We want to design algorithms to answer basic questions about directed graphs:

1 Is G strongly connected?

2 Is G directed acyclic?

3 Find all strongly connected components of G .

Turns out there are $O(n+m)$ algorithms for these, but they are more complicated than those for undirected graphs.

Graph Representations

Adjacency Matrix Extends naturally to directed graphs: $A_{ij} = \begin{cases} 1 & \text{if } ij \in E \\ 0 & \text{otherwise} \end{cases}$

note $ij \neq ji$ (directed)

* finish directed graphs

Single-Source Shortest Paths

We study Dijkstra's algorithm for shortest paths from source vertex to all other vertices.

Shortest Paths Input: A directed graph $G = (V, E)$ with a weight $w_e \in \mathbb{R}_{\geq 0}$ for each $e \in E$, and two arbitrary vertices $s, t \in V$

source target

Output: A shortest st-path P , where $\text{length}(P) = \sum_{e \in P} w_e$.

It will be more convenient to solve a more general problem.

shortest single-source path algo. Input: Directed $G = (V, E)$ with $w_e \in \mathbb{R}_{\geq 0}$ for $e \in E$, and some source $s \in V$.

Output: A shortest sv-path, for all $v \in V$.

This may seem **harder**, since each path could have $\leq n$ edges, and so output size could be $\leq n^2$.

→ turns out that there is a succinct representation of paths, and the single-source shortest path can be solved in the same complexity as shortest st-path problem.

Dijkstra's Algorithm

We have seen that BFS can be used to solve the single-source shortest path problem, where every edge has the same weight. so we essentially minimize the #edges in the path

We can reduce the non-constant $w_e \in \mathbb{R}_{\geq 0}$ case to the same-weight special case.

replace each edge of weight $w_e \in \mathbb{Z}_{\geq 0}$ by a path of w_e edges, forming G'

then $\exists s$ -path of length k in $G \Leftrightarrow \exists s$ -path of length k in G'

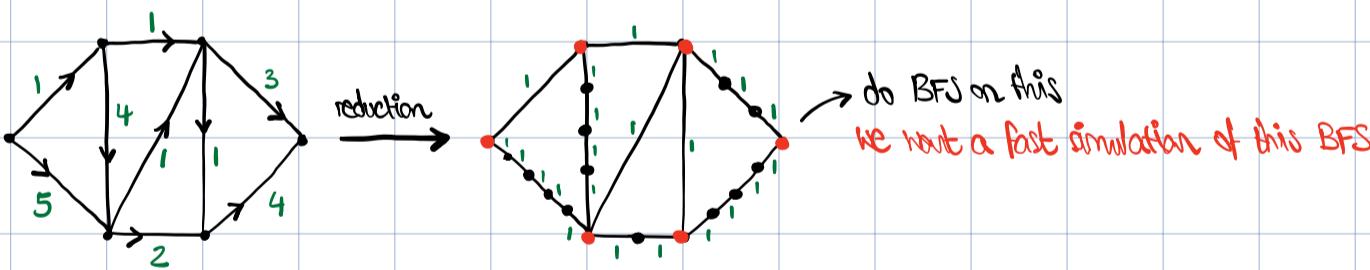
\Rightarrow linear time algorithm for $G' \Leftrightarrow$ linear time algorithm for G

Problem: reduction might be inefficient since $|V'| \geq |V|$.

In particular, $|V'| = n + \sum_{e \in E} (w_e - 1) \Rightarrow$ large w_e problematic.

Idea: To design an efficient algorithm through this reduction, instead of constructing G' explicitly and running BFS,

we simulate BFS on G' efficiently.



Physical Process Analogy

Construct graph G' with $w_e = 1, \forall e \in E'$.

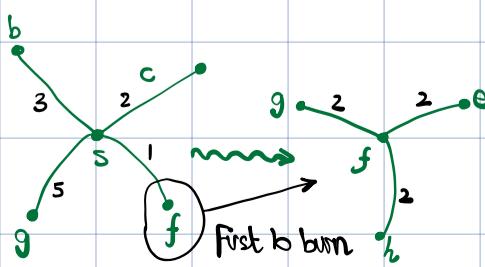
Choose source $s \in V$ and start a fire that spreads one edge per unit time.

\Rightarrow shortest path distance to $t \in V$ = amount of time until fire reaches t .

To simulate the process efficiently, we just need to keep identifying what the next vertex to be burned is, and when.

Initially, source $s \in V$ is burned.

We know b, c, g, f will be burned in time at most 3, 2, 5, 1.



f burns first. Then f burned at time 1 \Rightarrow g, h, e will burn in at most 2 more time.

\Rightarrow we update the upper bound time for g to be burned to 3.

This is Dijkstra's algorithm.

\rightarrow We continue by finding the next vertex burned, and update upper bounds on its neighbours.

Dijkstra's Algorithm (BFS Simulation)

dijkstra($G, s \in V$):

initialize {
 $\text{dist}(v) = \infty$ for every $v \in V$ initial upper bound ∞ for all vertices
 $\text{dist}(s) = 0$

$Q = \text{priority-queue}(V)$ put all vertices in a priority queue with priority value $\text{dist}(v)$
 notice $s \in Q$ as well.

while $Q \neq \emptyset$:

next to burn $\rightarrow u = \text{delete-min}(Q)$ deque vertex with shortest dist() value (next to be burned)

for each neighbour v of u :

shorter path update {
 if $\text{dist}(u) + w_{uv} < \text{dist}(v)$: shorter way to get to v through u
 $\text{dist}(v) = \text{dist}(u) + w_{uv}$
 decrease-key(Q, v) \rightarrow the distance always only decreases here, so we update the
 parent(v) = u priority of the element using decrease-key in Q

update the parent of v to be u } this is where P is constructed
on the shortest path found }

Proposition (Dijkstra): The algorithm finds shortest paths in $O(n(m+n\log n))$ or $O(m+n\log n)$, where $m = \sum_{v \in V} \text{outdeg}(v)$.

Proof (Correctness): The algorithm is an efficient implementation

of BFS on pseudo- G' . One method is to prove the reduction holds.

We proceed with a formal proof on an equivalent formulation.

1 The following formulation of Dijkstra's is equivalent to the one above:

dijkstra($G, s \in V$):

$\text{dist}(v) = \infty \quad \forall v \in V, \text{dist}(s) = 0$

$R = \emptyset$

while $R \neq V$:

pick $u \in V \setminus R$ that minimizes $\text{dist}(u)$

$R \leftarrow R \cup \{u\}$

for each $v \in E$:

$\text{dist}(v) = \min(\text{dist}(v), \text{dist}(u) + w_{uv})$

*Equivalence: exercise (easy)

equivalently, $u \in R$ is the vertex such that

$\text{dist}(u) = \min_{v \in R, y \in E} \{\text{dist}(y) + w_{yu}\}$

all y that have been burned

v that are not yet burned, neighbouring y

\Rightarrow choose next to burn

2 The algorithm above is correct.

We proceed by induction. Invariant: for any $v \in R$, $\text{dist}(v)$ is the shortest path distance from s to v .

Base Case: True when $R = \{s\}$.

Now assume that the invariant holds for current R . We want to show that the invariant holds after a new vertex u is added to R (in the $R \leftarrow R \cup \{u\}$ step).

We must show that $\text{dist}(u) = \min_{v \in R, y \in E} \{\text{dist}(y) + w_{yu}\}$ is the shortest path distance from s to u .

Let $y \in R$ be the vertex in a minimizer,

*What's v here?

i.e. $\text{dist}(u) = \text{dist}(y) + w_{yu} \leq \text{dist}(a) + w_{ab} \quad \forall a \in R \text{ and } b \notin R$.

(\Leftarrow) Since invariant holds for R , we have $y \in R \Rightarrow \text{dist}(y)$ shortest path from s to y

$\Rightarrow \text{dist}(u) \leq \text{dist}(y) + w_{yu}$ using sy-path and following yu-edge. Δ^{\leq}

(\geq) Now we show $\text{dist}(u) \geq \text{dist}(y) + w_{yu}$ for the minimizer $y \in R$.

Let P be any subpath. Since $s \in R$ and $u \in R$, there must be some edge $xz \in P$ with

$x \in R$ and $z \in R$

$\Rightarrow \text{length}(P) \geq \text{dist}(x) + w_{xz}$ lower bound (crucially uses fact that $w_e > 0 \forall e \in E$). \star later we explore this more difficult case using DP

But $y \in R$ minimizer $\Rightarrow \text{dist}(y) + w_{yu} \leq \text{dist}(x) + w_{xz}$ since

$$\text{dist}(u) = \text{dist}(y) + w_{yu} \leq \text{dist}(x) + w_{xb}$$

$\Rightarrow \text{length}(P) \geq \text{dist}(x) + w_{xz} \geq \text{dist}(y) + w_{yu}$. shortest path choice

Equality holds for any subpath $\Rightarrow \text{length}(P) = \text{dist}(u) \geq \text{dist}(y) + w_{yu}$. Δ^{\geq}

$\text{dist}(u)$ on shortest path

So $\text{dist}(u) \leq \text{dist}(y) + w_{yu}$ and $\text{dist}(u) \geq \text{dist}(y) + w_{yu} \Rightarrow \text{dist}(u) = \text{dist}(y) + w_{yu}$ \square Correctness

Proof (Complexity): Every $v \in V$ is enqueued exactly once (at the start) and dequeued exactly once (when it burns).

When dequeuing, we check every outbound edge once and may use $\text{dist}(u) + w_{uv}$ to decrease-key.

dijkstra($G, s \in V$):

$\text{dist}(v) = \infty$ for every $v \in V$ $O(n)$

$\text{dist}(s) = 0$ $O(1)$

$Q = \text{priority-queue}(V) \leftarrow O(n \log n)$ inserting n items into Q

suppose binary min-heap

while $Q \neq \emptyset$:

$u = \text{delete-min}(Q)$

Note: Fibonacci heaps are out of scope.

for each neighbour v of u :

if $\text{dist}(u) + w_{uv} < \text{dist}(v)$:

$\text{dist}(v) = \text{dist}(u) + w_{uv}$

$\text{decrease-key}(Q, v)$

$\text{parent}(v) = u$

$O(\deg(u))$ iterations

$O(1 \log n)$ for bubble-up

$$\Rightarrow T(n) = n \log n + \sum_{v=1}^n \left(\sum_{\substack{w \in E \\ \text{deg}(v)}} \log n \right) = n \log n + \sum_{v=1}^n \deg(v) \log n = (n + \underbrace{\sum_{v \in V} \deg(v)}_m) \log n.$$

\square Complexity

Allowing negative edge lengths makes the SS shortest-path problem considerably harder.

Before approaching this problem and introducing the **Bellman-Ford algorithm**, we introduce dynamic programming.

Intuition: although Dijkstra may not compute all distances correctly in one pass, it will compute distances to some of the vertices correctly (e.g. the first vertex on a shortest path).

Dynamic Programming (1)

On a high level, we can solve a problem by **dynamic programming** if there is a recurrence relation with only a small number of subproblems (i.e. polynomially many).

This is a **general and powerful technique** that is simple to use once learnt well.

A common toy problem is computing the Fibonacci sequence.

The function $F(n) = F(n-1) + F(n-2)$ is defined recursively. Using recursion gives a huge recursion tree
⇒ runtime $\mathcal{O}(1.618^n)$ (MATH 239).

Notice many subproblems are computed redundantly (we only have n subproblems) ↼ why waste time?

Top-Down Memorization

Store visited/computed values in an array $\text{visited}[i]$ ↗ or hashmap so we can avoid re-computation.
⇒ each subproblem computed once, and when it is computed, we do an $\mathcal{O}(1)$ lookup $\Rightarrow \mathcal{O}(n)$ computations.

Bottom-Up Computation

Straightforward bottom-up iterative approach $\Rightarrow \mathcal{O}(n)$ additions.

This is the framework of **DP**: store intermediate values to avoid re-computation.

From the top-down approach, it should be clear that if we write a recursion with only polynomially many subproblems with some additional polynomial-time processing, the problem can be solved in polynomial time.

⇒ designing efficient DP algorithms boils down to designing a nice recursion (like D&C algorithms).

In principle, we just need to find some topological ordering of subproblems.

A. Weighted Interval Scheduling :

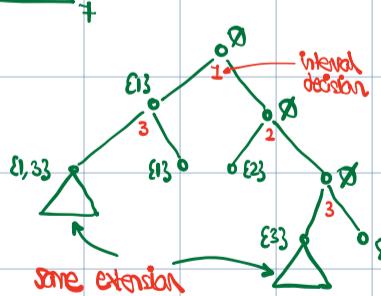
Input: n intervals $[s_i, f_i], \dots, [s_n, f_n]$ with weights $w_1, \dots, w_n \in \mathbb{R}$
Output: a subset n of disjoint intervals that maximizes $\sum_{i \in S} w_i$

We can think of the intervals as requests to book a room for time $[s_i, f_i]$ and cost w_i

⇒ we want to maximize our earnings by choosing good intervals

Naïve Exhaustive search.

- 1 Decide whether to choose interval 1.
choose 1 ⇒ reject 2 ⇒ decide on interval 3
not choose 1 ⇒ decide on interval 2
- 2 Proceed recursively.



⇒ lots of redundancy in tree

Notice: To determine how to extend partial solutions, what really matters is the last chosen interval of the current partial solution.

Nothing on the left matters, since they won't

intersect with anything on the right.

here, the boundary is the last interval

⇒ we just need to keep track of the "boundary" of the solution ⇔ there should be a recursion with only one parameter

boundary

Better Recurrence → use a good ordering of intervals and pre-compute some useful information

1 Sort intervals by starting time so $s_1 \leq s_2 \leq \dots \leq s_n$.

2 For i^{th} interval, use $\text{next}[i]$ to determine smallest j with $j > i$ and $s_j < s_i$.
If no such j exists, set $\text{next}[i] = n+1$ (end)

strict imp. because we want disjoint intervals
 i ends before j starts

⇒ for interval i , we have that intervals $\{i, \dots, \text{next}[i]-1\}$ overlap with i , and $\{\text{next}[i], \dots\}$ are disjoint from i .

Now we are ready to write a recurrence with one parameter, resulting in n subproblems.

Define $\text{opt}(i) = \max$ income that we can earn using intervals in $\{i, \dots, n\}$ only. Goal: compute $\text{opt}(1)$.

To compute $\text{opt}(1)$ there are only 2 possibilities:

1) Solutions that choose interval 1.

⇒ we cannot choose intervals in $\{2, \dots, \text{next}[1]-1\}$ since these overlap with interval 1

⇒ the optimal value for solutions choosing interval 1 is $w_1 + \text{opt}(\text{next}[1])$

2) Solutions that exclude interval 1.

⇒ $\text{opt}(2)$ is by definition the optimal value

⇒ $\text{opt}(1) = \max(\text{opt}(2), w_1 + \text{opt}(\text{next}[1]))$ and generally, $\text{opt}(i) = \max(\text{opt}(i+1), w_i + \text{opt}(\text{next}[i]))$ defines a recurrence.

Top-down weighted interval scheduling

1 Sort intervals by non-increasing starting time so $s_1 \leq \dots \leq s_n$

2 Compute $\text{next}[i]$ for $1 \leq i \leq n$, set $\text{visited}[i] = \text{false}$

3 Return $\text{opt}(1) \leftarrow \text{O}(n)$ we have n subproblems

} preprocess $\text{O}(n \log n)$ * why?

clearly, sorting is $\text{O}(n \log n)$.
remains to show that $\text{next}[]$ array can be computed in $\text{O}(n \log n)$.
consider using a binary search approach:

$\text{next}[i, A, n]:$ A sorted array
 $\text{mid} = n/2$
 $A_L, A_R = A$ (halved at middle)
if $A_L[i] \leq A[i]$:
 next must be in A_R
else:
 next must be in A_L

$\text{O}(n \log n)$ algorithm
⇒ $\text{O}(n \log n)$ when called over all elems

$\text{opt}(i):$ subroutine definition

if $i \geq n+1$: } base case
 return 0

if $\text{visited}[i]$ true: } $\text{opt}(i)$ already computed
 return $\text{answers}[i]$

$\text{answers}[i] = \max(\text{opt}(i+1), w_i + \text{opt}(\text{next}[i]))$ } recursive

$\text{visited}[i] = \text{true}$

return $\text{answers}[i]$

We can also write an equivalent bottom-up variant:

Bottom-Up Weighted Interval Scheduling

some pre-processing in $\text{O}(n \log n)$

$\text{opt}(n+1) = 0$

for ($i=n$; $i \geq 1$; $i \leftarrow$):

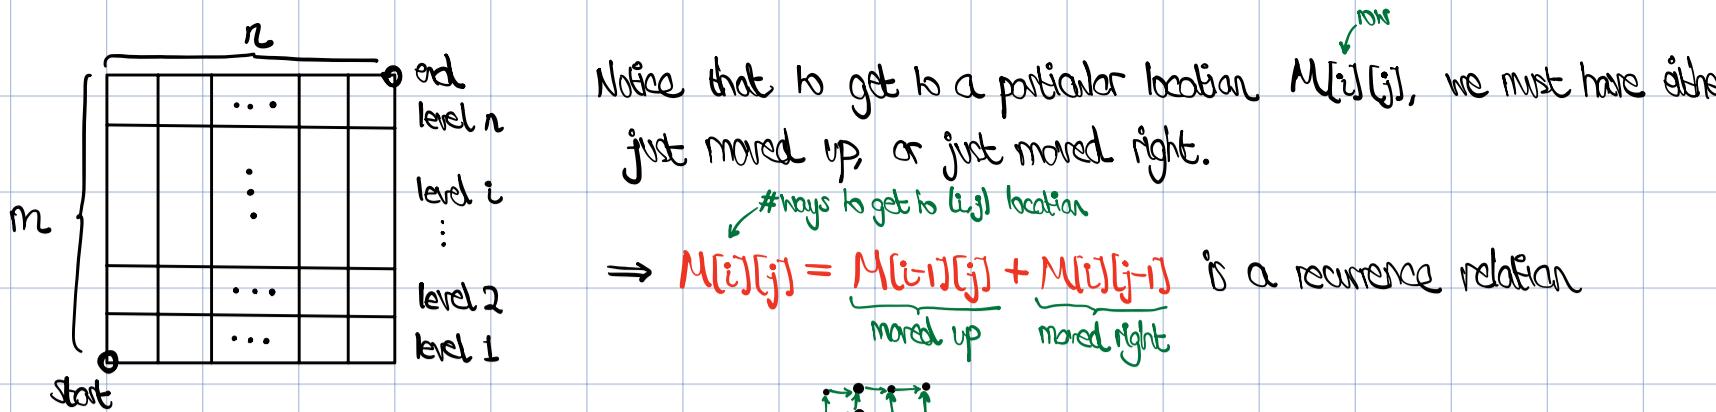
$\text{opt}(i) = \max(\text{opt}(i+1), w_i + \text{opt}(\text{next}[i]))$

} potentially more natural solution working backwards to find an answer

both of these would already be pre-computed

We can see why dynamic programming is general and powerful, because it is very systematic and yields very competitive algorithms.

B Grid Path Counting say we have an $m \times n$ grid and start at the bottom left. We can move up or right each step.
how many ways of getting from bottom left to top right?



Therefore compute dependencies look something like
Fortunately, this is very easy.

\Rightarrow we have reduced to a recurrence with n^2 subproblems.

Remark: We can likely derive a combinatorial/analytic solution, but we are not interested in that

C Subset-Sum Problem

Input: $a_1, \dots, a_n \in \mathbb{N}_0$ and $k \in \mathbb{Z}$

Output: some collection $S \subseteq \{1, \dots, n\}$ such that $\sum_{i \in S} a_i = k$, or report that no such S exists

For each a_i , we can choose to either include it or exclude it. Say we start with a_1 .

1) Include a_1 . Then we must choose some subset $S' \subseteq \{2, \dots, n\}$ such that $a_1 + \sum_{i \in S'} a_i = k$ so $\sum_{i \in S} a_i = k - a_1$.

2) Exclude a_1 . Then we need $S' \subseteq \{2, \dots, n\}$ with $\sum_{i \in S'} a_i = k$

Naive: Naively implement the above logic $\Rightarrow O(2^n)$ time since every subset will be enumerated.

We want to **reduce the search space** for the problem.

Observation: Say we have chosen two subsets $S_1 \subseteq \{1, \dots, l\}$ and $S_2 \subseteq \{l+1, \dots, n\}$ so far, and say that

$\sum_{i \in S_1} a_i = \sum_{i \in S_2} a_i$. Notice the remaining $S' \subseteq \{l+1, \dots, n\}$ is agnostic of the choice of S_1 or S_2 thus far.

\Rightarrow as long as subsets collected thus far have the same sum, we need not track them.

\Rightarrow we can just keep track of the partial sum in the recursion.

{ reduce

consider the subproblem $\text{subsum}[i, L]$ for $i \in \{1, \dots, n\}$ and $L \leq k$, where $\text{subsum}[i, L] = \text{true} \Leftrightarrow \exists S' \subseteq \{i, \dots, n\}$ with sum L .

\Rightarrow we would like to solve $\text{subsum}[1, k]$

We have a recurrence $\text{subsum}[i, L] = (\underbrace{\text{subsum}[i+1, L-a_i]}_{\text{include element } a_i} \text{ OR } \underbrace{\text{subsum}[i+1, L]}_{\text{exclude element } a_i})$

Top-down subset-sum

$\text{subsum}(i, L)$:

if $L=0$: return true \leftarrow choose empty set \emptyset

if $(i > n \text{ or } L < 0)$: return false \leftarrow ran out of numbers or partial sum too large

return $(\text{subsum}(i+1, L-a_i) \text{ or } \text{subsum}(i+1, L))$

\Rightarrow there are $O(nk)$ subproblems (n choices for i , k choices for L)
 \Rightarrow time complexity $O(nk)$.

Remark: $O(nk)$ is pseudo-polynomial since the runtime looks polynomial but depends on k .

If $k \in \Omega(2^n)$ this would be exponential, and would be even slower than exhaustive search.

This is likely unavoidable. We will soon see that the subset-sum problem is NP-complete.

* Exercise: bottom-up algorithm implementation (use matrix $M \in \mathbb{R}^{n \times k}$ for memorization).

optimize space and use a matrix $M \in \mathbb{R}^{2 \times k}$.

Notice that the bottom-up version has runtime $\Theta(nk)$ but top-down is $O(nk)$.

some inputs may return faster

Dynamic Programming and Graph Search

We can consider representing the search space of a DP problem as a graph, where each vertex is a subproblem, and edges are added according to the DP recurrence relation.

Then, the problem is equivalent to finding some directed path from a starting (desired) state to a "true base state".

Top-down memorization is basically doing a DFS on this subproblem graph.

Input: n items of weight w_i and value v_i , and some $W \in \mathbb{N}_0$.

D_{in} Knapsack Output: a subset $S \subseteq \{1, \dots, n\}$ with $\sum_{i \in S} w_i \leq W$ that maximizes $\sum_{i \in S} v_i$.

Think of a robber who can carry weight at most W , and wants to maximize the value of stolen goods.

Define subproblem $\text{knapsack}(i, W, V) = \text{true} \Leftrightarrow \exists S \subseteq \{1, \dots, n\}$ with weight $\leq W$ and total value V .

Two possibilities for object i : either include or exclude.

- 1) Include $\Rightarrow \text{knapsack}(i+1, W-w_i, V-v_i)$; 2) Exclude $\Rightarrow \text{knapsack}(i+1, W, V)$.
- weight of removing value of removing

Therefore $\text{knapsack}(i, W, V) = (\text{knapsack}(i+1, W-w_i, V-v_i) \text{ or } \text{knapsack}(i+1, W, V))$ recurrence relation.

Notice we can use a 3D tensor to encode this information.

Parameter / Dimensionality Reduction

Notice that, for the solution we intend to find, some subproblems dominate others.

In particular, if $\text{knapsack}(i, W, V')$ is true for some $V' > V$, we can ignore subproblem $\text{knapsack}(i, W, V)$.

\Rightarrow define $\text{knapsack}(i, W) = \max$ value achievable with $S \subseteq \{1, \dots, n\}$ of weight $\leq W$.

More precisely, $\text{knapsack}(i, W) = \max_{S \subseteq \{1, \dots, n\}} \left\{ \sum_{j \in S} v_j \mid \sum_{j \in S} w_j \leq W \right\}$.

↓ recurrence

$\text{knapsack}(i, W) = \max \left\{ \underbrace{v_i + \text{knapsack}(i+1, W-w_i)}_{\text{include item } i}, \underbrace{\text{knapsack}(i+1, W)}_{\text{exclude item } i} \right\}$. Base cases: $W < 0 \Rightarrow \text{return } -\infty$
 $i > n \Rightarrow \text{return } 0$

* Exercise: Implement pseudocode. $\Rightarrow O(nW)$ time (subproblem space size).

Remark: Knapsack is pseudo-polynomial and is actually NP-complete, proven in last chapter.

there are n stairs in a staircase. each step we can climb either 1 or 2 steps.

DP1 climbing-stairs how many distinct ways of climbing?

Notice $\text{climb}(i) = \underbrace{\text{climb}(i-1)}_{1 \text{ stair}} + \underbrace{\text{climb}(i-2)}_{2 \text{ stairs}}$ defines a recurrence relation with $\text{climb}(0) = 1$
 $\text{climb}(1) = 1$
 \Rightarrow reduce to Fibonacci.

3/17/2021
-Charlie Lu

DP2 house-robbing: some array A where $A[i] = \text{money in house } i$. Houses are arranged contiguously in a straight line.
med can only steal from non-adjacent houses; what's the most we can steal?

Notice we can define a recursive subproblem decomposition:

$$\Rightarrow \text{value}[i] = \max \left(\underbrace{\text{get-value}(i-1)}_{\text{skip } i}, \underbrace{A[i] + \text{get-value}(i-2)}_{\text{skip } i+1} \right)$$

2 choices: steal from $i+1$ and skip i , or steal from $[1..i]$



and skip $k+1$

giving us $O(n)$ time and $O(n)$ space with memoization

DP3 coin-change: given a target $t \in \mathbb{N}$ and a list of coin denominations A , find smallest number of coins needed to make t .

Notice greedy won't work here, since consider $t=105$ and $A=[1, 7, 98, 100]$.

Given some value v , we want to know what coin a_i to use to minimize total #coins.

Say $\text{coins}[v]$ represents the smallest #coins needed for value v .

then we have a recursive definition $\text{coins}[v] = \min_i (1 + \text{coins}[v - a_i])$ $\xrightarrow{\text{O}(n)}$ \Rightarrow we can design an $O(n)$ algorithm for coin-change.

```

need = dict()
need[0] = 0
def get_need(amt):
    if amt < 0: return ∞
    if amt in need: return need[amt]
    need[amt] = 1 + min{get_need(amt - a_i) | a_i ≤ amt}
return get_need(t)

```

Dynamic Programming (2)

In this section, we will use dynamic programming for basic sequence and string problems.

E Longest Increasing Subsequence

Given numbers a_1, \dots, a_n , a subsequence is a subset of the form a_{i_1}, \dots, a_{i_k} where $i_1 < \dots < i_k$.

The subsequence is increasing if $a_{i_1} < \dots < a_{i_k}$.

Subproblem: let $L(i) =$ length of a longest increasing subsequence starting at a_i formed using $\{a_i, \dots, a_n\}$

\Rightarrow there are n subproblems.

once we compute $L(1), \dots, L(n)$, our final answer is $\max(L(i))$.

Recurrence: say we are at a_i . Notice that we can prepend a_i to any subsequence $L(j)$ for $j > i$ and $a_j > a_i$.

(so $L(i)$ forms some subsequence $\underbrace{a_i, a_j, \dots}_{\text{found}}$ giving us $L(i) = \max_{\substack{j > i \\ a_j > a_i}} (1 + L(j)) = 1 + \max_{j > i} \{L(j) | a_j > a_i\}$).

* Implement

$\Rightarrow O(n^2)$ time.

Faster Algorithm: Notice some subproblems are "dominated" by others.

For each length k , we store the "best" position to start an increasing subsequence of length k .

\Rightarrow these "best" positions satisfy a monotone property.

\Rightarrow we can binary search and update these values in $O(\log n)$ time when considering a new a_i .

Idea: say we have computed $L(i+1), \dots, L(n)$ and we want to compute $L(i)$.

for some length k , consider some indices i_1, \dots, i_k such that $L(i_1) = \dots = L(i_k) = k$ and $i < i_1 < \dots < i_k$.

what's the best subproblem to keep for computations $L(i), L(i-1), \dots, L(1)$?

notice since we are extending using some elements in $\{i_1, \dots, i_k\}$, the start indices i_1, \dots, i_k are unimportant.

The starting values are what matter.

if $L(i_1) = L(i_2)$ and $a_{i_1} > a_{i_2}$, then subproblem $L(i_1)$ dominates subproblem $L(i_2)$, since

among increasing subseq s_i, s_{i_2} of length k
one with largest start value is easiest to be extended

} any increasing $s_i \subseteq \{a_1, \dots, a_n\}$ that can be extended by some s_{i_2} where s_{i_2} starts at a_{i_2} length k
can also be extended by s_{i_2} starting at a_{i_1} length k

$a_i \underbrace{a_{i_1} \dots}_k$ and $a_i \underbrace{a_{i_2} \dots}_k \Rightarrow$ we choose to extend by $a_{i_1} > a_{i_2}$

Define $\text{pos}[k] = \underset{j > i}{\operatorname{argmax}} \{a_j \mid L(j) = k\}$ where $L(i)$ is the current subproblem to be computed.

j such that a_j maximized with same length-k subsequence
best position $j > i$ to start a length-k increasing subsequence

Let $m_i = \max_{j > i} \{L(j)\}$ be the length of the longest future subsequence.

\Rightarrow when computing $L(i)$, we just need to consider $L(\text{pos}(i)), \dots, L(\text{pos}(m_i))$

for instance, say our sequence is $\underbrace{2, 7, 6, 1, 4, 8, 5, 3}_1, \underbrace{2, 7, 6, 1, 4, 8, 5, 3}_2, \underbrace{2, 7, 6, 1, 4, 8, 5, 3}_3$. Say we are computing $L(2)$. We have:

$L(3) = L(5) = 2 \rightarrow \text{pos}[2] = 3 \quad (a_3 = 6)$
 $L(4) = 3 \rightarrow \text{pos}[3] = 4 \quad (a_4 = 1)$
 $L(6) = L(7) = L(8) = 1 \rightarrow \text{pos}[1] = 6 \quad (a_6 = 8)$

} extract best future index to start a subseq. of that length
we only keep these subproblems

Claim (Monotone Property): Once we only keep the above subproblems, we have

$A[\text{pos}[1]] > A[\text{pos}[2]] > \dots > A[\text{pos}[m_i]]$ at subproblem $L(i)$

longest element to start length 1 longest element to start length 2

\rightarrow makes intuitive sense?

Proof: SFAC $\exists j$ with $A[\text{pos}[j]] \geq A[\text{pos}[j-1]]$. Let $a_{p_1} < \dots < a_{p_j}$ be an optimal increasing subsequence of length j

where $p_i = \text{pos}[j]$. Notice $a_{p_2} < \dots < a_{p_j}$ is an increasing subsequence of length $j-1$ and $a_{p_2} > a_{p_1}$.

Thus $A[\text{pos}[j-1]] \geq a_{p_2} > a_{p_1} = A[\text{pos}[j]] \Rightarrow A[\text{pos}[j]] < A[\text{pos}[j-1]]$. \square

Now we try to formulate a way to update the best subproblems using **binary search** when on a_i .

Suppose we have found best subproblems $\text{pos}[m_i], \dots, \text{pos}[1]$ in past computations. We have 3 cases:

1) $a_i < a_{\text{pos}[m_i]}$ best case. We extend the longest subsequence by 1, so $m_{i-1} = m_i + 1$ and $\text{pos}[m_{i-1}] = i$.

binary search to find such j \rightarrow 2) $a_{\text{pos}[j]} \leq a_i < a_{\text{pos}[j-1]}$ \rightarrow cannot use a_i to extend length-j subsequence, but can extend $j-1$.

since $a_i \geq a_{\text{pos}[j]}$ the new index i for a length-j subsequence is better than previous $\text{pos}[j]$

$\Rightarrow \text{pos}[j] \leftarrow i$.

3) $a_{\text{pos}[i]} \leq a_i \Rightarrow a_i$ is larger than all other subsequence starting values, so we can extend no subseq.

but can update $\text{pos}[1] \leftarrow i$.

$m=1, \text{pos}[1]=n$

for ($i=n-1; i \geq 1; i--$):

if $a_i < a_{\text{pos}[m]}$: set $m \leftarrow m+1$ and $\text{pos}[m] \leftarrow i$

else we binary search to find smallest j with $a_{\text{pos}[j]} \leq a_i < a_{\text{pos}[j-1]}$
and set $\text{pos}[j] \leftarrow i$

return m

the final algo. is very simple
but may be hard to invent.

\Rightarrow time $O(n \log n)$

F Longest Common Subsequence

We have two strings a_1, \dots, a_n and b_1, \dots, b_m where each a_i, b_j are symbols.

We want to find the largest k such that there are $i_1 < \dots < i_k$ and $j_1 < \dots < j_k$ such that $a_{i_1}, \dots, a_{i_k} = b_{j_1}, \dots, b_{j_k}$.

The E Longest Increasing Subsequence problem is a special case of this, where $b = \text{sorted}(a)$.

recurrence: let $C(i, j)$ be the length of a longest common subsequence of a_1, \dots, a_n and b_1, \dots, b_m . Want: $C(1, 1)$.

To compute $C(i, j)$ there are three cases:

1) $a_i = b_j$ \Rightarrow put a_i and b_j at the start of any existing common subsequence \Rightarrow remains to find $C(i+1, j+1)$
use both
 \Rightarrow let $S_1 = 1 + C(i+1, j+1)$. otherwise $S_1 = 0$.

2) not use a_i \Rightarrow find a longest common subsequence for a_{i+1}, \dots, a_n and b_1, \dots, b_m .
 \Rightarrow let $S_2 = C(i+1, j)$.

3) not use b_j \Rightarrow find a longest common subsequence for a_1, \dots, a_n and b_1, \dots, b_{j-1}, b_m .
 \Rightarrow let $S_3 = C(i, j+1)$.

Then we take $C(i, j) = \max\{S_1, S_2, S_3\}$ since these encompass all possibilities.

Time: there are mn subproblems, each doing $O(1)$ lookups (if we memoize) $\Rightarrow O(mn)$ time.

G Edit Distance

Two strings a_1, \dots, a_n and b_1, \dots, b_m .

We want to find smallest k such that k insert/delete/change operations can change a_1, \dots, a_n to b_1, \dots, b_m .

recurrence: similar to E LCS. let $D(i, j)$ be the edit distance of a_1, \dots, a_n and b_1, \dots, b_m . Want: $D(1, 1)$.

To compute $D(i, j)$ there are four subcases:

1) insert b_j to the current string (e.g. $\dots | abc \rightarrow \dots - a | b c$)

Then we match one more symbol of the target string and move on: if $j \leq m$ then $\text{sol}_1 = 1 + D(i, j+1)$ else $\text{sol}_1 = \infty$.

2) delete a_i from the current string when $i \leq n$ (e.g. $\dots | abc \rightarrow a \dots - b c$)

Then move forward one symbol in the current string: if $i \leq n$ then $\text{sol}_2 = 1 + D(i+1, j)$ else $\text{sol}_2 = \infty$.

3) change a_i to b_j when $i \leq n$ and $j \leq m$. Thus if $i \leq n$ and $j \leq m$ then $\text{sol}_3 = 1 + D(i+1, j+1)$ else $\text{sol}_3 = \infty$.

4) match a_i and b_j if $i \leq n$ and $j \leq m$ and $a_i = b_j$, so if $i \leq n$ and $j \leq m$ and $a_i = b_j$ then $\text{sol}_4 = D(i+1, j+1)$ else $\text{sol}_4 = \infty$

Set $D(i, j) = \min\{\text{sol}_1, \text{sol}_2, \text{sol}_3, \text{sol}_4\}$.

There are mn subproblems, each requiring a constant number of operations \Rightarrow time $O(mn)$.

Dynamic Programming on Graphs

Recall:

Allowing negative edge lengths makes the SS shortest-path problem considerably harder.

Before approaching this problem and introducing the **Bellman-Ford algorithm**, we introduce dynamic programming.

Intuition: although Dijkstra may not compute all distances correctly in one pass, it will compute distances to some of the vertices correctly (e.g. the first vertex on a shortest path).

Dijkstra's

We study algorithms to solve the following problems:

- 1) Given a ^{weighted} directed graph G , check if there exists a negative cycle, i.e. $\sum_{e \in C} w_e < 0$.
- 2) If G has no negative cycles, solve the single-source shortest path problem.

Bellman-Ford Algorithm

We use dynamic programming to compute the shortest path distance from $s \in V$ to every $v \in V$ using at most i edges from $i=1$ to $i=n-1$.

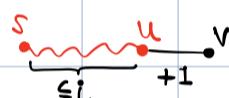
Subproblems: Let $D(v, i)$ be the shortest path distance from $s \in V$ to v using at most i edges.

Answers: For each $v \in V$, distance $D(v, n-1)$ is shortest from s to v , when G has no negative cycles.

Base case(s): $D(s, 0) = 0$ and $D(v, 0) = \infty$ for all $v \in V \setminus \{s\}$.

Recurrence: To compute $D(v, i+1)$, a path with $\leq i+1$ edges must be an extension of a path with $\leq i$ edges

from s to u , for some in-neighbour u of v .



\Rightarrow look for all paths $P_{s \rightarrow u}$ with $\text{len}(P_{s \rightarrow u}) \leq i$ to find the best candidate for extension to v .

We get the recurrence relation $D(v, i+1) = \min_{\substack{\text{existing path} \\ u: u \in V}} \{D(u, i) + w_{uv}\}$

Time complexity: Given $D(v, i)$ for all $v \in V$, takes $\text{in-deg}(v)$ to compute $D(v, i+1)$

\Rightarrow computing $D(w, i+1)$ for all $w \in V$ takes $O(\sum_{w \in V} \text{in-deg}(w)) = O(m)$ time.

\Rightarrow takes $O(mn)$ time to compute for all $1 \leq i \leq n-1$.

Space: naively $\Theta(n^2)$, reduce to $O(n)$ by noticing we only need $D(v, i)$ to compute $D(w, i+1)$.

$\forall v \in V$ $\forall w \in V$

$$\text{dist}(s) = 0, \quad \text{dist}(v) = \infty \quad \forall v \in V \setminus \{s\}$$

for $i=1$ to $i=n-1$:

for $u \in E$:

if $\text{dist}(u) + w_{uv} < \text{dist}(v)$:

$$\text{dist}(v) = \text{dist}(u) + w_{uv}$$

$$\text{parent}(v) = u$$

Bellman-Ford algorithm

Shortest Path Tree As in Dijkstra's, we would like to return a shortest path $s \rightarrow v$ by following edges

$(\text{parent}(v), v) \in E$. But Bellman-Ford has many outer-loop iterations, and it is not clear whether the edges

form a tree. In fact, the path $(\text{parent}(v), v)$ may contain cycles. We show these cycles must be negative.

Lemma: If there is a directed cycle C in edges $(\text{parent}(v), v)$, then C must be negative ($\sum_{e \in C} w_e < 0$).

Proof: Let C be the directed cycle, and write $C = v_1 \dots v_k v_1$ with $v_i, v_{i+1} \in E$.

Assume that $v_k v_1$ is the last edge formed by the Bellman-Ford algorithm in C , so the cycle C is formed when $\text{parent}(v_i) \leftarrow v_k$, while $\text{parent}(v_i) = v_{i-1}$ already for $2 \leq i \leq k$.

Consider values $\text{dist}(v_i)$ right before v_k becomes the parent of v_i .

Since $v_{i-1} = \text{parent}(v_i)$ we have $\text{dist}(v_i) \geq \text{dist}(v_{i-1}) + w_{v_{i-1} v_i}$ for $2 \leq i \leq k$.

Note we cannot have $\text{dist}(v_i) < \text{dist}(v_{i-1}) + w_{v_{i-1} v_i}$ since otherwise $\text{parent}(v_i)$ would be updated.

When we set $\text{parent}(v_i) = v_k$ we must have $\text{dist}(v_i) > \text{dist}(v_k) + w_{v_k v_i}$

$$\begin{aligned} \Rightarrow \sum_{i=1}^k \text{dist}(v_i) &= \sum_{i=2}^k \underbrace{\text{dist}(v_i)}_{\geq \text{dist}(v_{i-1}) + w_{v_{i-1} v_i}} + \text{dist}(v_1) \geq \sum_{i=1}^{k-1} \text{dist}(v_i) + \sum_{i=2}^k w_{v_{i-1} v_i} + \underbrace{\text{dist}(v_1)}_{> \text{dist}(v_k) + w_{v_k v_1}} \\ &> \sum_{i=1}^k \text{dist}(v_i) + \sum_{e \in C} w_e \end{aligned}$$

$\Rightarrow \sum_{e \in C} w_e < 0$ so C is a negative cycle. \square

Thus if G has no negative cycles, then edges $(\text{parent}(v), v)$ have no directed cycles.

G connected \Rightarrow every $v \in V$ has exactly one incoming edge in $(\text{parent}(v), v)$ and \nexists directed cycles

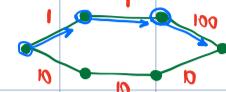
\Rightarrow edges $(\text{parent}(v), v)$ form a directed tree moving away from $s \in V$.

Greedy Algorithms

When does making the locally optimal choice at each stage result in a globally optimal solution?

In dynamic programming, we perform an exhaustive search on all subproblems, and may reconsider the previous stage's solutions to improve future stages.

Anti-example Pathfinding.



Locally optimal choices are bad.

without backtracking

not necessarily unique!

We say a problem exhibits the **greedy choice property** if locally optimal choices lead to a globally optimized solution.

Often, we express these problems recursively.

To prove that a problem satisfies the **greedy property**, we:

- 1 Show an optimal solution exists and define it
- 2 Show that the (greedy) recursive problem decomposition generates the optimal solution.

A Fractional Knapsack Problem: Recall **DP D Knapsack**,

Input: n items of weight w_i and value v_i , and some $W \in \mathbb{N}_0$.

Output: a subset $S \subseteq \{1, \dots, n\}$ with $\sum_{i \in S} w_i \leq W$ that maximizes $\sum_{i \in S} v_i$.

What if we are allowed to take fractions $p_i \in [0, 1]$ of each item such that $\sum_{i \in S} p_i v_i$ is maximized under $\sum_{i \in S} p_i w_i \leq W$?

This is the fractional knapsack problem. **Intuitively, filling the knapsack with highest value:weight ratio is best.**

Define densities $d_i = (\frac{v_i}{w_i})$ for $i \in \{1, \dots, n\}$ and sort $D = [d_i]_i$ in decreasing order.

Then choose high-density elements until we no longer can.

$$D = [\left(\frac{v_i}{w_i}\right)]_i \quad O(n)$$

$$D = \text{sorted}(D) \quad O(n \log n)$$

Collect elements from D until adding one more would be overweight } $O(n)$
add fractional part of that last element

$\Rightarrow O(n \log n)$ time, $O(1)$ space.

* HARD (try!) Find an $O(n)$ algorithm for the problem that adopts a weighted median algorithm

Remark: The proof for greedy property we use is by contradiction and is constructive.

We can also formulate a recursive proof:

- 1 Identify recursive structure
- 2 Prove that optimal structure preserves optimal substructure

Claim: The fractional knapsack problem has the **greedy choice property**.

Proof: Say w_i, v_i are sorted such that $\frac{v_i}{w_i} > \frac{v_{i+1}}{w_{i+1}}$.

Let $G = (p_1, \dots, p_n)$ be the fractions produced greedily, and

$\text{OPT} = \{G' = (q_1, \dots, q_n) \mid G' \text{ optimal}\}$. We claim $G \in \text{OPT}$, SFAC mt.

By assumption, $\sum q_i v_i > \sum p_i v_i$ and so $(p_i)_i \neq (q_i)_i$.

Let i be the first index such that $q_i \neq p_i$.

By the design of our algorithm, we must have $p_i > q_i$. $\boxed{G \text{ takes a larger proportion of this higher density element}}$

By optimality of G' we must have some $j > i$ with $p_j < q_j$. $\rightarrow G'$ takes larger prop. of some lower density elem.

Idea: Define G' that takes slightly more i and less j . show $G' > G$.

Say $G' = (q'_1, \dots, q'_n)$ where $q'_k = q_k$ for $k \neq i, j$.

Define $q'_i = q_i + \epsilon$ and $q'_j = q_j - \epsilon \frac{w_i}{w_j}$.

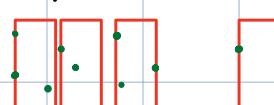
$$\text{Then } \sum_{k=1}^n q'_k v_k = \sum_{k \neq i, j} q_k v_k + (q_i + \epsilon)v_i + (q_j - \epsilon \frac{w_i}{w_j})v_j$$

$$= \sum_{k \neq i, j} q_k v_k + q_i v_i + q_j v_j + \epsilon v_i - \epsilon \frac{v_j}{w_j} w_i > \sum_{k=1}^n q_k v_k$$

$$\text{and } \sum_{k=1}^n q'_k w_k = \sum_{k=1}^n q_k w_k < W$$

\rightarrow we have contradicted optimality of G' . $\Leftrightarrow G \in \text{OPT}$. \square

B minimum-rectangle-cover: some array $A = ((x_i, y_i))$ of points, some $W \in \mathbb{R}_{>0}$ width. Find minimum number of rectangles of width $\leq W$ that covers A .



We can use a greedy algorithm here.

We claim that this algorithm produces the correct result.

$A = \text{sort } A \text{ by } x\text{-value}$
cover the first point with rectangle R of width w
"remove" all points covered by R
do recursively

\Rightarrow time $O(n \lg n)$, space $O(1)$.

Substructure: Consider the decomposition $A = (A_1, \dots, A_k) \cup (A_{k+1}, \dots, A_n)$ such that $x_k - x_1 \leq w$ but $x_{k+1} - x_1 > w$.

Then $\text{min-rect-cover}(A) = 1 + \text{min-rect-cover}((A_{k+1}, \dots, A_n))$.

Proof: SFAC not true, so $m(A) \neq 1 + m(A')$. If $m(A) < 1 + m(A')$ then $m(A) \leq m(A')$ but if $x_{k+1} - x_1 > w$ then this cannot be true.

If $m(A) > 1 + m(A')$ we can construct a cover taking $m(A')$ and adding a rectangle covering $(A_1, \dots, A_k) \Rightarrow m(A) \leq 1 + m(A')$. ↗

Therefore the recursive substructure holds. \square

Since the algorithm uses this substructure decomposition, we have proven correctness.

C min-halbe-sum-operations

* finish greedy algos.

Amortized Analysis

So far, we have been looking at static problems: solve a problem of size n .

eg. heap use cost
in bitmaps

We want to shift our focus to general sequences of operations, where we aim to analyze the cost per operation.

Instead of being given n items, we might want to analyze the cost of n consecutive search/insert/delete operations.

The "amortized cost" over n operations is roughly

$$\frac{\sum_{i=1}^n \text{cost(op}_i\text{)}}{n} \quad \left. \begin{array}{l} \text{total cost} \\ \text{operations} \end{array} \right\} \text{our actual definition will be slightly more flexible}$$

Aggregate Method

A std::vector dynamic resizing (aggregate method)

"weakest" but most intuitive

In C++, we know the std::vector structure provides dynamic arrays. How can we perform amortized analysis?

1 resizing: Say our resizing strategy when pushing into an array of size n was to: 1 allocate size $n+1$ array
2 copy old elements into new

Starting with an empty array and performing n insertions, we have

$$C_{Am} = \frac{1}{n} \sum_{i=0}^{n-1} i = \frac{1}{n} \cdot \frac{n(n-1)}{2} = \frac{n-1}{2} \quad \text{amortized cost per operation}$$

of course, this resizing strategy is bad.

double resizing: Instead, we double array size whenever we reach capacity.

Let $k \in \mathbb{N}$ be such that $2^k < n \leq 2^{k+1}$, so our array has capacity 2^{k+1} . Then we have

$$C_{Am} = \frac{1}{n} \sum_{i=0}^k 2^i = \frac{1}{n} \left(\frac{2^{k+1}-1}{2-1} \right) \xrightarrow{\text{drop } -1} \frac{1}{n} 2^{k+1} = \frac{1}{n} 2 \cdot 2^k < \frac{1}{n} 2n = 2 \quad \text{amortized cost per operation}$$

Clearly, this strategy is better.

* Exercise: Hash tables have $O(1)$ expected amortized cost with doubling strategy.

But why do we need amortized analysis?

B Bitmask (motivating example): Consider the algorithm:

* Exercise: amortized cost of operations

```
increment(A):
    i ← 0
    while i < len(A) and A[i] = 1:
        A[i] = 0 ← flip the bit 1 → 0
        i++
    if i < len(A):
        A[i] = 1 ← flip 0 → 1
```

What if we have mixed sequences of operations? Amortized bounds: assign an amortized cost per operation type,

"preserve the sum" over all operations.

$$\sum_{\text{op.}} \text{amortized cost} \geq \sum_{\text{op.}} \text{actual cost}$$

worst upper bound on actual cost

we define amortized cost such that we form information about total cost

an amortized bound somehow

C binary trees: consider create-empty/insert/delete on binary trees.

Recall: create-empty: $O(1)$ worstcase
insert/delete: $\Theta(\log n)$

so say we have $c, i, d \in \mathbb{N}$ of each operation, starting empty.

Notice $d \leq i$ since we can only delete at most as many times as we insert.

$$\Rightarrow \text{total cost } O(c + i \log n + d \log n) = O(c + i \log n) \text{ so amortized bounds}$$

* but n changes so how can we analyze properly? order of operations also matters
(e.g. create → i → i → ... → i → d → d → ... → d expensive, create → i → d → i → ... → d cheap)

create-empty $O(1)$
insert $\Theta(\log n)$.
delete Θ \rightsquigarrow dub with insert

Accounting Method

Imagine you have a bank account, and that every operation can store credit in that account (always ≥ 0 balance).

- allow operations to add credit to account.
- allow operations to pay for time using stored credit.

{ eg. binary tree insertion adds $\log n^*$ money to account
deletion removes $\log n^*$ money from account}

* Ex. show valid amortized bound

- * Calculate amortized operation cost = actual operation cost + deposits - withdrawals \rightarrow we can form amortized bound.

C (continued): we claim that $\text{insert} = \Theta(\log n)$ and $\text{delete} = \Theta$ amortized.

Structure: add 1 coin worth $\Theta(\log n^*)$ per insert
consume 1 coin for each delete } amortized cost of operation = actual cost + deposits - withdrawals

Generalize: say n is the current structure size.

Insert: deposit 1 coin worth $\log n$, delete : withdraw one coin worth $\log n$.
 \Rightarrow amortized $\log n + \log n - 0$ \Rightarrow amortized $\log n + 0 - \log n$
 $\in \Theta(\log n)$ amortized

Invariant: a structure of size n always consists of an account A with $\{\log 1, \log 2, \dots, \log(n)\} \subseteq A$.

* include $\log n$ or not?

Proof: Induction. True for $n=0$. Say true for n .

If we insert, then, $A \leftarrow A \cup \{\log(n+1)\}$ so holds. If we delete, then $A \leftarrow A \setminus \{\log(n)\}$ holds. \square

D hash-table doubling (insertion-only)

When we insert an item, put a coin worth $C = \Theta(1)$ on the item.

When doubling, use existing coins to pay for doubling operation.

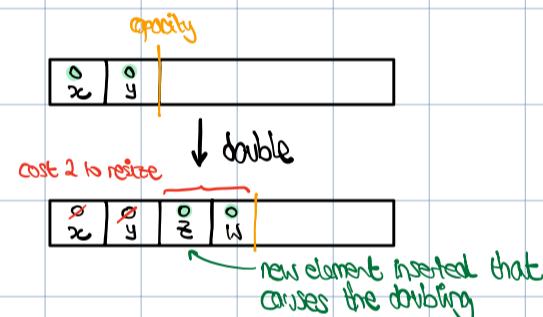
lost $\frac{1}{2}$ items have coins *

\Rightarrow table has at most half coins (at most $\frac{n}{2}$ coins).

amortized doubling: $\Theta(n) - C \frac{n}{2} = 0$ for c large enough constant

actual withdrawals

* we would have to prove some spatial property/invariant
to prove that each operation charged at most once



Charging Method

- allow operations to charge cost retroactively to the past (not future).

- * Calculate amortized operation cost = actual cost - total charge to past + total charge in future.

* Ex. show valid amortized bound

D hash-table doubling (insertion-only)

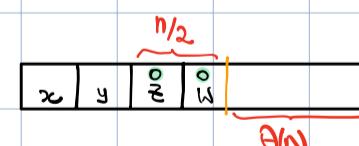
For every doubling operation, charge all the insertions made since the last doubling.

\leadsto uniformly charge each $\Theta(1)$.

\rightarrow key: each insert is only charged at most once over all time

\Rightarrow amortized insert = $\Theta(1) - 0 + \Theta(1) \in \Theta(1)$

amortized double = $\Theta(n) - \frac{1}{2}\Theta(1) + 0 \in \Theta(0)$ zero cost.



D hash-table doubling and halving (insertion and deletion): double at capacity, halve at $\frac{1}{4}$ capacity

\Rightarrow after any doubling or halving, our structure is $\frac{1}{2}$ full. Say $m = \Theta(n)$ capacity, where n items.

Suppose we are at the "nice" equilibrium state (50% full). Perform some sequence of insert/delete.

At some point, we reach a state of either 25% or 100%.

25%: we must have deleted 25% of items, i.e. $\geq \frac{m}{4}$ deletions.

100%: we must have inserted 50% of items, i.e. $\geq \frac{m}{2}$ insertions.

\Rightarrow charge halving to $\geq \frac{m}{4}$ deletions since last resize so charge $\Theta(1)$ per op. } each resize only charges part time window ✗
 charge doubling to $\geq \frac{m}{2}$ insertions since last resize so charge $\Theta(1)$ per op. } of operations made (so each op charged at most once)

* unlike accounting, we only prove a temporal property/invariant here.

\Rightarrow each op. charged at most once is trivial to observe

Potential Method

idea is to smooth out the analysis of some algorithm

Define a potential function Φ mapping data-structure configuration $\rightarrow N_{\geq 0}$. Typically we use size $\in N_{\geq 0}$

Therefore $\Phi: N_{\geq 0} \rightarrow N_{\geq 0}$. We also constrain Φ to satisfy $\Phi(0) = 0$.

$n \mapsto \Phi(n)$

empty structure potential

potential difference $\Delta\Phi(n)$

* calculate amortized operation cost = actual cost + $(\Phi(n) - \Phi(n-1))$.

$\Phi(\text{after op}) - \Phi(\text{before op})$

\Rightarrow amortized cost of n operations = $\sum_{i=1}^n$ amortized cost of i^{th} operation

valid amortized bound

$$= \sum_{i=1}^n \text{actual } i^{\text{th}} \text{ cost} + \sum_{i=1}^n (\Phi(i) - \Phi(i-1))$$

$$= \text{actual } n \text{ operations cost} + \Phi(n) - \Phi(0) \geq \text{actual } n \text{ operations cost}.$$

Choosing a good Φ is hard. In general, we want expensive i^{th} op. to have $\Delta\Phi(i) < 0$ to such an extent that

it nullifies the effect of the actual expensive cost.

E. Binary counter

increment(A):
 $\xleftarrow{\text{array of bits}}$
 $i \leftarrow 0$
 while $i < \text{len}(A)$ and $A[i] = 1$:
 $A[i] = 0$ \leftarrow flip the bit $i \rightarrow 0$
 $i \leftarrow i + 1$
 if $i < \text{len}(A)$:
 $A[i] = 1$ \leftarrow flip $0 \rightarrow 1$

e.g. $0011010111 \xrightarrow{\text{increment}} 0011011000$ at most n bits change

In an amortized sense, how many bits change?

We can use the potential method.

Notice an increment has a very dear cost: $O(1 + \#(\text{trailing } 1s))$. We want constant amortized cost.

Try: $\Phi = \# \text{ trailing } 1s$. But then 11110 has cost $O(1) + \underbrace{\Phi(1111)}_{\Theta(n)} - \underbrace{\Phi(11110)}_{O(1)} = \Theta(n)$ amortized bound.
 (won't work)

Define $\Phi(\text{bitstring}) = \# 1\text{-bits}$. Then notice $\Delta\Phi$ is at most 1 since we add at most one 1-bit.

However for expensive operations (e.g. $01111 \rightarrow 10000$) we have $\Delta\Phi$ large negative (expensive ops).

In general, t trailing 1s \Rightarrow increment destroys t 1s, creates one 1, so $\Phi(i) = \Phi(i-1) - t + 1$

$$\Rightarrow \Delta\Phi(i) = -t + 1$$

amortized op. cost = $\underbrace{(1+t)}_{\text{actual cost}} - \underbrace{t+1}_{\Delta\Phi(\text{op})} = 2$ constant amortized.

Ex. Consider the amortized cost of splitting on a 2-3 tree. The "problem cases" are where the tree has too many 3-nodes.

So we define $\Phi(\text{state}) = \# 3\text{-nodes}$.

One strategy to define suitable ϕ is to look for **bad state properties**.

F Stack-1 (Potential)

Consider a stack with three operations:

- push(S, x) $\Theta(1)$
- pop(S) $\Theta(1)$
- multi-pop(S, k) $\Theta(\min(|S|, k))$

} what are amortized operation costs?

Define potential $\gamma(S) = |S|$. Clearly $\gamma(0) = 0$ and $\gamma(S) \geq 0$ so γ is a valid potential function.

$$\text{POP: amortized cost} = \underbrace{\text{actual cost}}_{C_{\text{pop constant}}} + \underbrace{(\gamma(S_{\text{new}}) - \gamma(S_{\text{old}}))}_{\frac{|S_{\text{new}}|}{1} - \frac{|S_{\text{old}}|}{1}} = C_{\text{pop}} + |S_{\text{new}}| - |S_{\text{old}}| = 0.$$

$$\text{PUSH: amortized cost} = \underbrace{\text{actual cost}}_{C_{\text{push constant}}} + \underbrace{(\gamma(S_{\text{new}}) - \gamma(S_{\text{old}}))}_{\frac{|S_{\text{new}}|}{1} - \frac{|S_{\text{old}}|}{1}} = C_{\text{push}} + |S_{\text{old}}| + 1 - |S_{\text{old}}| = 2.$$

$$\text{MULTI-POP: actual cost} + \underbrace{\gamma(S_{\text{new}}) - \gamma(S_{\text{old}})}_{\frac{|S_{\text{old}}| - \min(|S_{\text{old}}, k|)}{\# \text{elements popped}}} = 0.$$

F Stack-2 (Accounting)

Intuitively, elements can only be popped if they were inserted in the past. Thus, charge insertions now for their potential future deletion. Define:

- PUSH: charge \$1, deposit \$1 $\Rightarrow a_{\text{cost}} = \underbrace{\text{actual cost}}_1 + \underbrace{\text{deposits}}_1 - \underbrace{\text{withdrawals}}_0 = \2
- POP: withdraw \$1 $\Rightarrow a_{\text{cost}} = 1 + 0 - 1 = \0 .
- MULTI-POP: withdraw $\$ \min(|S|, k)$ $\Rightarrow a_{\text{cost}} = \min(|S|, k) + 0 - \min(|S|, k) = \0 .

It suffices to show that our bank balance $B \geq 0$ regardless of operation order. *Easy exercise

General Insertion/Deletion Structures

When we have a structure where single-element insertion costs $C_{\text{insert}}(n)$ and deletion costs $C_{\text{delete}}(n)$, we can generally charge the more expensive operation the cost of both.

Things become tricky when the cost of deletion varies (e.g. is context-dependent).

DYNAMIC BFS Distances

Consider directed G with fixed V and dynamic E (initially $E=\emptyset$).

Let $s \in V$ be some fixed source (constant over the sequence of operations). We can only add to E .

After an insertion $(u, v) \in E$ we update:

- 1. The BFS tree rooted at source $s \in V$
- 2. The distance $\text{dist}[x]$ from $s \rightarrow x$ for $x \in V$

Since naive BFS recomputation is slow, we want a better update algorithm.



Efficient BFS Update

Notice that on insertion $u \in E$ has 2 possibilities: 1) $\text{dist}[v]$ reduces. 2) $\text{dist}[v]$ is unchanged.

The new candidate path to v goes through u :

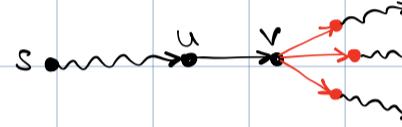
\Rightarrow we update all out-neighbours of v using a directed BFS

```

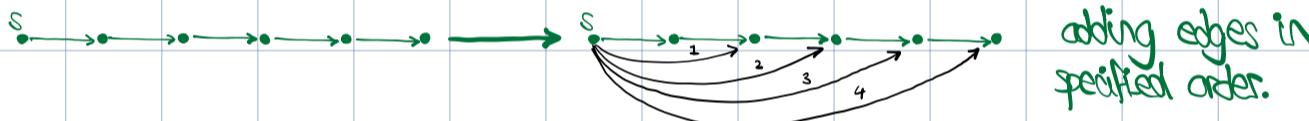
if dist[v] <= dist[u]+1: return ← no change condition
dist[v] ← dist[u]+1
Q = queue(v)
while Q ≠ ∅:
    u = Q.pop()
    for outgoing edge (u, v) ∈ E:
        if dist[v] > dist[u]+1:
            dist[v] ← dist[u]+1
            Q.push(v)
    
```

← update(G) Function

Notice this algorithm may have cost $\Omega(m)$ for particular updates (worst case).



Remark: Notice amortized update cost is not $O(1)$ since we can form an adversarial example:

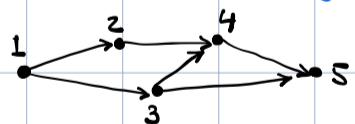


Lemma: G directed as above, $E=\emptyset$. For all $v \in V$ such that $\exists s-v$ path, say $\text{dist}[v] \in \{0, 1, \dots, D\}$ for some $D \in \mathbb{N}$ constant independent of $|V|, |E|$, over m edge insertions. Then $\text{update}(G)$ $O(1)$ amortized.

Incremental Topological Ordering

Let G be directed and acyclic (DAG) such that V is fixed and $|V|=n$.

A topological ordering of G is an enumeration $[v_1, \dots, v_n]$ of V such that $v_i, v_j \in E \Rightarrow i < j$.



In general, topological orderings are not unique.

Suppose $T = (v_1, \dots, v_n)$ is a topological ordering for G we must maintain.

Consider the insertEdge(u, v) operation that:

- 1) adds $u \in E$ (with the guarantee that G remains a DAG)
- 2) if T satisfies $\text{pos}[u] < \text{pos}[v]$ do nothing
- 3) else edge $u \rightarrow v$ violates T , so repair T by reordering.

correctness

Lemma: Upon insertEdge(u, v), calling reorder(u, v, T) maintains a valid top. ordering T .

Proof: Let G be original DAG, and $G' = G + uv$, and T' be the resulting reordered T .

If $\text{pos}[u] < \text{pos}[v]$ then $T' = T$ and T valid on $G \Rightarrow T'$ valid on G' .

$\text{reorder}(u, v, T):$
 $i = \text{pos}[u], j = \text{pos}[v] \text{ so } u \in E \text{ but } i > j$
 $S = \{T[k] \mid j \leq k \leq i\}$
 $R = \{x \in S \mid \exists u \rightarrow x \text{ path in } G\}$ } O(n)
 $\text{construct } T' \text{ such that:}$
 $R \text{ appears after } S \setminus R \text{ in } T'$
 $\text{everything else follows } T$ } O(n)
 $T \leftarrow T'$

reorder algorithm for T

Say $\text{pos}(u) \geq \text{pos}(v)$.

Define $S = \{\tau[k] \mid \text{pos}(v) \leq k \leq \text{pos}(u)\}$ and $R = \{x \in S \mid \exists u-x \text{ path in } G\}$.

For $y \in S \setminus R$ notice $\exists u-y$ path in G , and $\text{pos}(u) \geq \text{pos}(y)$ in $\tau \Rightarrow \exists y-u$ path in G .

\Rightarrow relative ordering of $y \in S \setminus R$ and u does not matter.

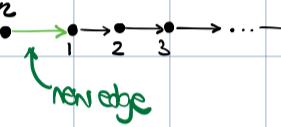
Let $z \in R$ so $\exists u-z$ path in $G \Rightarrow \exists z-y$ path in G (else $\overbrace{u \rightarrow z \rightarrow y}^{\text{worse case } \exists y-z \text{ path}}$).

$\Rightarrow \underbrace{\text{pos}(y)}_{S \setminus R} < \underbrace{\text{pos}(z)}_R$ suffices to satisfy relations between R and $S \setminus R$

Therefore, moving $S \setminus R$ before R in τ' whilst maintaining internal R and $S \setminus R$ orders from τ

defines a valid topological ordering τ' on G . \square

Lemma: $\text{reorder}(w, \tau)$ is worst-case $\mathcal{O}(n)$.

Proof: Considers  requiring $\mathcal{O}(n)$ elements moved. \square

Lemma: Over m operations, the amortized cost of $\text{reorder}()$ is $O(1)$.

Proof: Show that each operation is $O(1)$. \square

Fibonacci Heap

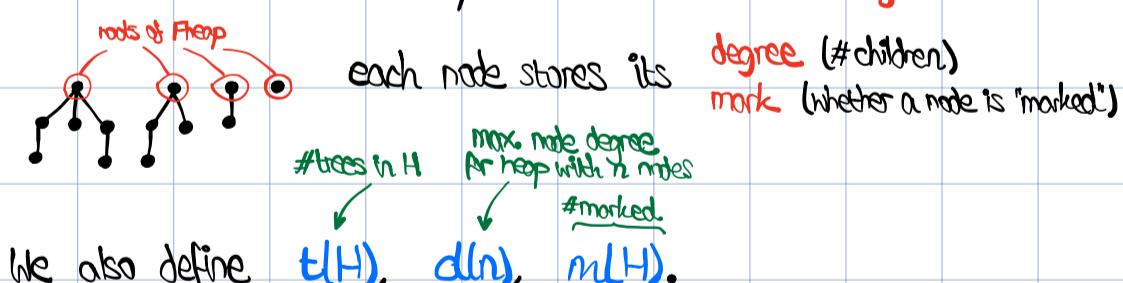
We introduce and analyse the amortized cost for Fibonacci heap operations.

The **min-oriented FibHeap** supports:

| $\text{find-min}()$ | $\text{insert}(x)$ | $\text{delete-min}()$ | $\text{decrease-key}(x, \text{newkey})$ | $\text{union}(H_1, H_2)$ |
|---------------------|--------------------|-----------------------|---|--------------------------|
| $O(1)$ | $O(1)$ | $O(\log n)$ | $O(1)$ | $O(1)$ ← amortized |

Notice the runtime of the operations above is quite different from a binary heap.

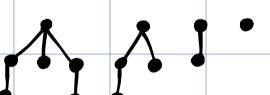
The main idea is an FibHeap H is a **collection of trees** each satisfying the heap property.



Lemma: $d(n) \leq \log_\gamma(n) \in O(\log n)$ where γ is the golden ratio.

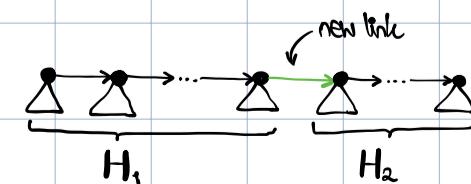
Proof: Beyond scope. \square

find-min(): We can store/memoize the minimum value of $H \Rightarrow O(1)$ find-min.

insert(x): Add x as a new root to H  $\Rightarrow O(1)$ insertion.

union(H_1, H_2): Collect all trees of H_1 and H_2 into a new $H_1 \cup H_2$.

If trees are linked using a LL, this can be $O(1)$. (and update min)



- delete-min**):
1. Delete minimum node and add its child subtrees as new trees in H . $\mathcal{O}(\log n)$ since at most $d(n)$ nodes are children
 2. Combine trees so no roots of trees have the same degree.
 3. Update min.
- \rightarrow Iterate through roots and collect degrees in table D .
Upon collision, merge colliders and update D with new degree.
- $\mathcal{O}(\log n)$ because only $\mathcal{O}(\log n)$ trees after step 2
- } at most $d(n) + t(H)$ roots
 and at most $\mathcal{O}(d(n))$ combines
 $\Rightarrow \mathcal{O}(\log n + t(H))$

\Rightarrow delete-min has worst-case $\mathcal{O}(\log n + t(H))$.

decrease-key(x, K): decrease $\text{key}[x]$ to K where $K \leq \text{key}[x]$.

```

key[x] ← K
y ← parent[x]
if y not None and key[x] < key[y]:
    move x to be a new root
    degree[y]--
    parent[x] = None
    mark[x] = False
    cut x from y
while y not a root:
    if mark[y] is False, set mark[y] = True and exit
    else:
        y is marked → fragile (could be cut)
        z ← parent[y]
        cut y from z (so y becomes a new root)
        y ← z (continue on z)
    update H.min
  
```

marked nodes are nodes that have lost one child

\Rightarrow they are somewhat "fragile"

also notice a node cannot lose >1 child without itself being cut to the root, so nodes cannot lose children too aggressively.

Lemma: The amortized cost for delete-min over k operations is $\mathcal{O}(\log n)$.

Proof: This is an interesting analysis. We use the potential method.

For delete-min, notice { "roots are a mess" that must be cleaned later.
marked nodes are "fragile" and may get cut easily. }

$\Rightarrow t(H)$ and $m(H)$ encode bad state properties.

(One strategy to define suitable ϕ is to look for bad state properties.)

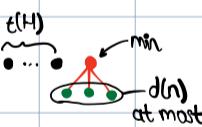
Thus define $\phi(H) = t(H) + 2m(H)$ so clearly $\phi: S \rightarrow \mathbb{N}_{\geq 0}$ and $\phi(0) = 0$.

We aim to analyze $C_{\text{am}} = C_{\text{actual}} + \Delta\phi$ for delete-min.

Notice upon delete-min, we create $\mathcal{O}(d(n))$ new roots $\Rightarrow t(H_{\text{new}}) - t(H_{\text{old}}) \leq d(n)$.

$m(H)$ remains constant at worst case, so $m(H_{\text{new}}) - m(H_{\text{old}}) \leq 0$.

$\Rightarrow \Delta\phi = \phi(H_{\text{new}}) - \phi(H_{\text{old}}) \leq d(n)$.



After deleting the (dangling) min. node, we have at most $t(H) + d(n) - 1$ trees in H_{new} .

Then we combine some-degree trees. Notice upon combining, $t(H)$ reduces by 1 and $m(H)$ at worst remains constant (cannot increase).

Claim: $t(H_{\text{new}}) \in \mathcal{O}(\log n)$ after "combining" some-degree roots.

Notice $d(n) \in \mathcal{O}(\log n)$ and each root has a distinct degree in $[0, \dots, d(n)] \Rightarrow \mathcal{O}(\log n)$ roots. \square

Therefore, if $t(H)$ is large, then $t(H_{\text{new}}) \in \mathcal{O}(\log n) \Rightarrow \Delta\phi$ sufficiently kills $t(H)$.

$$\Rightarrow C_{\text{am}} = \mathcal{O}(\log n + t(H)) + 2(m(H_{\text{new}}) - m(H)) + (t(H_{\text{new}}) - t(H))$$

$$\leq \Theta(\log n) + \Theta(t(H)) + \Theta(\log n) - \Theta(t(H)) \in \Theta(\log n). \quad \square$$

Showing these cancel needs a bit more algorithmic rigour

NP-Hardness, Completeness, and Problem Reduction

Once we have learned more and more algorithms, they become our building blocks and we may not need to design new algorithms from scratch every time. **new problem $\xrightarrow{\text{reduction}}$ known problem**
we are interested in finding problem reductions

Decision Problems

To formalize the notion of a reduction, it is convenient to restrict our attention to decision problems, where the outputs are YES and NO.

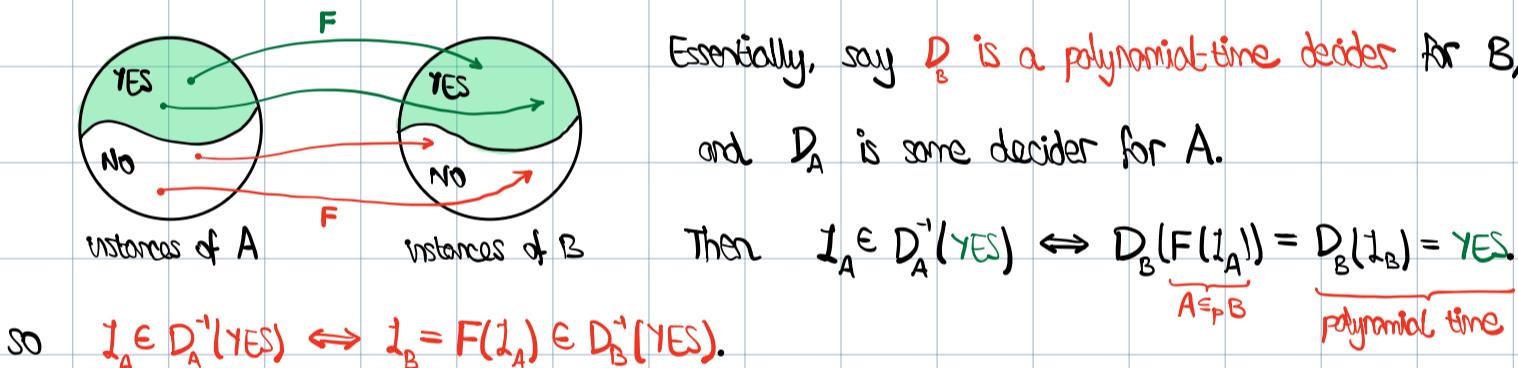
Ex. Finding max. graph matching $\xrightarrow{\text{decision}}$ "does G have a size- k matching?"

For all the problems we will consider, $\xrightarrow{\text{solve decision problem in polynomial time}} \xrightarrow{\text{solve search problem in polynomial time}}$
 \downarrow
use decision algorithm as a subroutine

Let A be a decision problem. We say A is **polynomial-time reducible** to a decision problem B if there is a polynomial-time algorithm F that transforms $I_A \xrightarrow{F} I_B$ such that $I_A = \text{YES} \iff I_B = F(I_A) = \text{YES}$.

We use $A \leq_p B$ to denote the existence of such a reduction.

Intuitively, this says A is not more difficult than B in terms of polynomial-time solvability.



Conversely, say $A \leq_p B$ and it is known that A cannot be solved in polynomial time. Then B cannot either.

Lemma (P-reduction): A and B decision problems. Then $A \leq_p B$ and B polynomial $\Rightarrow A$ polynomial.

Lemma (NP-reduction): A and B decision problems. Then $A \leq_p B$ and A non-polynomial $\Rightarrow B$ non-polynomial.

How can we prove that a problem is not solvable using a polynomial-time algorithm?

Ex. Consider traveling salesman problem TSP, and say this is almost-surely NP.

To show problem C is almost-surely NP, we can show $TSP \leq_p C$.

We will start with the following 3 problems and show they are equivalent in terms of polynomial-time solvability.

1 Maximum Clique (Clique): A subset $S \subseteq V$ is a **clique** if $u, v \in S \Rightarrow u, v \in E$.

Input: $G = (V, E)$ and $k \in \mathbb{Z}$. Output: Is there a clique in G with $\geq k$ vertices?

2 Maximum Independent Set (IS): A subset $S \subseteq V$ is **independent** if $u, v \in S \Rightarrow u, v \notin E$.

Input: $G = (V, E)$ and $k \in \mathbb{Z}$. Output: Is there an IS in G with $\geq k$ vertices? yes or no

3 Maximum Vertex Cover (VC): A subset $S \subseteq V$ is a **vertex cover** if $u, v \in E \Rightarrow \{u, v\} \cap S \neq \emptyset$.

Input: $G = (V, E)$ and $k \in \mathbb{Z}$. Output: Is there a vertex cover for G with $\leq k$ vertices?

Proposition: Clique \leq_p IS and IS \leq_p Clique.

Proof: Clique \rightarrow IS needs us to change edges to non-edges, and IS \rightarrow Clique needs us to change non-edges to edges.

Clique \leq_p IS. To solve clique on G , construct complement \bar{G} and solve IS on \bar{G} .

so $uv \in E \Leftrightarrow uv \notin \bar{E}$. Clearly, the reduction (constructing \bar{G}) is polynomial. Δ

IS \leq_p Clique. Some construction. Δ

Therefore $\{G, k\} \in \text{Clique}'(\text{YES}) \Leftrightarrow \{\bar{G}, k\} \in \text{IS}(\text{YES})$. \square

Lemma: $S \subseteq V$ is a vertex cover for $G \Leftrightarrow V \setminus S$ is independent in G .

Proof: (\Rightarrow) $S \subseteq V$ vertex cover \Rightarrow for $uv \in E$ we have $\{u, v\} \cap S \neq \emptyset$. SFAC $V \setminus S$ not independent.

$\Rightarrow \exists u, v \in V \setminus S$ with $uv \in E$. So $\exists u, v \in E$ with $u, v \in V \setminus S$ so $\{u, v\} \cap S = \emptyset$. $\not\models \Delta$

(\Leftarrow) $V \setminus S$ independent \Rightarrow all $uv \in E$ satisfy $\{u, v\} \cap S \neq \emptyset$ since $u, v \in V \setminus S \Rightarrow uv \notin E$. $\Delta \square$

Proposition: VC \leq_p IS and IS \leq_p VC.

Proof: Since $S \subseteq V$ vertex cover $\Leftrightarrow V \setminus S$ independent, we have that $\exists |S|$ vertex cover $\Leftrightarrow \exists |V \setminus S|$ IS.

can always be made larger

G has a VC of size $\leq k \Leftrightarrow G$ has an IS of size $\geq n-k$.

\Rightarrow the reduction maps $\{G, k\}$ VC $\mapsto \{\bar{G}, n-k\}$ IS. Clearly polynomial time. \square

Lemma (\leq_p Transitive): If $A \leq_p B$ and $B \leq_p C$ then $A \leq_p C$.

Proof: Say $A \xrightarrow{F} B$ and $B \xrightarrow{H} C$. Then $A \xrightarrow{H \circ F} C$ and $H \circ F$ polynomial reduction. \square

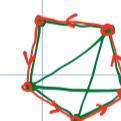
Therefore, Clique, IS and VC are equivalent in polynomial-time solvability.

The above is an example of simple problem reductions. We see more reductions below.

1 Hamiltonian Cycle (HC): A cycle is a **Hamiltonian cycle** if it touches every vertex exactly once.

Input: $G = (V, E)$ undirected.

Output: Does G have a Hamiltonian cycle?



2 Hamiltonian Path (HP): A path is a **Hamiltonian path** if it touches every vertex exactly once.

Input: $G = (V, E)$ undirected.

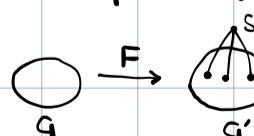
Output: Does G have a Hamiltonian path?



It is not surprising that $HC =_p HP$ but is a good exercise to formalize the reductions.

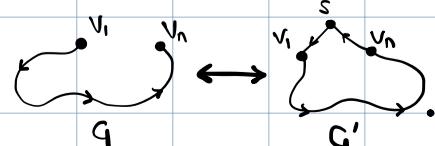
Proposition: $HP \leq_p HC$.

Proof: Given $G = (V, E)$ for HP, construct $G' = (V \cup \{s\}, E')$ by adding a new vertex s that neighbours every $v \in V$.



$\Rightarrow G$ has $HP = v_1, \dots, v_n$, then $\overbrace{s, v_1, \dots, v_n, s}^{EE'} \overbrace{v_1, \dots, v_n, s}^{EE'}$ is a HC in G' .

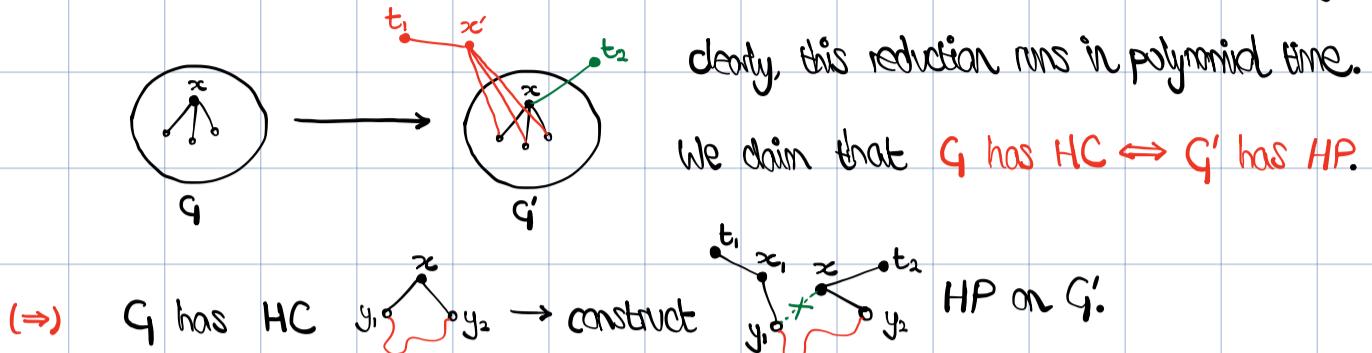
\Leftarrow G' has HC \Rightarrow there is a HC $\overbrace{s, v_1, \dots, v_n, s}^{EE'}$ by choosing to start at $s \in V' \Rightarrow v_1, \dots, v_n$ HP in G' . \square



Remark: It may be tempting to ask why the above $\Rightarrow \text{HC} \leq_p \text{HP}$ directly. Note that we have built a reduction G for HP $\xrightarrow{F} G'$ for HC, so given G for HP, we know how to build G' for HC. However, given G' for HC, we don't yet know how to reduce to G for HP (reverse).

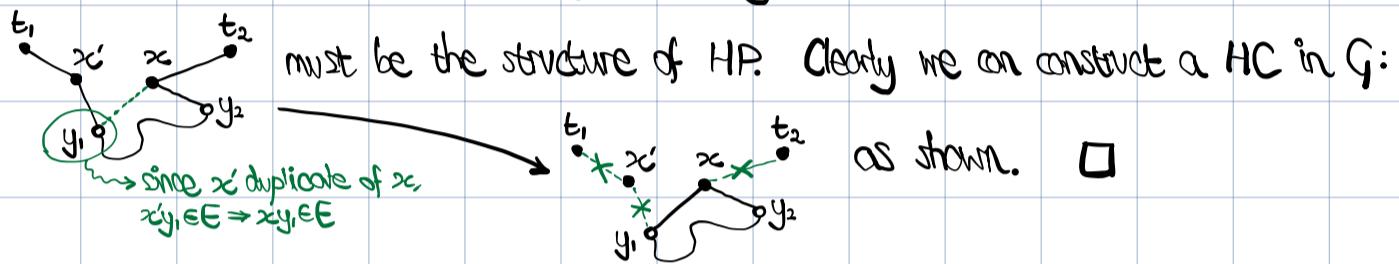
Proposition: $\text{HC} \leq_p \text{HP}$.

Proof: Given $G = (V, E)$ for HC, choose arbitrary $x \in V$ and construct G' by duplicating x and two vertices:



(\Rightarrow) G has HC $\xrightarrow{y_1, y_2}$ construct HP on G' .

(\Leftarrow) G' has HP. Since $t_1, t_2 \in V'$ are degree-1, they must be endpoints of HP and $t_1, x' \in E$.



One technique to do reduction is to show that problem A is a special case of problem B.

This technique is called **specialization**.

Proposition: $\text{HC} \leq_p \text{TSP}$.



Proof: Given $G = (V, E)$ for HC, construct weighted G' such that $G' = (V, E')$ is complete, and for $uv \in E'$ we have $w_{uv} = \begin{cases} 1 & \text{if } uv \in E \\ 2 & \text{if } uv \notin E \end{cases}$. Then G has HC $\Leftrightarrow G'$ solves TSP with a tour of cost n . \square

3-SAT and nontrivial reductions

The 3-SAT problem is important in the theory of NP-completeness.

In this problem we have n boolean variables x_1, \dots, x_n that can each be set to True or False.

We are also given a formula in CNF (Q245) whose literals are either x_i or $\bar{x}_i = \neg x_i$.

$$\underbrace{(x_1 \vee x_2 \vee \bar{x}_3) \wedge (\bar{x}_2 \vee \bar{x}_3) \wedge (\bar{x}_1 \vee x_2 \vee x_3)}_{\text{clause}} \wedge (x_1 \vee x_3). \quad \left. \begin{array}{l} \text{each clause has at most 3 literals} \\ \text{and the formula only has } x_1, x_2, x_3 \end{array} \right\}$$

3-Satisfiability (3-SAT): Input: CNF-formula where each clause has at most 3 literals.

Output: Is there a truth valuation for x_1, \dots, x_n that satisfies all the clauses?

This problem looks very different from the ones seen so far. Finding reductions between seemingly unrelated problems often requires new ideas.

Theorem: $3\text{-SAT} \leq_p \text{IS}$.

Proof: Given a 3-SAT formula, we would like to construct a graph G such that

the formula is satisfiable $\Leftrightarrow G$ has an independent set of a certain size. Let F denote the CNF.

Idea: For each literal in F , construct a new vertex in G .

We want satisfying truth valuation \Leftrightarrow independent set in G .

For F (CNF) to be True, at least one literal in every clause of F must be true,

but we must be consistent in literal valuations across clauses.

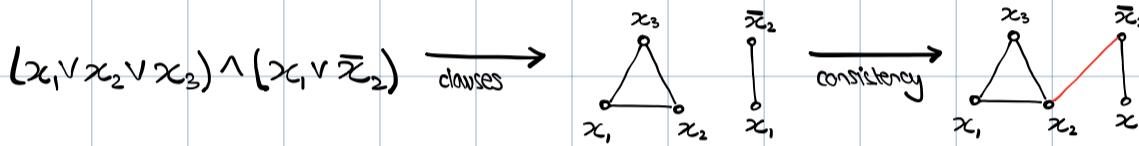
(both x_i and \bar{x}_i cannot be True, exactly one is True)

\Rightarrow for each appearance of x_i and \bar{x}_i , add an edge between them,

so x_i and \bar{x}_i cannot exist in the same independent set.

We also add edges between literals of the same clause \rightsquigarrow the idea here seems to be to reduce redundant "true" valuations where possible

so each clause can have at most one literal in the independent set.



$\Rightarrow G$ has two types of edges: clause-edges and consistency edges.

Claim: Say F has k clauses. Then F satisfiable \Leftrightarrow there is an IS of size k .

The intuition is above: we constructed G such that each clause can have at most one vertex in the IS, and if x_i is in the IS then no other occurrence of x_i can be.

(\Rightarrow) Say there is a satisfying truth valuation. Then at least one literal in each clause is True \Rightarrow choose exactly one per clause to form a set $S = \{x_{a_1}, \dots, x_{a_k}\}$ of size k .

We show S is independent. Suffices to show that $uv \in E \Rightarrow \{u, v\} \notin S$. Say $uv \in E$.

1 uv clause-edge. Then $\{u, v\} \subseteq S$ since at most one vertex per clause is in S .

2 uv consistency-edge. The valuation $\{x_{a_1}, \dots, x_{a_k}\}$ well-defined $\Rightarrow \{x_i, \bar{x}_i\} \notin S$.

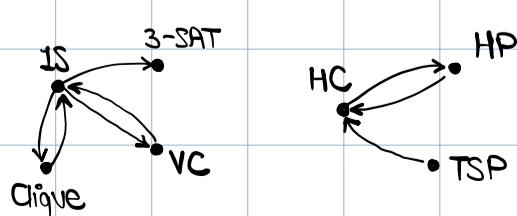
Therefore S independent of size k . \square

(\Leftarrow) Let $S \subseteq V$ be independent of size k . By construction, each clause has exactly one vertex (literal) in S , and $x_i \in S \Rightarrow \bar{x}_i \notin S$ since S independent and $x_i, \bar{x}_i \in E$.

Setting each $x_i \in S$ True gives a truth valuation where each clause of F is True \Rightarrow satisfied. \square

The graph construction is clearly polynomial in n , so $3\text{-SAT} \leq_p \text{IS}$ from the construction above. \square

We have introduced the notion of polynomial time reduction. So far, we have



where $A \rightarrow B$ means $B \leq_p A$.

In principle, we can keep growing this computational web by analysing reductions for more and more problems.

If we define the relation $A =_P B$ to mean $A \leq_P B$ and $B \leq_P A$, we can show that $=_P$ defines an equiv. relation.

What if we could identify the hardest problem X of a large class, so for a new problem Y , it suffices to show that $X \leq_P Y$ to show that Y is at least as hard as others in the class?

NP class

Let A be a decision problem. Determining whether an instance I_A is satisfiable may be hard, but given some "solution" it may be easy to check if the solution satisfies I_A .

Ex. $A = HC$, $I_A = G$ graph. Given a cycle C we can "easily" check if it is Hamiltonian.

Let X be a problem, and represent each instance of X by some binary string s .

Then $X \in NP$ if there is a polynomial time verifier B_X such that

s is a YES-instance \iff there is a proof t of length $\text{poly}(|s|)$ such that $B_X(s, t)$ is YES.

$\underbrace{\text{the problem } s = I_X \text{ is satisfiable}}$ $\xrightarrow{\text{binary string}}$ $\underbrace{(t \text{ typically a solution to } I_A)}$ $\xrightarrow{\text{typically poly-time algorithm to verify that solution } t \text{ "works"}}$

Ex. Consider VC problem. Then

$X = VC$ the problem (vertex cover of size $\leq k$)

$s =$ (binary representation of some graph G) the input string

$t =$ (binary representation of $S \subseteq V$ with $|S| \leq k$)

$B_X(s, t)$ checks if t is a vertex cover for G

} * Exercise: write B_X to show that $VC \in NP$.

* Ex. VC, IS, Clique, HC, HP, subset-sum, 3-SAT $\in NP$.

Remark: Say the problem is to determine if a graph G is non-Hamiltonian (i.e. no Hamiltonian cycles).
Non-example

Whether $nHC \in NP$ is unsolved. We don't know much better than to enumerate all candidate cycles and check.

Remark: The problem of whether a graph has no bipartite matching of size $> k$ can be efficiently validated by checking any vertex cover of size $\leq k \Rightarrow$ a fast verifier exists. Thus $nBM \in NP$ and $BM \in NP$.

A problem X with polynomial verifiers for YES- and NO-instances are in $NP \cap co-NP$.

Clearly, $P \subseteq NP$, where P is the class of problems solvable in polynomial time.

* as long as there is a short solution, a non-deterministic machine will "magically" find the solution.

NP is the class of problems solvable by a non-deterministic polytime (NP) algorithm.

Whether $P = NP$ or $PCNP$ is unsolved.

NP-completeness

X is a "hardest" problem in NP

A problem X is NP-complete if $Y \leq_P X$ for all $Y \in NP$.

Proposition: $P = NP \iff$ some NP-complete problem can be solved in polynomial time. * Trivial by defn.

Once we know one problem is NP-complete, we can use reductions to show that other problems are too.

But how do we prove NP-completeness for that first problem?

Theorem (Cook-Levin): 3-SAT is NP-complete. * We prove this below.

Therefore, since $3\text{-SAT} \leq_p LS =_p VC =_p \text{Clique}$, we have: Corollary: LS, VC and Clique are NP-complete.

To show some problem X is NP-complete, we must show $X \in \text{NP}$ then show $Y \leq_p X$ for some NP-complete Y . The experience of proving $Y \leq_p X$ is often magical. We don't even know what Y to choose to reduce, and then we must find some polynomial reduction of Y into our target X .

Proof: We will prove that 3-SAT is NP-complete. We will introduce some intermediate problems to do so. (Cook-Levin)

Circuit-SAT Input: circuit with AND/OR/NOT gates with some subset of unknown boolean inputs. Output: is there a truth valuation to the unknown subset that outputs True?

We can assume the circuit is a DAG with each AND/OR gate having two incoming edges.

Claim: Circuit-SAT is NP-complete.

Proof sketch: The statement is very abstract, so we must start from definitions. We show $X \leq_p \text{C-SAT} \wedge X \in \text{NP}$.

$X \in \text{NP} \stackrel{\text{def.}}{\Rightarrow} \exists B_x \text{ polytime verifier for YES-instances such that } s \text{ is a YES-instance,}$
 $\exists t \xrightarrow{\text{proof}}$ short such that $B_x(s, t)$ is YES. What is an algorithm?

If B_x runs in $\text{poly}(|s|)$ then there is some machine executing B_x in $\text{poly}(|s|)$ time

\Rightarrow takes at most $\text{poly}(|s|)$ atomic operations to execute B_x

\Rightarrow there is a circuit of size $\text{poly}(|s|)$ implementing B_x .

this reduction is sketchy; we have to return to Turing machines and do a much more tedious reduction here

The main conceptual idea is that a circuit is as general as an algorithm.

The reduction above can be carried out in polynomial time $\text{poly}(|s|)$. } * figure out details

so s YES-instance \Rightarrow some proof t makes $B_x(s, t)$ YES \Rightarrow circuit outputs True.

} \Rightarrow C-SAT reduction

s NO-instance \Rightarrow no input t to the circuit outputs True.

Thus $X \leq_p \text{C-SAT}$, so C-SAT is NP-complete. \triangle

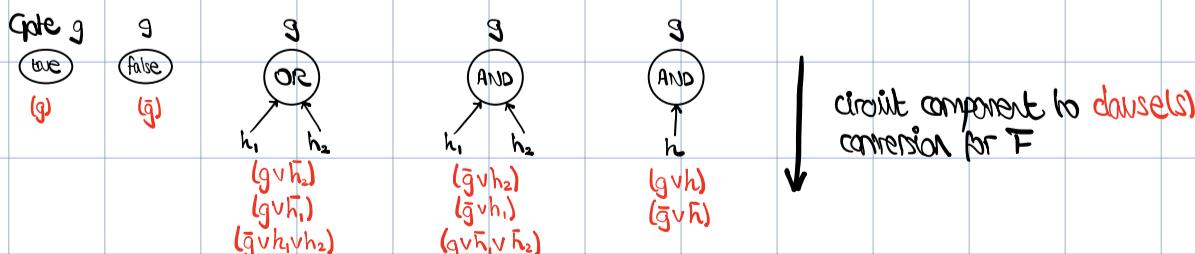
It remains to show that 3-SAT CNFs are as expressive as circuits from C-SAT.

Claim: C-SAT \leq_p 3-SAT.

Proof: Given a circuit of n gates, we want a polytime reduction to a meaningful 3-CNF F for 3-SAT.

let C denote the circuit. For each gate in C , we create one variable in F .

Then, we add clauses to F and modify it as shown:



Finally, to ensure that C outputs True only if C is True, add an output clause (\top) to F .

This is a faithful simulation of C using F , and can be constructed in polytime using local circuit replacement at each construction step. $\triangle \square$

Hard Graph Problems

We use reductions to show that Hamiltonian cycle and graph colouring are NP-complete.

Hamiltonian Cycle We introduce an intermediate problem for the reduction.

Directed Hamiltonian Cycle (DHC)

Input: Directed $G = (V, E)$.

Output: Does G have a directed Hamiltonian cycle?

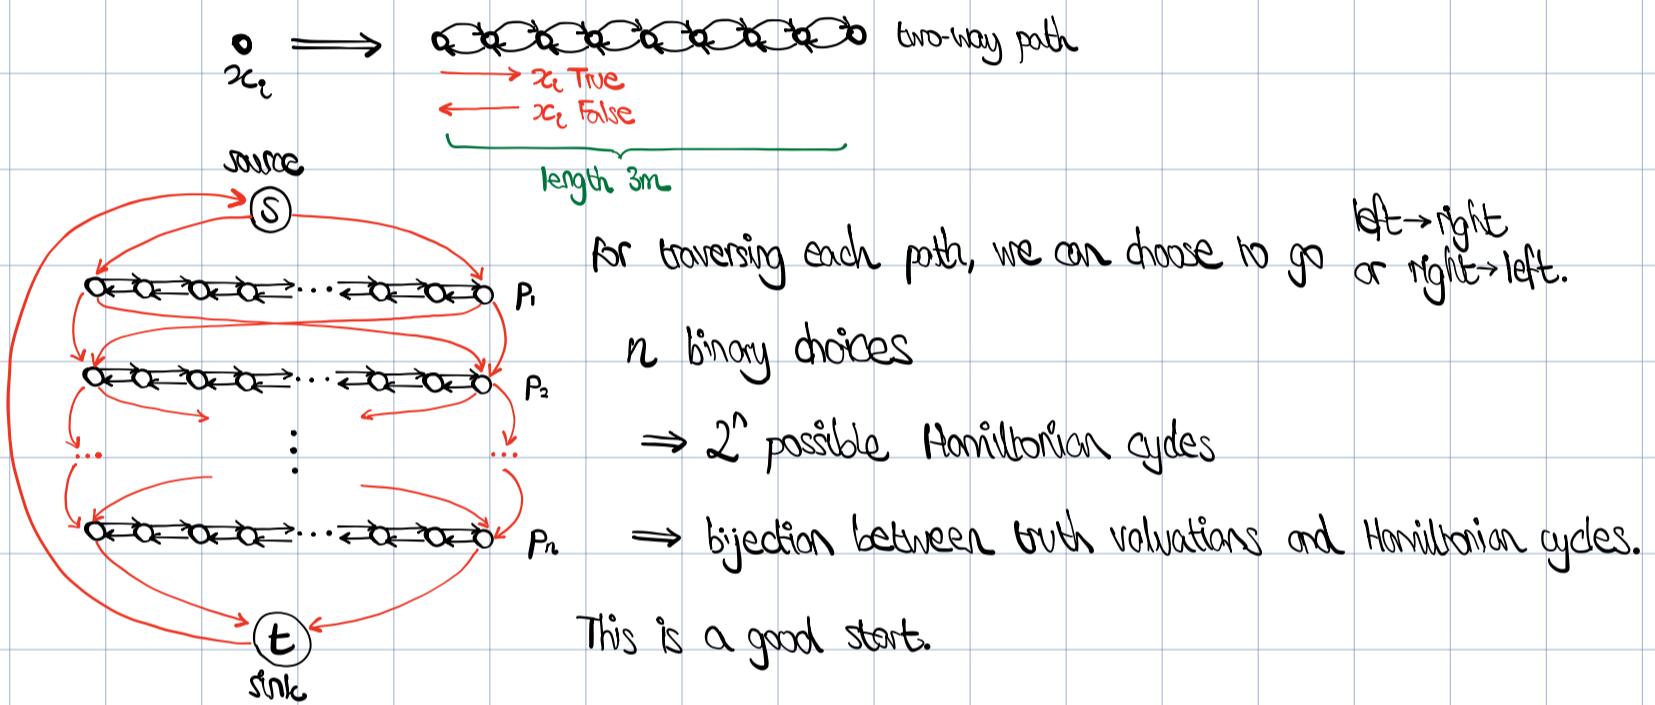
We show $3\text{-SAT} \leq_p \text{DHC} \leq_p \text{HC}$.

Theorem: DHC is NP-complete.

Proof: Clearly $\text{DHC} \in \text{NP}$. We show $3\text{-SAT} \leq_p \text{DHC}$.

Given a 3-SAT instance of n variables x_1, \dots, x_n and m clauses c_1, \dots, c_m we would like to construct directed G , such that F satisfiable $\Leftrightarrow G$ has a Hamiltonian cycle.

Idea: associate a long two-way path to a variable and traveling the path one way corresponds to setting the variable to True, other way corresponds to setting False.



Now, we would like to add some clause structures to "kill" all HCs that do not correspond

to arrangements satisfying F . We decompose each path P_i : 

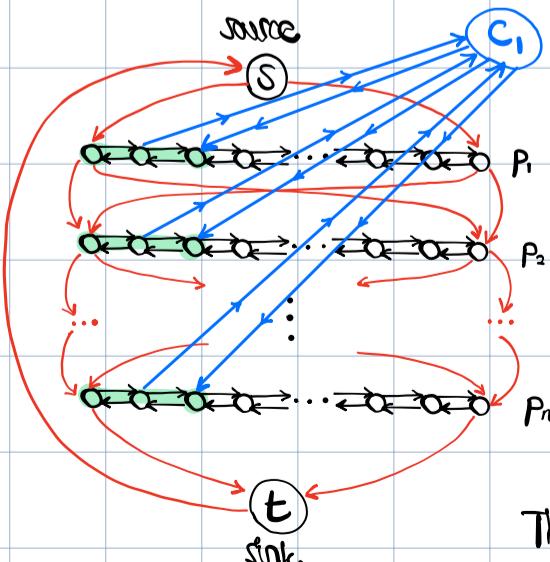
For some clause, e.g. $(x_1 \vee \bar{x}_2 \vee x_3)$, we want at least one of the literals to be True, so

we want HCs to either:

- (1) more L \rightarrow R in P_1
- (2) more R \rightarrow L in P_2
- (3) more L \rightarrow R in P_3

} we want a way to track which clauses our potential HC choice is able to satisfy

\Rightarrow add clause vertices to G , and connect them to variable paths where meaningful.



if x_i appears in clause C_j , connect the corresponding decomposition in P_i to C_j L \rightarrow R.

if \bar{x}_i appears in clause C_j , connect R \rightarrow L.

traverse P_i L \rightarrow R and $x_i \in C_j \Rightarrow C_j$ visited.

traverse P_i R \rightarrow L and $\bar{x}_i \in C_j \Rightarrow C_j$ visited.

That's the whole construction for graph G .

Clearly, G can be constructed in polynomial time.

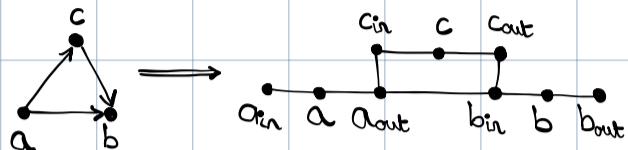
The argument above provides a rough sketch as to prove F satisfiable $\Leftrightarrow G$ has a DHC.

The details of formalizing this claim are left to the reader. \square

To prove $DHC \leq_p HC$, we will use length-3 paths to simulate directed edges.

Proposition: $DHC \leq_p HC$.

Proof: Given directed G for DHC, construct undirected G' for HC using v_{in} and v_{out} vertices for direction.



Clearly, G' can be constructed in polynomial time.

G has a DHC $\Leftrightarrow G'$ has a HC.

(\Rightarrow) Immediately true by construction: follow the cycle.

(\Leftarrow) Start at v_{in} . We must go to v , since otherwise v becomes a dead-end. From v , the only place left to go is v_{out} . Then we must go to some w_{in} by construction. Process continues.

Reconstruct DHC on G based on the above structure by traversing $v \rightarrow w$. \square

Corollary: TSP is NP-complete. Trivial: we proved earlier $HC \leq_p TSP$.

Graph Colouring

Input: Undirected G , and $k \in \mathbb{N}$.

Output: Does G admit a k -colouring?

- $k=1$: True $\Leftrightarrow E = \emptyset$ no edges.
- $k=2$: True $\Leftrightarrow G$ bipartite.
- $k=3$: NP-complete.

graph colouring problem is very useful in resource allocation problems, e.g. interval colouring

Theorem: 3-colouring is NP-complete

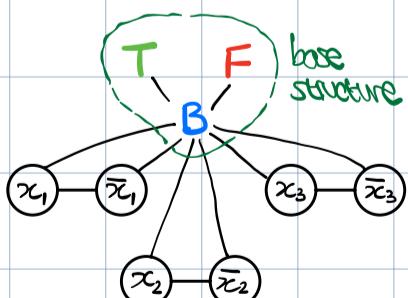
Proof: Clearly $3C \in NP$. We prove $3-SAT \leq_p 3C$. Say 3-SAT instance is CNF F with x_1, \dots, x_n variables and c_1, \dots, c_m clauses.

We want to construct G such that F satisfiable $\Leftrightarrow G$ 3-colourable.

Naturally, let one colour represent True and another represent False. Add $x_i, \bar{x}_i \in E$ to ensure that not both are same truth valuation.

To enforce that x_i and \bar{x}_i get the True/False colours,

connect them to a base vertex structure (since we have 3 colours).

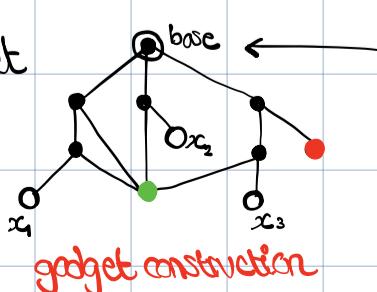


Now, we would like to add some clause structures to "kill"

the 3-colourings that don't correspond to satisfying F .

We want some "gadget" for clauses such that the gadget is 3-colourable \Leftrightarrow at least one of the clause literals are True.

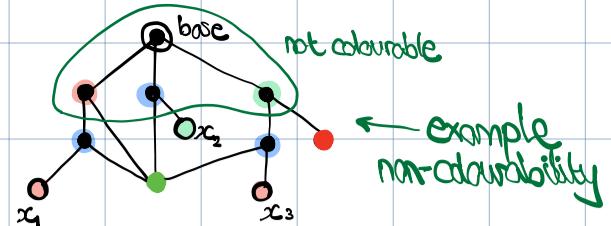
Given $(x_1 \vee \bar{x}_2 \vee x_3)$ we can construct such that \circ base is only colourable if the clause is True.



Note: The colour of \circ base does not matter: the only thing that matters is the colourability of the constructed gadget.

Clearly, this reduction can be done in polynomial time.

By trial-and-error, we can create similar gadgets for other clause configurations.



Remains to prove $F \text{ satisfiable} \Leftrightarrow G \text{ has a 3-colouring}$. Exercise. \square

Σ_3 -SAT Input: 3-CNF F and some $U \subseteq \{0, 1, 2, 3\}$.

Output: Does there exist a truth valuation for F such that $\#\text{True literals in each clause} \in U$ for every clause of F ?

Lemma: Σ_3 -SAT \in NP.

Proof: Say s is an instance, and t is some truth valuation.

Clauses-wise, evaluate literals using valuation t , and verify that $\#\text{True} \in U$.

This verifier is polytime. \square

Lemma: $\{0\}$ -3-SAT is solvable in polynomial time.

Proof: Start with clause c_i of F , assuming $F = \bigwedge_{i=1}^m c_i$ decomposition.

Enforce all literals in c_i to evaluate to False and track the necessary truth valuation in a hashmap.

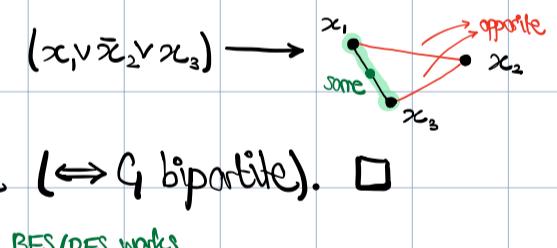
Iterate over all clauses. If there is a collision, then F not $\{0\}$ -3-SAT. Else, True. \square

Lemma: $\{0, 3\}$ -3-SAT is solvable in polynomial time.

Proof: let F be a CNF. Notice for each clause in F , all literals must have the same truth valuation.

Build a signed graph G with $V = \{x_1, \dots, x_n\}$ encoding this constraint:

and complete the graph by growing to include all x_i .



Then G is 2-colourable $\Leftrightarrow F$ $\{0, 3\}$ -3-SAT, solvable in polytime ($\Leftrightarrow G$ bipartite). \square

Theorem: $\{0, 1, 2, 3\}$ -3-SAT is NP-complete.

Proof: We show $3\text{-SAT} \leq_p \{0, 1, 2, 3\}$ -3-SAT. Recall $3\text{-SAT} = \{1, 2, 3\}$ -3-SAT.

Given a 3-SAT instance F with variables x_1, \dots, x_n and clauses C_1, \dots, C_m , we want to find a polynomial reduction to a new CNF F' for $\{0, 1, 2, 3\}$ -3-SAT.

For a clause C_j , define $\gamma(C_j) = \#\text{True literals in } C_j$ under some truth valuation for F .

Intuitively if $\gamma(C_j) \in \{1, 2, 3\}$ for $C_j \in F$, then $\gamma(C'_j) = 3 - \gamma(C_j) \in \{0, 1, 2\}$ where

$$C'_j = (\bar{x}_1 \vee \bar{x}_2 \vee \bar{x}_3) \text{ if } C_j = (x_1 \vee x_2 \vee x_3).$$

Therefore we apply the above transformation for each $C_i \in F$ and form $F' = \bigwedge_{i=1}^m C'_i$. Polytime.

F is $\{1, 2, 3\}$ -3-SAT $\Leftrightarrow F'$ is $\{0, 1, 2, 3\}$ -3-SAT.

(\Rightarrow) Let v be a valuation for F so $\gamma(C_j) \in \{1, 2, 3\}$ for $C_j \in F$. Notice v is well-defined for F' .

Then $\gamma(C_j) \in \{1, 2, 3\} \Rightarrow \gamma(C'_j) = 3 - \gamma(C_j) \in \{0, 1, 2\}$. \triangle Some argument for (\Leftarrow). $\triangle \square$

We consider an extension to the Independent Set (IS) problem. Let G be a graph.

An independent set $I \subseteq V$ is **maximally independent** if it cannot be enlarged to a larger IS.

Extend-MIS-out Input: G graph, $S \subseteq V$ set.

Output: Is there a maximally independent $I \subseteq V$ with $I \cap S = \emptyset$?

Lemma: E-MIS-out \in NP.

Proof: Let (G, S) be a problem instance, and $I \subseteq V$ be a short proof.

- 1 Check $I \subseteq V$ is independent. Polytime from IS.
- 2 Check $S \cap I = \emptyset$. Polytime.
- 3 For all $v \in V \setminus I$ check that $I \cup \{v\}$ is not independent. n^2 polytime.

The above defines a polytime verifier $B_{E\text{-MIS-out}}$. \square

Theorem: Extend-MIS-out is NP-complete.

Proof: We provide a reduction from 3-SAT.

let F be a CNF for 3-SAT with variables x_1, \dots, x_n and clauses C_1, \dots, C_m .

Similar to the reduction $3\text{-SAT} \rightarrow \text{IS}$, construct consistency edges between every x_i, \bar{x}_i occurrence.

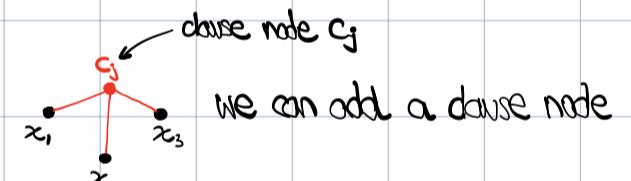
$$(x_1 \vee x_2 \vee x_3) \wedge (\bar{x}_1 \vee \bar{x}_2) \xrightarrow{\text{consistency}} \begin{array}{c} x_3 \\ \circ \\ x_1 \\ \circ \\ x_2 \\ \circ \\ \bar{x}_1 \\ \bar{x}_2 \end{array}$$

In IS, we asked if G had an IS of size m . \nearrow enforce exactly one true literal per clause and have m total clauses

Here, we need a different way to check this size condition.

In particular, each clause needs one True evaluation \Rightarrow

such that if $x_1, x_2, x_3 \in c_j$ then $x_1 c_j, x_2 c_j, x_3 c_j \in E$.



Add clause nodes to S . The above reduction is clearly polytime.

We claim F satisfiable $\Leftrightarrow G$ has a maximal IS outside of S .

(\Rightarrow) F satisfiable \Rightarrow there is some truth valuation such that C_1, \dots, C_m each have at least one True literal. Add all corresponding True nodes to I .

Since each clause c_j has at least one neighbour in I , we have $I \cup \{c_j\}$ dependent for

any $c_j \Rightarrow$ there is a maximal IS for G outside $S = \{C_1, \dots, C_m\}$. \triangle

(\Leftarrow) Let I be a maximal IS for G , and $S = \{C_1, \dots, C_m\}$, so $I \cap S = \emptyset$.

SFAC I does not define a satisfying valuation for F . Then $\exists j$ with c_j False
so no neighbours of c_j (i.e. the clause literals) are True \Rightarrow none are in I .

Therefore $I \cup \{c_j\}$ independent, contradicting maximality outside S . $\nmid \triangle$

Thus $3\text{-SAT} \leq_p E\text{-MIS-out}$, so E-MIS-out is NP-complete. \square

Hard Partitioning Problems

→ generalization of bipartite matching

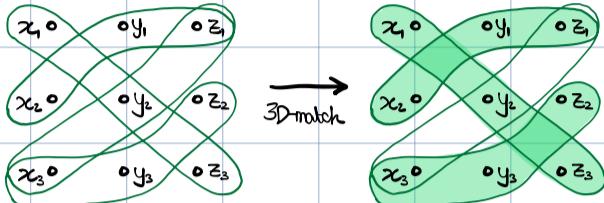
We will show that **3D-matching** and **subset-sum** are NP-complete.

3D-matching Input: Disjoint X, Y, Z each of size n , and some set $T \subseteq X \times Y \times Z$ of triples.

Output: Is there a subset of n triples in T s.t. each $x \in X \cup Y \cup Z$ is contained in exactly one of the triples?

Ex. Say $X = \{x_1, \dots, x_n\}$, $Y = \{y_1, \dots, y_n\}$ and $Z = \{z_1, \dots, z_n\}$. Then $X \cup Y \cup Z = \{x_1, \dots, x_n, y_1, \dots, y_n, z_1, \dots, z_n\}$.

Given T , we want $S \subseteq T$ with $|S|=n$ such that every $w \in X \cup Y \cup Z$ appears in exactly one triple in S .



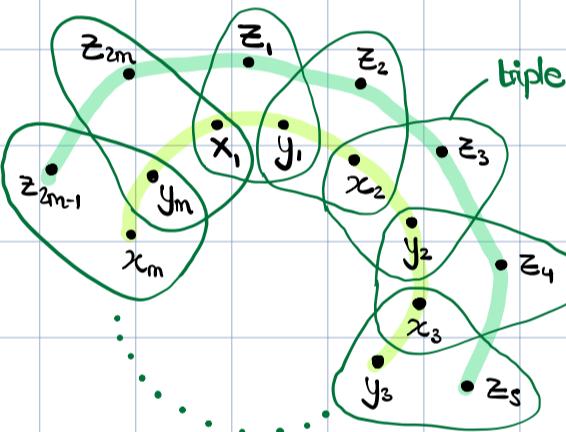
Theorem: **3DM** is NP-complete.

Proof: Clearly **3DM** ∈ NP. We will show that **3-SAT** ≤_p **3DM**.

Given 3-SAT CNF F with variables v_1, \dots, v_n and clauses C_1, \dots, C_m , we want to build a 3DM instance.

As in the Hamiltonian cycle problem, we would like to build some "gadgets" to capture binary decisions.

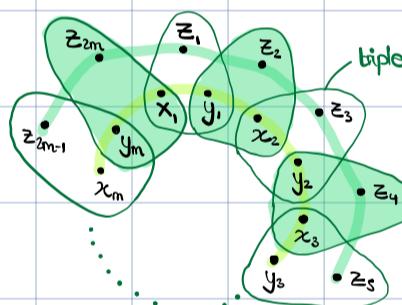
Gadget construction. For each variable v_i , we build the following gadget:



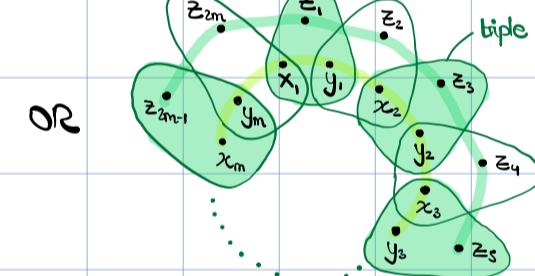
⇒ gadget captures binary decision for v_i .

It remains to add some clause structure.

In our construction, we will ensure that each x_i, y_j are only contained in two triples each, and are not contained in any other triples ⇒ only 2 possibilities for matchings in the gadget



corresponds to $v_i = \text{True}$

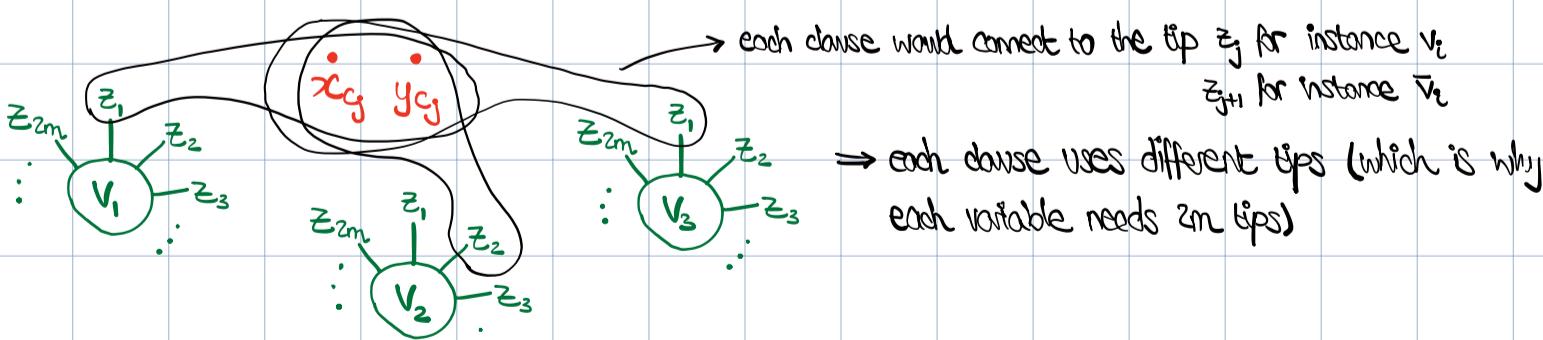


corresponds to $v_i = \text{False}$

We use gadget tips z_i to enforce clause structure. Say $c_j = (v_1 \vee \bar{v}_2 \vee v_3)$.

We create new vertices x_{C_j}, y_{C_j} for c_j . We want x_{C_j}, y_{C_j} to be matched ⇔ c_j evaluates to True.

Thus one of gadgets v_1, v_3 must be True or v_2 must be False.



so clause nodes x_{C_j}, y_{C_j} can only be matched if clause c_j evaluates to True.

For uncovered tips, we create dummy vertices such that these tips are matched.

Total $2nm - m = (2n-1)m$ uncovered tips ⇒ add $(2n-1)m \cdot 2mn < 4m^2 n^2$ dummy triples.

We can construct this in polytime.

F satisfiable \Leftrightarrow there is a perfect 3D-matching.

(\Rightarrow) Say F has a satisfying valuation. If $v_i = T$, then v_i gadget covered using odd tips, else even.

All c_j satisfied \Rightarrow each c_j has a literal satisfied \Rightarrow there will be a tip available for v_i to cover x_{c_j}, y_{c_j} .
other uncovered tips use dummy tips. \triangle

(\Leftarrow) Perfect 3DM \Rightarrow each clause pair x_{c_j}, y_{c_j} covered \Rightarrow reverse-construct truth valuation. \triangle

Therefore $3\text{-SAT} \leq_p 3\text{DM}$ so 3DM NP-complete. \square

Subset-sum

Input: $a_1, \dots, a_n \in \mathbb{N}$ and $k \in \mathbb{Z}$.

Output: Is there a subset $S \subseteq \{1, \dots, n\}$ such that $\sum_{i \in S} a_i = k$?

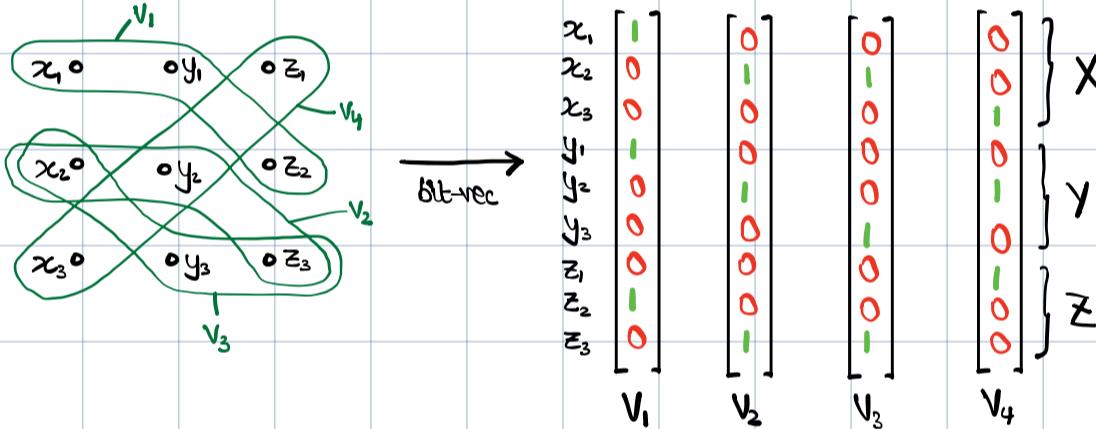
This is a problem about numbers, so some new ideas are needed to connect to previous (combinatorial) problems.

Theorem: Subset-sum is NP-complete.

Proof: Clearly subset-sum \in NP. We show $3\text{DM} \leq_p \text{subset-sum}$.

Given an instance (X, Y, Z, T) of 3DM, we construct an instance of 3DM.

The idea is quite natural: $t \in T \mapsto \text{bit vector} \mapsto \text{number}$. We show an example first.



We essentially create masks for each $t_i \in T$ that identify the X, Y, Z elements in t_i .

Clearly, the construction can be done in polytime.

Denote $V = \{v_1, \dots, v_{|T|}\} \subseteq \mathbb{R}^{3n}$ to be our set of constructed bit-vectors.

Clearly, we see that if $\exists S \subseteq V$ such that $\sum_{v \in S} v = \mathbf{1} \in \mathbb{R}^{3n}$ then each row-sum is 1,

so exactly one of each x_i, y_i, z_i was included \Rightarrow there is a perfect 3DM.

$\Rightarrow \exists S \subseteq V$ with $\sum_{v \in S} v = \mathbf{1} \in \mathbb{R}^{3n} \Leftrightarrow (X, Y, Z, T)$ has a perfect 3D-matching.

We would like to reduce this subproblem to subset-sum.

Won't quite work yet! { Idea: Think of each $v_i \in V$ as the binary representation of some number,

so $(x_i, y_i, z_i) \mapsto 2^i + 2^{i+j} + 2^{i+k}$ so \exists matching \Leftrightarrow some subset sums to $\sum_{i=1}^{3n} 2^i$.

Problem: Above won't work because bits may be "carried" during addition: $\begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix} + \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix} = \begin{bmatrix} 0 \\ 1 \\ 0 \end{bmatrix}$ but

we work around this by working in a large enough base b such that $2^i + 2^i = 2^i \leftarrow \begin{bmatrix} 0 \\ 1 \end{bmatrix}$

carrying can never happen.

Say there are at most $m = |T|$ vectors constructed. Choose base $b = m+1$ so a carry never occurs.

Thus we define our reduction $\gamma: T \rightarrow \mathbb{Z}$ and define $k = \sum_{i=1}^{3n} b^i$ the target.

1 $\gamma(T)$ is polytime.

Recall we define time complexity in terms of **input size in bits**, not as $n = |X| = |Y| = |Z|$.

Thus although m may be exponential in n , the input size is $\mathcal{O}(m)$ since T is part of the input.

For $t \in T$, clearly $\gamma(t)$ is polytime, so $\gamma(T)$ polytime.

2 (X, Y, Z, T) has a perfect 3D-matching \Leftrightarrow there is $S \subseteq \gamma(T)$ with $\sum_{s \in S} s = k$.

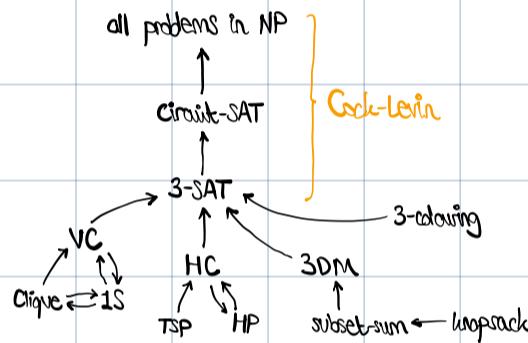
(\Rightarrow) Proven above. Δ

(\Leftarrow) Each position i in the base- b representation records the number of times we cover i -th vertex in the 3DM instance. All positions 1 base- b \Rightarrow perfect 3DM. Δ

Therefore $3DM \leq_p$ subset-sum, so subset-sum is NP-complete. \square

Corollary: Knapsack is NP-complete.

So far, we have built the following relations:



It is an interesting phenomenon to observe that 2 is usually easy but 3 is hard.

2-SAT is easy, but 3-SAT is hard. 2-colouring is easy but 3-colouring is hard. 2DM easy, 3DM hard.

Usually "2" is easy since once a decision is made, the problem space becomes much smaller (fixed).

However, with "3", we are still left with a binary decision.

Decision v.s. Search Problems

Does restricting our scope to decision problems limit our characterization of NP-completeness?

Many search problems can be easily reduced to decision problems. We briefly formalize this.

A problem X is said to be **self-reducible** if $X_{\text{search}} \xrightarrow{F} X_{\text{decision}}$ for polytime reduction F .

Theorem: 3-SAT is self-reducible.

\rightarrow polytime in the runtime of the solver for X_{decision}

Proof: Let F be a 3-CNF with variables x_1, \dots, x_n and clauses C_1, \dots, C_m . Say γ is a solver for $3\text{-SAT}_{\text{decision}}$.

1) Set $x_1 = 0$ and obtain new formula F_1 . Check $\gamma(F_1)$.

If $\gamma(F_1)$ True, recursively find assignments for x_2, \dots, x_n that satisfy F_1 .

2) Else, F_1 was not satisfiable, so try with $x_1 = 1$ to obtain F_2 . Check $\gamma(F_2)$.

If $\gamma(F_2)$ True, recursively find assignments for x_2, \dots, x_n that satisfy F_2 .

3) F_1 and F_2 not satisfiable \Rightarrow no assignment for x_i satisfies $F \Rightarrow F$ not satisfiable.

The above algorithm makes at most 3^n (polynomially many) calls to $\gamma \stackrel{\text{def}}{\Rightarrow} 3\text{-SAT self-reducible}$. \square

Theorem: HC is self-reducible.

Proof: let G be an instance for HC, and γ be a solver for HC decision. Consider the following algorithm:

for $e \in E$:

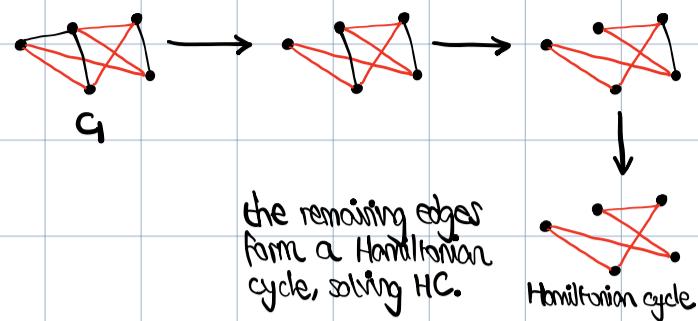
define $G' = G - e$ (remove the edge)

check if $\gamma(G')$ is True

$\gamma(G')$ True \Rightarrow update $G = G'$ and continue

$\gamma(G')$ False \Rightarrow skip and continue (since e must be in all HCs of G)

stop when G has $\leq |V|$ edges remaining



The algorithm calls γ at most $\Theta(|E|)$ times, and $|E| \in \text{poly}(|V|) \Rightarrow \text{HC} \xrightarrow{\text{P}} \text{HC}_{\text{decision}}$. \square

Theorem (Self-Reducibility): Let A be an arbitrary decision problem.

1) $A \in \text{NP} \Rightarrow A \xrightarrow{\text{P}} A_{\text{search}}$. 2) A is NP-complete $\Rightarrow A_{\text{search}} \xrightarrow{\text{P}} A$.

Proof (1): Trivial. To solve A , run A_{search} solver once to obtain a short proof for A , then verify in polytime. Δ

(2): Hard, beyond scope. $\Delta \square$

Therefore, for NP-complete problems, search and decision are identical in terms of solvability complexity.